

Configuration Explained

An A-Z guide to Pacemaker's Configuration Options

, Written by the Pacemaker project contributors

Configuration Explained: An A-Z guide to Pacemaker's Configuration Options

by

Abstract

The purpose of this document is to definitively explain the concepts used to configure Pacemaker. To achieve this, it will focus exclusively on the XML syntax used to configure Pacemaker's Cluster Information Base (CIB).

Copyright © 2009-2019 The Pacemaker project contributors This material may only be distributed subject to the terms and conditions set forth in the GNU Free Documentation License (GFDL), V1.2 or later (the latest version is presently available at <http://www.gnu.org/licenses/fdl.txt>).

Table of Contents

Preface	xi
Document Conventions	xi
Typographic Conventions	xi
Pull-quote Conventions	xii
Notes and Warnings	xiii
We Need Feedback!	xiii
1. Read-Me-First	1
The Scope of this Document	1
What Is <i>Pacemaker</i> ?	1
Cluster Architecture	2
Pacemaker Architecture	2
Node Redundancy Designs	3
2. Cluster-Wide Configuration	5
Configuration Layout	5
CIB Properties	5
Cluster Options	6
3. Cluster Nodes	12
Defining a Cluster Node	12
Where Pacemaker Gets the Node Name	12
Node Attributes	12
Setting and querying node attributes	12
Special node attributes	13
4. Cluster Resources	15
What is a Cluster Resource?	15
Resource Classes	15
Open Cluster Framework	15
Linux Standard Base	16
Systemd	16
Upstart	17
System Services	17
STONITH	17
Nagios Plugins	17
Resource Properties	18
Resource Options	18
Resource Meta-Attributes	19
Setting Global Defaults for Resource Meta-Attributes	21
Resource Instance Attributes	21
Resource Operations	23
Operation Properties	23
Monitoring Resources for Failure	25
Monitoring Resources When Administration is Disabled	25
Setting Global Defaults for Operations	26
When Implicit Operations Take a Long Time	26
Multiple Monitor Operations	26
Disabling a Monitor Operation	27
5. Resource Constraints	28
Scores	28
Infinity Math	28
Deciding Which Nodes a Resource Can Run On	28
Location Properties	28
Asymmetrical "Opt-In" Clusters	30

Symmetrical "Opt-Out" Clusters	30
What if Two Nodes Have the Same Score	31
Specifying the Order in which Resources Should Start/Stop	31
Ordering Properties	31
Optional and mandatory ordering	32
Placing Resources Relative to other Resources	32
Colocation Properties	33
Mandatory Placement	33
Advisory Placement	34
Colocation by Node Attribute	34
Resource Sets	34
Ordering Sets of Resources	35
Ordered Set	35
Ordering Multiple Sets	36
Resource Set OR Logic	37
Colocating Sets of Resources	38
6. Fencing	41
What Is Fencing?	41
Why Is Fencing Necessary?	41
Fence Devices	41
Fence Agents	42
When a Fence Device Can Be Used	42
Limitations of Fencing Resources	42
Special Options for Fencing Resources	42
Unfencing	47
Fence Devices Dependent on Other Resources	47
Configuring Fencing	48
Example Fencing Configuration	48
Fencing Topologies	51
Example Dual-Layer, Dual-Device Fencing Topologies	52
Remapping Reboots	57
7. Alerts	58
Alert Agents	58
Alert Recipients	58
Alert Meta-Attributes	58
Alert Instance Attributes	59
Alert Filters	60
Using the Sample Alert Agents	61
Writing an Alert Agent	61
8. Rules	64
Rule Properties	64
Node Attribute Expressions	64
Date/Time Expressions	66
Date Specifications	67
Durations	67
Example Time-Based Expressions	68
Using Rules to Determine Resource Location	69
Location Rules Based on Other Node Properties	70
Using <code>score-attribute</code> Instead of <code>score</code>	70
Using Rules to Define Options	71
Using Rules to Control Resource Options	71
Using Rules to Control Resource Defaults	72
Using Rules to Control Cluster Options	72
9. Advanced Configuration	74

Specifying When Recurring Actions are Performed	74
Handling Resource Failure	74
Failure Counts	74
Failure Response	75
Moving Resources	76
Moving Resources Manually	76
Moving Resources Due to Connectivity Changes	77
Migrating Resources	80
Tracking Node Health	81
Node Health Attributes	81
Node Health Strategy	81
Measuring Node Health	82
Reloading Services After a Definition Change	82
10. Advanced Resource Types	85
Groups - A Syntactic Shortcut	85
Group Properties	86
Group Options	86
Group Instance Attributes	86
Group Contents	86
Group Constraints	86
Group Stickiness	86
Clones - Resources That Can Have Multiple Active Instances	86
Anonymous versus Unique Clones	87
Promotable clones	87
Clone Properties	87
Clone Options	87
Clone Contents	88
Clone Instance Attributes	89
Clone Constraints	89
Clone Stickiness	91
Clone Resource Agent Requirements	92
Monitoring Promotable Clone Resources	98
Determining Which Instance is Promoted	98
Bundles - Isolated Environments	99
Bundle Prerequisites	99
Bundle Properties	99
Bundle Container Properties	100
Bundle Network Properties	101
Bundle Storage Properties	102
Bundle Primitive	103
Bundle Node Attributes	104
Bundle Meta-Attributes	104
Limitations of Bundles	104
11. Reusing Parts of the Configuration	106
Reusing Resource Definitions	106
Configuring Resources with Templates	106
Using Templates in Constraints	108
Using Templates in Resource Sets	108
Reusing Rules, Options and Sets of Operations	109
Tagging Configuration Elements	110
Configuring Tags	110
Using Tags in Constraints and Resource Sets	111
12. Utilization and Placement Strategy	112
Utilization attributes	112

Placement Strategy	113
Allocation Details	114
Which node is preferred to get consumed first when allocating resources?	114
Which node has more free capacity?	114
Which resource is preferred to be assigned first?	114
Limitations and Workarounds	115
13. Access Control Lists (ACLs)	116
ACL Prerequisites	116
ACL Configuration	116
ACL Roles	116
ACL Targets and Groups	117
ACL Examples	118
14. Status — Here be dragons	121
Node Status	121
Transient Node Attributes	122
Operation History	122
Simple Operation History Example	124
Complex Operation History Example	124
15. Multi-Site Clusters and Tickets	126
Challenges for Multi-Site Clusters	126
Conceptual Overview	126
Ticket	126
Dead Man Dependency	127
Cluster Ticket Registry	127
Configuration Replication	127
Configuring Ticket Dependencies	127
Managing Multi-Site Clusters	128
Granting and Revoking Tickets Manually	128
Granting and Revoking Tickets via a Cluster Ticket Registry	129
General Management of Tickets	130
For more information	130
A. Sample Configurations	131
Empty	131
Simple	131
Advanced Configuration	132
B. Revision History	134
Index	136

List of Figures

1.1. Example Cluster Stack	2
1.2. Internal Components	3
1.3. Active/Passive Redundancy	4
1.4. Shared Failover	4
1.5. N to N Redundancy	4
5.1. Visual representation of the four resources' start order for the above constraints	35
5.2. Visual representation of the start order for two ordered sets of unordered resources	36
5.3. Visual representation of the start order for the three sets defined above	37
5.4. Visual representation of the above example (resources are placed from left to right)	40

List of Tables

2.1. CIB Properties	6
2.2. Cluster Options	7
3.1. Node attributes with special significance	13
4.1. Properties of a Primitive Resource	18
4.2. Meta-attributes of a Primitive Resource	19
4.3. Properties of an Operation	24
5.1. Attributes of a rsc_location Element	29
5.2. Attributes of a rsc_order Element	31
5.3. Attributes of a rsc_colocation Constraint	33
5.4. Attributes of a resource_set Element	35
6.1. Additional Properties of Fencing Resources	43
6.2. Properties of Fencing Levels	51
7.1. Meta-Attributes of an Alert	59
7.2. Environment variables passed to alert agents	61
8.1. Attributes of a rule Element	64
8.2. Attributes of an expression Element	65
8.3. Built-in Node Attributes	65
8.4. Attributes of a date_expression Element	66
8.5. Attributes of a date_spec Element	67
8.6. Attributes of a duration Element	68
9.1. Common Options for a <i>ping</i> Resource	78
9.2. Allowed Values for Node Health Attributes	81
9.3. Node Health Strategies	81
10.1. Properties of a Group Resource	86
10.2. Properties of a Clone Resource	87
10.3. Clone-specific configuration options	87
10.4. Additional colocation constraint options for promotable clone resources	89
10.5. Additional colocation set options relevant to promotable clone resources	90
10.6. Additional ordered set options relevant to promotable clone resources	91
10.7. Role implications of OCF return codes	92
10.8. Environment variables supplied with Clone notify actions	92
10.9. Extra environment variables supplied for promotable clones	94
10.10. XML Attributes of a bundle Element	100
10.11. XML attributes of a docker, podman, or rkt Element	100
10.12. XML attributes of a network Element	101
10.13. Attributes of a port-mapping Element	101
10.14. Attributes of a storage-mapping Element	102
13.1. Properties of an ACL Role	116
13.2. Properties of an ACL Permission	116
13.3. Properties of an ACL Target	117
13.4. Properties of an ACL Group	117
13.5. Properties of an ACL Role Reference	118
14.1. Authoritative Sources for State Information	121
14.2. Node Status Fields	121
14.3. Contents of an lrm_rsc_op job	122

List of Examples

2.1. An empty configuration	5
3.1. Example Corosync cluster node entry	12
3.2. Result of using <code>crm_attribute</code> to specify which kernel <code>pcmk-1</code> is running	13
4.1. A system resource definition	18
4.2. An OCF resource definition	18
4.3. An LSB resource with cluster options	21
4.4. An example OCF resource with instance attributes	21
4.5. Displaying the metadata for the Dummy resource agent template	22
4.6. An OCF resource with a non-default start timeout	23
4.7. An OCF resource with a recurring health check	25
4.8. An OCF resource with custom timeouts for its implicit actions	26
4.9. An OCF resource with two recurring health checks, performing different levels of checks specified via <code>OCF_CHECK_LEVEL</code>	27
4.10. Example of an OCF resource with a disabled health check	27
5.1. Opt-in location constraints for two resources	30
5.2. Opt-out location constraints for two resources	30
5.3. Constraints where a resource prefers two nodes equally	31
5.4. Optional and mandatory ordering constraints	32
5.5. Mandatory colocation constraint for two resources	34
5.6. Mandatory anti-colocation constraint for two resources	34
5.7. Advisory colocation constraint for two resources	34
5.8. A set of 3 resources	34
5.9. A chain of ordered resources	35
5.10. A chain of ordered resources expressed as a set	36
5.11. Ordered sets of unordered resources	36
5.12. Advanced use of set ordering - Three ordered sets, two of which are internally unordered	36
5.13. Resource Set "OR" logic: Three ordered sets, where the first set is internally unordered with "OR" logic	37
5.14. Colocation chain as individual constraints, where A is placed first, then B, then C, then D	38
5.15. Equivalent colocation chain expressed using <code>resource_set</code>	38
5.16. Using colocated sets to specify a shared dependency	39
5.17. Colocation in which the members of the middle set have no interdependencies, and the last set listed applies only to instances in the master role	39
6.1. Obtaining a list of Fence Agent Parameters	48
6.2. An IPMI-based Fencing Resource	51
6.3. Fencing topology with different devices for different nodes	52
7.1. Simple alert configuration	58
7.2. Alert configuration with recipient	58
7.3. Alert configuration with meta-attributes	59
7.4. Alert configuration with instance attributes	59
7.5. Alert configuration to receive only node events and fencing events	60
7.6. Alert configuration to be called when certain node attributes change	60
7.7. Sending cluster events as SNMP traps	61
7.8. Sending cluster events as e-mails	61
8.1. True if now is any time in the year 2005	68
8.2. Equivalent expression	68
8.3. 9am-5pm Monday-Friday	68
8.4. 9am-6pm Monday through Friday or anytime Saturday	68
8.5. 9am-5pm or 9pm-12am Monday through Friday	69
8.6. Mondays in March 2005	69
8.7. A full moon on Friday the 13th	69

8.8. Prevent resource "webserver" from running on node3	70
8.9. Prevent resource "webserver" from running on node3 using rule	70
8.10. A sample nodes section for use with score-attribute	70
8.11. Defining different resource options based on the node name	71
8.12. Change resource-stickiness during working hours	72
8.13. Schedule a maintenance window for 9 to 11 p.m. CDT Sept. 20, 2019	72
9.1. Specifying a Base for Recurring Action Intervals	74
9.2. An example ping cluster resource that checks node connectivity once every minute	78
9.3. Don't run a resource on unconnected nodes	78
9.4. Run only on nodes connected to three or more ping targets.	79
9.5. Prefer the node with the most connected ping nodes	79
9.6. How the cluster translates the above location constraint	79
9.7. A more complex example of choosing a location based on connectivity	79
9.8. The DRBD agent's logic for supporting reload	82
9.9. The DRBD Agent Advertising Support for the reload Operation	83
9.10. Parameter that can be changed using reload	83
10.1. A group of two primitive resources	85
10.2. How the cluster sees a group resource	85
10.3. Some constraints involving groups	86
10.4. A clone that runs a web server on all nodes	88
10.5. Some constraints involving clones	89
10.6. Constraints involving promotable clone resources	89
10.7. Colocate C and D with A's and B's master instances	90
10.8. Start C and D after first promoting A and B	91
10.9. Notification variables	93
10.10. Monitoring both states of a promotable clone resource	98
10.11. Explicitly preferring node1 to be promoted to master	98
10.12. A bundle for a containerized web server	99
11.1. Resource template for a migratable Xen virtual machine	106
11.2. Xen primitive resource using a resource template	106
11.3. Equivalent Xen primitive resource not using a resource template	107
11.4. Xen resource overriding template values	107
11.5. Referencing rules from other constraints	109
11.6. Referencing attributes, options, and operations from other resources	110
11.7. Tag referencing three resources	110
11.8. Constraint using a tag	111
11.9. Equivalent constraints without tags	111
12.1. Specifying CPU and RAM capacities of two nodes	112
12.2. Specifying CPU and RAM consumed by several resources	112
14.1. A bare-bones status entry for a healthy node cl-virt-1	121
14.2. A set of transient node attributes for node cl-virt-1	122
14.3. A record of the apcstonith resource	122
14.4. A monitor operation (determines current state of the apcstonith resource)	124
14.5. Resource history of a pingd clone with multiple jobs	124
15.1. Constraint that fences node if ticketA is revoked	128
15.2. Constraint that demotes rscl if ticketA is revoked	128
15.3. Ticket constraint for multiple resources	128
A.1. An Empty Configuration	131
A.2. A simple configuration with two nodes, some cluster options and a resource	131
A.3. An advanced configuration with groups, clones and STONITH	132

Preface

Document Conventions

This manual uses several conventions to highlight certain words and phrases and draw attention to specific pieces of information.

Typographic Conventions

Four typographic conventions are used to call attention to specific words and phrases. These conventions, and the circumstances they apply to, are as follows.

`Mono-spaced Bold`

Used to highlight system input, including shell commands, file names and paths. Also used to highlight keys and key combinations. For example:

To see the contents of the file `my_next_bestselling_novel` in your current working directory, enter the `cat my_next_bestselling_novel` command at the shell prompt and press **Enter** to execute the command.

The above includes a file name, a shell command and a key, all presented in mono-spaced bold and all distinguishable thanks to context.

Key combinations can be distinguished from an individual key by the plus sign that connects each part of a key combination. For example:

Press **Enter** to execute the command.

Press **Ctrl+Alt+F2** to switch to a virtual terminal.

The first example highlights a particular key to press. The second example highlights a key combination: a set of three keys pressed simultaneously.

If source code is discussed, class names, methods, functions, variable names and returned values mentioned within a paragraph will be presented as above, in `mono-spaced bold`. For example:

File-related classes include `filesystem` for file systems, `file` for files, and `dir` for directories. Each class has its own associated set of permissions.

`Proportional Bold`

This denotes words or phrases encountered on a system, including application names; dialog-box text; labeled buttons; check-box and radio-button labels; menu titles and submenu titles. For example:

Choose `System` → `Preferences` → `Mouse` from the main menu bar to launch `Mouse Preferences`. In the `Buttons` tab, select the `Left-handed mouse` check box and click `Close` to switch the primary mouse button from the left to the right (making the mouse suitable for use in the left hand).

To insert a special character into a `gedit` file, choose `Applications` → `Accessories` → `Character Map` from the main menu bar. Next, choose `Search` → `Find...` from the `Character Map` menu bar, type the name of the character in the `Search` field and click `Next`. The character you sought will be highlighted in the `Character Table`. Double-click

this highlighted character to place it in the Text to copy field and then click the Copy button. Now switch back to your document and choose Edit → Paste from the gedit menu bar.

The above text includes application names; system-wide menu names and items; application-specific menu names; and buttons and text found within a GUI interface, all presented in proportional bold and all distinguishable by context.

Mono-spaced Bold Italic or *Proportional Bold Italic*

Whether mono-spaced bold or proportional bold, the addition of italics indicates replaceable or variable text. Italics denotes text you do not input literally or displayed text that changes depending on circumstance. For example:

To connect to a remote machine using ssh, type **ssh *username@domain.name*** at a shell prompt. If the remote machine is `example.com` and your username on that machine is john, type **ssh john@example.com**.

The **mount -o remount *file-system*** command remounts the named file system. For example, to remount the `/home` file system, the command is **mount -o remount /home**.

To see the version of a currently installed package, use the **rpm -q *package*** command. It will return a result as follows: ***package-version-release***.

Note the words in bold italics above: *username*, *domain.name*, *file-system*, *package*, *version* and *release*. Each word is a placeholder, either for text you enter when issuing a command or for text displayed by the system.

Aside from standard usage for presenting the title of a work, italics denotes the first use of a new and important term. For example:

Publican is a *DocBook* publishing system.

Pull-quote Conventions

Terminal output and source code listings are set off visually from the surrounding text.

Output sent to a terminal is set in mono-spaced roman and presented thus:

```
books      Desktop  documentation  drafts  mss    photos  stuff  svn
books_tests Desktop1  downloads      images  notes  scripts svgs
```

Source-code listings are also set in mono-spaced roman but add syntax highlighting as follows:

```
package org.jboss.book.jca.ex1;

import javax.naming.InitialContext;

public class ExClient
{
    public static void main(String args[])
        throws Exception
    {
        InitialContext iniCtx = new InitialContext();
        Object          ref    = iniCtx.lookup("EchoBean");
        EchoHome         home  = (EchoHome) ref;
```

```
        Echo            echo    = home.create();

        System.out.println("Created Echo");

        System.out.println("Echo.echo('Hello') = " + echo.echo("Hello"));
    }
}
```

Notes and Warnings

Finally, we use three visual styles to draw attention to information that might otherwise be overlooked.

Note

Notes are tips, shortcuts or alternative approaches to the task at hand. Ignoring a note should have no negative consequences, but you might miss out on a trick that makes your life easier.

Important

Important boxes detail things that are easily missed: configuration changes that only apply to the current session, or services that need restarting before an update will apply. Ignoring a box labeled “Important” will not cause data loss but may cause irritation and frustration.

Warning

Warnings should not be ignored. Ignoring warnings will most likely cause data loss.

We Need Feedback!

You should over ride this by creating your own local Feedback.xml file.

Chapter 1. Read-Me-First

The Scope of this Document

This document is intended to be an exhaustive reference for configuring Pacemaker.

To achieve this, it focuses on the XML syntax used to configure the CIB. For those that are allergic to XML, multiple higher-level front-ends (both command-line and GUI) are available. These tools will not be covered at all in this document ¹.

Users may be interested in other parts of the Pacemaker documentation set [<https://www.clusterlabs.org/pacemaker/doc/>], such as *Clusters from Scratch*, a step-by-step guide to setting up an example cluster, and *Pacemaker Administration*, a guide to maintaining a cluster.

What Is *Pacemaker*?

Pacemaker is a high-availability *cluster resource manager* — software that runs on a set of hosts (a *cluster* of *nodes*) in order to preserve integrity and minimize downtime of desired services (*resources*). ² It is maintained by the ClusterLabs [<https://www.ClusterLabs.org/>] community.

Pacemaker's key features include:

- Detection of and recovery from node- and service-level failures
- Ability to ensure data integrity by fencing faulty nodes
- Support for one or more nodes per cluster
- Support for multiple resource interface standards (anything that can be scripted can be clustered)
- Support (but no requirement) for shared storage
- Support for practically any redundancy configuration (active/passive, N+1, etc.)
- Automatically replicated configuration that can be updated from any node
- Ability to specify cluster-wide relationships between services, such as ordering, colocation and anti-colocation
- Support for advanced service types, such as *clones* (services that need to be active on multiple nodes), *stateful resources* (clones that can run in one of two modes), and containerized services
- Unified, scriptable cluster management tools

Fencing

Fencing, also known as *STONITH* (an acronym for Shoot The Other Node In The Head), is the ability to ensure that it is not possible for a node to be running a service. This is accomplished via *fence devices* such as intelligent power switches that cut power to the target, or intelligent network switches that cut the target's access to the local network.

¹ I hope, however, that the concepts explained here make the functionality of these tools more easily understood.

² *Cluster* is sometimes used in other contexts to refer to hosts grouped together for other purposes, such as high-performance computing (HPC), but Pacemaker is not intended for those purposes.

Pacemaker represents fence devices as a special class of resource.

A cluster cannot safely recover from certain failure conditions, such as an unresponsive node, without fencing.

Cluster Architecture

At a high level, a cluster can be viewed as having these parts (which together are often referred to as the *cluster stack*):

- **Resources:** These are the reason for the cluster's being — the services that need to be kept highly available.
- **Resource agents:** These are scripts or operating system components that start, stop, and monitor resources, given a set of resource parameters. These provide a uniform interface between Pacemaker and the managed services.
- **Fence agents:** These are scripts that execute node fencing actions, given a target and fence device parameters.
- **Cluster membership layer:** This component provides reliable messaging, membership, and quorum information about the cluster. Currently, Pacemaker supports Corosync [<http://www.corosync.org/>] as this layer.
- **Cluster resource manager:** Pacemaker provides the brain that processes and reacts to events that occur in the cluster. These events may include nodes joining or leaving the cluster; resource events caused by failures, maintenance, or scheduled activities; and other administrative actions. To achieve the desired availability, Pacemaker may start and stop resources and fence nodes.
- **Cluster tools:** These provide an interface for users to interact with the cluster. Various command-line and graphical (GUI) interfaces are available.

Most managed services are not, themselves, cluster-aware. However, many popular open-source cluster filesystems make use of a common *Distributed Lock Manager* (DLM), which makes direct use of Corosync for its messaging and membership capabilities and Pacemaker for the ability to fence nodes.

Figure 1.1. Example Cluster Stack

Pacemaker Architecture

Pacemaker itself is composed of multiple daemons that work together:

- pacemakerd
- pacemaker-attrd
- pacemaker-based
- pacemaker-controld
- pacemaker-execd
- pacemaker-fenced
- pacemaker-schedulerd

Figure 1.2. Internal Components

The Pacemaker master process (`pacemakerd`) spawns all the other daemons, and respawns them if they unexpectedly exit.

The *Cluster Information Base* (CIB) is an XML [<https://en.wikipedia.org/wiki/XML>] representation of the cluster's configuration and the state of all nodes and resources. The *CIB manager* (`pacemaker-based`) keeps the CIB synchronized across the cluster, and handles requests to modify it.

The attribute manager (`pacemaker-attrd`) maintains a database of attributes for all nodes, keeps it synchronized across the cluster, and handles requests to modify them. These attributes are usually recorded in the CIB.

Given a snapshot of the CIB as input, the *scheduler* (`pacemaker-schedulerd`) determines what actions are necessary to achieve the desired state of the cluster.

The *local executor* (`pacemaker-execd`) handles requests to execute resource agents on the local cluster node, and returns the result.

The *fencer* (`pacemaker-fenced`) handles requests to fence nodes. Given a target node, the fencer decides which cluster node(s) should execute which fencing device(s), and calls the necessary fencing agents (either directly, or via requests to the fencer peers on other nodes), and returns the result.

The *controller* (`pacemaker-controld`) is Pacemaker's coordinator, maintaining a consistent view of the cluster membership and orchestrating all the other components.

Pacemaker centralizes cluster decision-making by electing one of the controller instances as the *Designated Controller* (DC). Should the elected DC process (or the node it is on) fail, a new one is quickly established. The DC responds to cluster events by taking a current snapshot of the CIB, feeding it to the scheduler, then asking the executors (either directly on the local node, or via requests to controller peers on other nodes) and the fencer to execute any necessary actions.

Old daemon names

The Pacemaker daemons were renamed in version 2.0. You may still find references to the old names, especially in documentation targeted to version 1.1.

Old name	New name
<code>attrd</code>	<code>pacemaker-attrd</code>
<code>cib</code>	<code>pacemaker-based</code>
<code>crmd</code>	<code>pacemaker-controld</code>
<code>lrmd</code>	<code>pacemaker-execd</code>
<code>stonithd</code>	<code>pacemaker-fenced</code>
<code>pacemaker_remoted</code>	<code>pacemaker-remoted</code>

Node Redundancy Designs

Pacemaker supports practically any node redundancy configuration [https://en.wikipedia.org/wiki/High-availability_cluster#Node_configurations] including *Active/Active*, *Active/Passive*, *N+1*, *N+M*, *N-to-1* and *N-to-N*.

Active/passive clusters with two (or more) nodes using Pacemaker and DRBD [https://en.wikipedia.org/wiki/Distributed_Replicated_Block_Device:] are a cost-effective high-availability solution for many situations. One of the nodes provides the desired services, and if it fails, the other node takes over.

Figure 1.3. Active/Passive Redundancy

Pacemaker also supports multiple nodes in a shared-failover design, reducing hardware costs by allowing several active/passive clusters to be combined and share a common backup node.

Figure 1.4. Shared Failover

When shared storage is available, every node can potentially be used for failover. Pacemaker can even run multiple copies of services to spread out the workload.

Figure 1.5. N to N Redundancy

Chapter 2. Cluster-Wide Configuration

Configuration Layout

The cluster is defined by the Cluster Information Base (CIB), which uses XML notation. The simplest CIB, an empty one, looks like this:

Example 2.1. An empty configuration

```
<cib crm_feature_set="3.0.7" validate-with="pacemaker-1.2" admin_epoch="1" epoch="1">
  <configuration>
    <crm_config/>
    <nodes/>
    <resources/>
    <constraints/>
  </configuration>
  <status/>
</cib>
```

The empty configuration above contains the major sections that make up a CIB:

- **cib**: The entire CIB is enclosed with a `cib` tag. Certain fundamental settings are defined as attributes of this tag.
- **configuration**: This section—the primary focus of this document—contains traditional configuration information such as what resources the cluster serves and the relationships among them.
 - **crm_config**: cluster-wide configuration options
 - **nodes**: the machines that host the cluster
 - **resources**: the services run by the cluster
 - **constraints**: indications of how resources should be placed
- **status**: This section contains the history of each resource on each node. Based on this data, the cluster can construct the complete current state of the cluster. The authoritative source for this section is the local executor (`pacemaker-execd` process) on each cluster node, and the cluster will occasionally repopulate the entire section. For this reason, it is never written to disk, and administrators are advised against modifying it in any way.

In this document, configuration settings will be described as *properties* or *options* based on how they are defined in the CIB:

- Properties are XML attributes of an XML element.
- Options are name-value pairs expressed as `nvpair` child elements of an XML element.

Normally, you will use command-line tools that abstract the XML, so the distinction will be unimportant; both properties and options are cluster settings you can tweak.

CIB Properties

Certain settings are defined by CIB properties (that is, attributes of the `cib` tag) rather than with the rest of the cluster configuration in the `configuration` section.

The reason is simply a matter of parsing. These options are used by the configuration database which is, by design, mostly ignorant of the content it holds. So the decision was made to place them in an easy-to-find location.

Table 2.1. CIB Properties

Field	Description
admin_epoch	When a node joins the cluster, the cluster performs a check to see which node has the best configuration. It asks the node with the highest (admin_epoch, epoch, num_updates) tuple to replace the configuration on all the nodes — which makes setting them, and setting them correctly, very important. admin_epoch is never modified by the cluster; you can use this to make the configurations on any inactive nodes obsolete. <i>Never set this value to zero.</i> In such cases, the cluster cannot tell the difference between your configuration and the "empty" one used when nothing is found on disk.
epoch	The cluster increments this every time the configuration is updated (usually by the administrator).
num_updates	The cluster increments this every time the configuration or status is updated (usually by the cluster) and resets it to 0 when epoch changes.
validate-with	Determines the type of XML validation that will be done on the configuration. If set to none, the cluster will not verify that updates conform to the DTD (nor reject ones that don't). This option can be useful when operating a mixed-version cluster during an upgrade.
cib-last-written	Indicates when the configuration was last written to disk. Maintained by the cluster; for informational purposes only.
have-quorum	Indicates if the cluster has quorum. If false, this may mean that the cluster cannot start resources or fence other nodes (see no-quorum-policy below). Maintained by the cluster.
dc-uuid	Indicates which cluster node is the current leader. Used by the cluster when placing resources and determining the order of some events. Maintained by the cluster.

Cluster Options

Cluster options, as you might expect, control how the cluster behaves when confronted with certain situations.

They are grouped into sets within the `crm_config` section, and, in advanced configurations, there may be more than one set. (This will be described later in the section on Chapter 8, Rules where we will show how to have the cluster use different sets of options during working hours than during weekends.) For now, we will describe the simple case where each option is present at most once.

You can obtain an up-to-date list of cluster options, including their default values, by running the `man pacemaker-schedulerd` and `man pacemaker-controld` commands.

Table 2.2. Cluster Options

Option	Default	Description
cluster-name		An (optional) name for the cluster as a whole. This is mostly for users' convenience for use as desired in administration, but this can be used in the Pacemaker configuration in rules (as the #cluster-name node attribute). It may also be used by higher-level tools when displaying cluster information, and by certain resource agents (for example, the <code>ocf:heartbeat:GFS2</code> agent stores the cluster name in filesystem meta-data).
dc-version		Version of Pacemaker on the cluster's DC. Determined automatically by the cluster. Often includes the hash which identifies the exact Git changeset it was built from. Used for diagnostic purposes.
cluster-infrastructure		The messaging stack on which Pacemaker is currently running. Determined automatically by the cluster. Used for informational and diagnostic purposes.
no-quorum-policy	stop	What to do when the cluster does not have quorum. Allowed values: <ul style="list-style-type: none"> • <code>ignore</code>: continue all resource management • <code>freeze</code>: continue resource management, but don't recover resources from nodes not in the affected partition • <code>stop</code>: stop all resources in the affected cluster partition • <code>suicide</code>: fence all nodes in the affected cluster partition
batch-limit	0	The maximum number of actions that the cluster may execute in parallel across all nodes. The "correct" value will depend on the speed and load of your network and cluster nodes. If zero, the cluster will impose a dynamically calculated limit only when any node has high load.
migration-limit	-1	The number of live migration actions that the cluster is allowed to execute in parallel on a node. A value of -1 means unlimited.
symmetric-cluster	TRUE	Can all resources run on any node by default?
stop-all-resources	FALSE	Should the cluster stop all resources?
stop-orphan-resources	TRUE	Should deleted resources be stopped? This value takes precedence over <code>is-managed</code> (i.e. even unmanaged resources will be stopped if deleted from the configuration when this value is TRUE).
stop-orphan-actions	TRUE	Should deleted actions be cancelled?
start-failure-is-fatal	TRUE	Should a failure to start a resource on a particular node prevent further start attempts on that node? If FALSE,

Option	Default	Description
		the cluster will decide whether the same node is still eligible based on the resource's current failure count and <code>migration-threshold</code> (see the section called "Handling Resource Failure").
<code>enable-startup-probes</code>	TRUE	Should the cluster check for active resources during startup?
<code>maintenance-mode</code>	FALSE	Should the cluster refrain from monitoring, starting and stopping resources?
<code>stonith-enabled</code>	TRUE	<p>Should failed nodes and nodes with resources that can't be stopped be shot? If you value your data, set up a STONITH device and enable this.</p> <p>If true, or unset, the cluster will refuse to start resources unless one or more STONITH resources have been configured. If false, unresponsive nodes are immediately assumed to be running no resources, and resource takeover to online nodes starts without any further protection (which means <i>data loss</i> if the unresponsive node still accesses shared storage, for example). See also the <code>requires</code> meta-attribute in the section called "Resource Options".</p>
<code>stonith-action</code>	reboot	Action to send to STONITH device. Allowed values are <code>reboot</code> and <code>off</code> . The value <code>poweroff</code> is also allowed, but is only used for legacy devices.
<code>stonith-timeout</code>	60s	How long to wait for STONITH actions (reboot, on, off) to complete
<code>stonith-max-attempts</code>	10	How many times fencing can fail for a target before the cluster will no longer immediately re-attempt it.
<code>stonith-watchdog-timeout</code>	0	If nonzero, rely on hardware watchdog self-fencing. If positive, assume unseen nodes self-fence within this much time. If negative, and the <code>SBD_WATCHDOG_TIMEOUT</code> environment variable is set, use twice that value.
<code>concurrent-fencing</code>	FALSE	Is the cluster allowed to initiate multiple fence actions concurrently?
<code>fence-reaction</code>	stop	How should a cluster node react if notified of its own fencing? A cluster node may receive notification of its own fencing if fencing is misconfigured, or if fabric fencing is in use that doesn't cut cluster communication. Allowed values are <code>stop</code> to attempt to immediately stop pacemaker and stay stopped, or <code>panic</code> to attempt to immediately reboot the local node, falling back to <code>stop</code> on failure. The default is likely to be changed to <code>panic</code> in a future release. (<i>since 2.0.3</i>)
<code>priority-fencing-delay</code>		Enforce specified delay for the fencings that are targeting the lost nodes with the highest total resource priority in case we don't have the majority of the nodes in our cluster partition, so that the more significant nodes potentially win any fencing match, which is especially meaningful under split-brain of 2-node cluster. A promoted resource instance takes the base priority + 1 on calculation if the base priority

Option	Default	Description
		is not 0. If all the nodes have equal priority, then any <code>pcmk_delay_base/max</code> configured for the corresponding fencing resources will be applied. Otherwise as long as it's set, even if to 0, it takes precedence over any configured <code>pcmk_delay_base/max</code> . By default, priority fencing delay is disabled. (<i>since 2.0.4</i>)
<code>cluster-delay</code>	60s	Estimated maximum round-trip delay over the network (excluding action execution). If the DC requires an action to be executed on another node, it will consider the action failed if it does not get a response from the other node in this time (after considering the action's own timeout). The "correct" value will depend on the speed and load of your network and cluster nodes.
<code>dc-deadtime</code>	20s	How long to wait for a response from other nodes during startup. The "correct" value will depend on the speed/load of your network and the type of switches used.
<code>cluster-ipc-limit</code>	500	The maximum IPC message backlog before one cluster daemon will disconnect another. This is of use in large clusters, for which a good value is the number of resources in the cluster multiplied by the number of nodes. The default of 500 is also the minimum. Raise this if you see "Evicting client" messages for cluster daemon PIDs in the logs.
<code>pe-error-series-max</code>	-1	The number of PE inputs resulting in ERRORS to save. Used when reporting problems. A value of -1 means unlimited (report all).
<code>pe-warn-series-max</code>	-1	The number of PE inputs resulting in WARNINGS to save. Used when reporting problems. A value of -1 means unlimited (report all).
<code>pe-input-series-max</code>	-1	The number of "normal" PE inputs to save. Used when reporting problems. A value of -1 means unlimited (report all).
<code>placement-strategy</code>	default	How the cluster should allocate resources to nodes (see Chapter 12, <i>Utilization and Placement Strategy</i> [112]). Allowed values are <code>default</code> , <code>utilization</code> , <code>balanced</code> , and <code>minimal</code> .
<code>node-health-strategy</code>	none	How the cluster should react to node health attributes (see the section called "Tracking Node Health"). Allowed values are <code>none</code> , <code>migrate-on-red</code> , <code>only-green</code> , <code>progressive</code> , and <code>custom</code> .
<code>enable-acl</code>	FALSE	Whether access control lists (ACLs) (see Chapter 13, ACLs) can be used to authorize modifications to the CIB.
<code>node-health-base</code>	0	The base health score assigned to a node. Only used when <code>node-health-strategy</code> is <code>progressive</code> .

Option	Default	Description
node-health-green	0	The score to use for a node health attribute whose value is green. Only used when node-health-strategy is progressive or custom.
node-health-yellow	0	The score to use for a node health attribute whose value is yellow. Only used when node-health-strategy is progressive or custom.
node-health-red	0	The score to use for a node health attribute whose value is red. Only used when node-health-strategy is progressive or custom.
cluster-recheck-interval	15min	Pacemaker is primarily event-driven, and looks ahead to know when to recheck the cluster for failure timeouts and most time-based rules. However, it will also recheck the cluster after this amount of inactivity. This has two goals: rules with <code>date_spec</code> are only guaranteed to be checked this often, and it also serves as a fail-safe for certain classes of scheduler bugs. A value of 0 disables this polling; positive values are a time interval.
shutdown-lock	false	The default of false allows active resources to be recovered elsewhere when their node is cleanly shut down, which is what the vast majority of users will want. However, some users prefer to make resources highly available only for failures, with no recovery for clean shutdowns. If this option is true, resources active on a node when it is cleanly shut down are kept "locked" to that node (not allowed to run elsewhere) until they start again on that node after it rejoins (or for at most shutdown-lock-limit, if set). Stonith resources and Pacemaker Remote connections are never locked. Clone and bundle instances and the master role of promotable clones are currently never locked, though support could be added in a future release. Locks may be manually cleared using the <code>--refresh</code> option of <code>crm_resource</code> (both the resource and node must be specified; this works with remote nodes if their connection resource's target-role is set to Stopped, but not if Pacemaker Remote is stopped on the remote node without disabling the connection resource). (<i>since 2.0.4</i>)
shutdown-lock-limit	0	If shutdown-lock is true, and this is set to a nonzero time duration, locked resources will be allowed to start after this much time has passed since the node shutdown was initiated, even if the node has not rejoined. (This works with remote nodes only if their connection resource's target-role is set to Stopped.) (<i>since 2.0.4</i>)
remove-after-stop	FALSE	<i>Advanced Use Only:</i> Should the cluster remove resources from the LRM after they are stopped? Values other than the default are, at best, poorly tested and potentially dangerous.
startup-fencing	TRUE	<i>Advanced Use Only:</i> Should the cluster shoot unseen nodes? Not using the default is very unsafe!

Option	Default	Description
election-timeout	2min	<i>Advanced Use Only:</i> If you need to adjust this value, it probably indicates the presence of a bug.
shutdown-escalation	20min	<i>Advanced Use Only:</i> If you need to adjust this value, it probably indicates the presence of a bug.
join-integration-timeout	3min	<i>Advanced Use Only:</i> If you need to adjust this value, it probably indicates the presence of a bug.
join-finalization-timeout	30min	<i>Advanced Use Only:</i> If you need to adjust this value, it probably indicates the presence of a bug.
transition-delay	0s	<i>Advanced Use Only:</i> Delay cluster recovery for the configured interval to allow for additional/related events to occur. Useful if your configuration is sensitive to the order in which ping updates arrive. Enabling this option will slow down cluster recovery under all conditions.

Chapter 3. Cluster Nodes

Defining a Cluster Node

Each node in the cluster will have an entry in the nodes section containing its UUID, uname, and type.

Example 3.1. Example Corosync cluster node entry

```
<node id="101" uname="pcmk-1"/>
```

In normal circumstances, the admin should let the cluster populate this information automatically from the communications and membership data.

Where Pacemaker Gets the Node Name

Traditionally, Pacemaker required nodes to be referred to by the value returned by `uname -n`. This can be problematic for services that require the `uname -n` to be a specific value (e.g. for a licence file).

This requirement has been relaxed for clusters using Corosync 2.0 or later. The name Pacemaker uses is:

1. The value stored in `corosync.conf` under **ring0_addr** in the **nodelist**, if it does not contain an IP address; otherwise
2. The value stored in `corosync.conf` under **name** in the **nodelist**; otherwise
3. The value of `uname -n`

Pacemaker provides the `crm_node -n` command which displays the name used by a running cluster.

If a Corosync **nodelist** is used, `crm_node --name-for-id number` is also available to display the name used by the node with the corosync **nodeid** of *number*, for example: `crm_node --name-for-id 2`.

Node Attributes

Pacemaker allows node-specific values to be specified using *node attributes*. A node attribute has a name, and may have a distinct value for each node.

While certain node attributes have specific meanings to the cluster, they are mainly intended to allow administrators and resource agents to track any information desired.

For example, an administrator might choose to define node attributes for how much RAM and disk space each node has, which OS each uses, or which server room rack each node is in.

Users can configure rules that use node attributes to affect where resources are placed.

Setting and querying node attributes

Node attributes can be set and queried using the `crm_attribute` and `attrd_updater` commands, so that the user does not have to deal with XML configuration directly.

Here is an example of what XML configuration would be generated if an administrator ran this command:

Example 3.2. Result of using `crm_attribute` to specify which kernel `pcmk-1` is running

```
# crm_attribute --type nodes --node pcmk-1 --name kernel --update $(uname -r)

<node id="1" uname="pcmk-1">
  <instance_attributes id="nodes-1-attributes">
    <nvpair id="nodes-1-kernel" name="kernel" value="3.10.0-862.14.4.el7.x86_64"/>
  </instance_attributes>
</node>
```

To read back the value that was just set:

```
# crm_attribute --type nodes --node pcmk-1 --name kernel --query
scope=nodes name=kernel value=3.10.0-862.14.4.el7.x86_64
```

By specifying `--type nodes` the admin tells the cluster that this attribute is persistent across reboots. There are also transient attributes which are kept in the status section and are "forgotten" whenever the node leaves the cluster. Administrators can use this section by specifying `--type status`.

Special node attributes

Certain node attributes have special meaning to the cluster.

Node attribute names beginning with `#` are considered reserved for these special attributes. Some special attributes do not start with `#`, for historical reasons.

Certain special attributes are set automatically by the cluster, should never be modified directly, and can be used only within rules; these are listed under the section called "Node Attribute Expressions" [64].

For true/false values, the cluster considers a value of "1", "y", "yes", "on", or "true" (case-insensitively) to be true, "0", "n", "no", "off", "false", or unset to be false, and anything else to be an error.

Table 3.1. Node attributes with special significance

Name	Description
<code>fail-count-*</code>	Attributes whose names start with <code>fail-count-</code> are managed by the cluster to track how many times particular resource operations have failed on this node. These should be queried and cleared via the <code>crm_failcount</code> or <code>crm_resource --cleanup</code> commands rather than directly.
<code>last-failure-*</code>	Attributes whose names start with <code>last-failure-</code> are managed by the cluster to track when particular resource operations have most recently failed on this node. These should be cleared via the <code>crm_failcount</code> or <code>crm_resource --cleanup</code> commands rather than directly.
<code>maintenance</code>	Similar to the <code>maintenance-mode</code> cluster option, but for a single node. If true, resources will not be started or stopped on the node, resources and individual clone instances running on the node will become unmanaged, and any recurring operations for those will be cancelled.
<code>probe_complete</code>	This is managed by the cluster to detect when nodes need to be reprobbed, and should never be used directly.

Name	Description
resource-discovery-enabled	If the node is a remote node, fencing is enabled, and this attribute is explicitly set to false (unset means true in this case), resource discovery (probes) will not be done on this node. This is highly discouraged; the <code>resource-discovery</code> location constraint property is preferred for this purpose.
shutdown	This is managed by the cluster to orchestrate the shutdown of a node, and should never be used directly.
site-name	If set, this will be used as the value of the <code>#site-name</code> node attribute used in rules. (If not set, the value of the <code>cluster-name</code> cluster option will be used as <code>#site-name</code> instead.)
standby	If true, the node is in standby mode. This is typically set and queried via the <code>crm_standby</code> command rather than directly.
terminate	If the value is true or begins with any nonzero number, the node will be fenced. This is typically set by tools rather than directly.
#digests-*	Attributes whose names start with <code>#digests-</code> are managed by the cluster to detect when unfencing needs to be redone, and should never be used directly.
#node-unfenced	When the node was last unfenced (as seconds since the epoch). This is managed by the cluster and should never be used directly.

Warning

Restarting pacemaker on a node that is in single-node maintenance mode will likely lead to undesirable effects. If `maintenance` is set as a transient attribute, it will be erased when pacemaker is stopped, which will immediately take the node out of maintenance mode and likely get it fenced. Even if permanent, if pacemaker is restarted, any resources active on the node will have their local history erased when the node rejoins, so the cluster will no longer consider them running on the node and thus will consider them managed again, leading them to be started elsewhere. This behavior might be improved in a future release.

Chapter 4. Cluster Resources

What is a Cluster Resource?

A resource is a service made highly available by a cluster. The simplest type of resource, a *primitive* resource, is described in this section. More complex forms, such as groups and clones, are described in later sections.

Every primitive resource has a *resource agent*. A resource agent is an external program that abstracts the service it provides and present a consistent view to the cluster.

This allows the cluster to be agnostic about the resources it manages. The cluster doesn't need to understand how the resource works because it relies on the resource agent to do the right thing when given a `start`, `stop` or `monitor` command. For this reason, it is crucial that resource agents are well-tested.

Typically, resource agents come in the form of shell scripts. However, they can be written using any technology (such as C, Python or Perl) that the author is comfortable with.

Resource Classes

Pacemaker supports several classes of agents:

- OCF
- LSB
- Upstart
- Systemd
- Service
- Fencing
- Nagios Plugins

Open Cluster Framework

The OCF standard ¹ is basically an extension of the Linux Standard Base conventions for init scripts to:

- support parameters,
- make them self-describing, and
- make them extensible

OCF specs have strict definitions of the exit codes that actions must return. ²

¹ See <https://github.com/ClusterLabs/OCF-spec/tree/master/ra> . The Pacemaker implementation has been somewhat extended from the OCF specs.

² The resource-agents source code includes the `ocf-tester` script, which can be useful in this regard.

The cluster follows these specifications exactly, and giving the wrong exit code will cause the cluster to behave in ways you will likely find puzzling and annoying. In particular, the cluster needs to distinguish a completely stopped resource from one which is in some erroneous and indeterminate state.

Parameters are passed to the resource agent as environment variables, with the special prefix `OCF_RESKEY_`. So, a parameter which the user thinks of as `ip` will be passed to the resource agent as `OCF_RESKEY_ip`. The number and purpose of the parameters is left to the resource agent; however, the resource agent should use the `meta-data` command to advertise any that it supports.

The OCF class is the most preferred as it is an industry standard, highly flexible (allowing parameters to be passed to agents in a non-positional manner) and self-describing.

For more information, see the reference [http://www.linux-ha.org/wiki/OCF_Resource_Agents] and the *Resource Agents* section of *Pacemaker Administration*.

Linux Standard Base

LSB resource agents are more commonly known as *init scripts*. If a full path is not given, they are assumed to be located in `/etc/init.d`.

Commonly, they are provided by the OS distribution. In order to be used with a Pacemaker cluster, they must conform to the LSB specification.³

Warning

Many distributions or particular software packages claim LSB compliance but ship with broken init scripts. For details on how to check whether your init script is LSB-compatible, see the *Resource Agents* section of *Pacemaker Administration*. Common problematic violations of the LSB standard include:

- Not implementing the `status` operation at all
- Not observing the correct exit status codes for `start/stop/status` actions
- Starting a started resource returns an error
- Stopping a stopped resource returns an error

Important

Remember to make sure the computer is *not* configured to start any services at boot time — that should be controlled by the cluster.

Systemd

Some newer distributions have replaced the old "SysV" [<http://en.wikipedia.org/wiki/Init#SysV-style>] style of initialization daemons and scripts with an alternative called Systemd [<http://www.freedesktop.org/wiki/Software/systemd>].

Pacemaker is able to manage these services *if they are present*.

³ See http://refspecs.linux-foundation.org/LSB_3.0.0/LSB-Core-generic/LSB-Core-generic/inisrptact.html for the LSB Spec as it relates to init scripts.

Instead of init scripts, systemd has *unit files*. Generally, the services (unit files) are provided by the OS distribution, but there are online guides for converting from init scripts.⁴

Important

Remember to make sure the computer is *not* configured to start any services at boot time — that should be controlled by the cluster.

Upstart

Some newer distributions have replaced the old "SysV" [<http://en.wikipedia.org/wiki/Init#SysV-style>] style of initialization daemons (and scripts) with an alternative called Upstart [<http://upstart.ubuntu.com/>].

Pacemaker is able to manage these services *if they are present*.

Instead of init scripts, upstart has *jobs*. Generally, the services (jobs) are provided by the OS distribution.

Important

Remember to make sure the computer is *not* configured to start any services at boot time — that should be controlled by the cluster.

System Services

Since there are various types of system services (systemd, upstart, and lsb), Pacemaker supports a special *service* alias which intelligently figures out which one applies to a given cluster node.

This is particularly useful when the cluster contains a mix of systemd, upstart, and lsb.

In order, Pacemaker will try to find the named service as:

1. an LSB init script
2. a Systemd unit file
3. an Upstart job

STONITH

The STONITH class is used exclusively for fencing-related resources. This is discussed later in Chapter 6, Fencing.

Nagios Plugins

Nagios Plugins⁵ allow us to monitor services on remote hosts.

⁴ For example, <http://0pointer.de/blog/projects/systemd-for-admins-3.html>

⁵ The project has two independent forks, hosted at <https://www.nagios-plugins.org/> and <https://www.monitoring-plugins.org/>. Output from both projects' plugins is similar, so plugins from either project can be used with pacemaker.

Pacemaker is able to do remote monitoring with the plugins *if they are present*.

A common use case is to configure them as resources belonging to a resource container (usually a virtual machine), and the container will be restarted if any of them has failed. Another use is to configure them as ordinary resources to be used for monitoring hosts or services via the network.

The supported parameters are same as the long options of the plugin.

Resource Properties

These values tell the cluster which resource agent to use for the resource, where to find that resource agent and what standards it conforms to.

Table 4.1. Properties of a Primitive Resource

Field	Description
id	Your name for the resource
class	The standard the resource agent conforms to. Allowed values: lsb, nagios, ocf, service, stonith, systemd, upstart
type	The name of the Resource Agent you wish to use. E.g. IPaddr or Filesystem
provider	The OCF spec allows multiple vendors to supply the same resource agent. To use the OCF resource agents supplied by the Heartbeat project, you would specify heartbeat here.

The XML definition of a resource can be queried with the `crm_resource` tool. For example:

```
# crm_resource --resource Email --query-xml
```

might produce:

Example 4.1. A system resource definition

```
<primitive id="Email" class="service" type="exim"/>
```

Note

One of the main drawbacks to system services (LSB, systemd or Upstart) resources is that they do not allow any parameters!

Example 4.2. An OCF resource definition

```
<primitive id="Public-IP" class="ocf" type="IPaddr" provider="heartbeat">
  <instance_attributes id="Public-IP-params">
    <nvpair id="Public-IP-ip" name="ip" value="192.0.2.2"/>
  </instance_attributes>
</primitive>
```

Resource Options

Resources have two types of options: *meta-attributes* and *instance attributes*. Meta-attributes apply to any type of resource, while instance attributes are specific to each resource agent.

Resource Meta-Attributes

Meta-attributes are used by the cluster to decide how a resource should behave and can be easily set using the `--meta` option of the `crm_resource` command.

Table 4.2. Meta-attributes of a Primitive Resource

Field	Default	Description
priority	0	If not all resources can be active, the cluster will stop lower priority resources in order to keep higher priority ones active.
target-role	Started	What state should the cluster attempt to keep this resource in? Allowed values: <ul style="list-style-type: none"> Stopped: Force the resource to be stopped Started: Allow the resource to be started (and in the case of promotable clone resources, promoted to master if appropriate) Slave: Allow the resource to be started, but only in Slave mode if the resource is promotable Master: Equivalent to Started
is-managed	TRUE	Is the cluster allowed to start and stop the resource? Allowed values: true, false
maintenance	FALSE	Similar to the maintenance-mode cluster option, but for a single resource. If true, the resource will not be started, stopped, or monitored on any node. This differs from is-managed in that monitors will not be run. Allowed values: true, false
resource-stickiness	1 for individual clone instances, 0 for all other resources	A score that will be added to the current node when a resource is already active. This allows running resources to stay where they are, even if they would be placed elsewhere if they were being started from a stopped state.
requires	quorum for resources with a class of stonith, otherwise unfencing if unfencing is active in the cluster, otherwise fencing if stonith-enabled is true, otherwise quorum	Conditions under which the resource can be started Allowed values: <ul style="list-style-type: none"> nothing: can always be started quorum: The cluster can only start this resource if a majority of the configured nodes are active fencing: The cluster can only start this resource if a majority of the configured nodes are active <i>and</i> any failed or unknown nodes have been fenced unfencing: The cluster can only start this resource if a majority of the configured nodes are active <i>and</i> any failed or unknown nodes have been fenced <i>and</i> only on nodes that have been unfenced

Field	Default	Description
migration-threshold	INFINITY	How many failures may occur for this resource on a node, before this node is marked ineligible to host this resource. A value of 0 indicates that this feature is disabled (the node will never be marked ineligible); by contrast, the cluster treats INFINITY (the default) as a very large but finite number. This option has an effect only if the failed operation specifies <code>on-fail</code> as <code>restart</code> (the default), and additionally for failed start operations, if the cluster property <code>start-failure-is-fatal</code> is <code>false</code> .
failure-timeout	0	How many seconds to wait before acting as if the failure had not occurred, and potentially allowing the resource back to the node on which it failed. A value of 0 indicates that this feature is disabled.
multiple-active	stop_start	What should the cluster do if it ever finds the resource active on more than one node? Allowed values: <ul style="list-style-type: none"> <code>block</code>: mark the resource as unmanaged <code>stop_only</code>: stop all active instances and leave them that way <code>stop_start</code>: stop all active instances and start the resource in one location only
allow-migrate	TRUE for <code>ocf:pacemaker:remote-resources</code> , FALSE otherwise	Whether the cluster should try to "live migrate" this resource when it needs to be moved (see the section called "Migrating Resources")
container-attribute-target		Specific to bundle resources; see the section called "Bundle Node Attributes"
remote-node		The name of the Pacemaker Remote guest node this resource is associated with, if any. If specified, this both enables the resource as a guest node and defines the unique name used to identify the guest node. The guest must be configured to run the Pacemaker Remote daemon when it is started. WARNING: This value cannot overlap with any resource or node IDs.
remote-port	3121	If <code>remote-node</code> is specified, the port on the guest used for its Pacemaker Remote connection. The Pacemaker Remote daemon on the guest must be configured to listen on this port.
remote-addr	value of <code>remote-node</code>	If <code>remote-node</code> is specified, the IP address or hostname used to connect to the guest via Pacemaker Remote. The Pacemaker Remote daemon on the

Field	Default	Description
		guest must be configured to accept connections on this address.
remote-connect-timeout	60s	If remote-node is specified, how long before a pending guest connection will time out.

As an example of setting resource options, if you performed the following commands on an LSB Email resource:

```
# crm_resource --meta --resource Email --set-parameter priority --parameter-value 100
# crm_resource -m -r Email -p multiple-active -v block
```

the resulting resource definition might be:

Example 4.3. An LSB resource with cluster options

```
<primitive id="Email" class="lsb" type="exim">
  <meta_attributes id="Email-meta_attributes">
    <nvpair id="Email-meta_attributes-priority" name="priority" value="100"/>
    <nvpair id="Email-meta_attributes-multiple-active" name="multiple-active" value="block"/>
  </meta_attributes>
</primitive>
```

In addition to the cluster-defined meta-attributes described above, you may also configure arbitrary meta-attributes of your own choosing. Most commonly, this would be done for use in rules. For example, an IT department might define a custom meta-attribute to indicate which company department each resource is intended for. To reduce the chance of name collisions with cluster-defined meta-attributes added in the future, it is recommended to use a unique, organization-specific prefix for such attributes.

Setting Global Defaults for Resource Meta-Attributes

To set a default value for a resource option, add it to the `rsc_defaults` section with `crm_attribute`. For example,

```
# crm_attribute --type rsc_defaults --name is-managed --update false
```

would prevent the cluster from starting or stopping any of the resources in the configuration (unless of course the individual resources were specifically enabled by having their `is-managed` set to `true`).

Resource Instance Attributes

The resource agents of some resource classes (`lsb`, `systemd` and `upstart` *not* among them) can be given parameters which determine how they behave and which instance of a service they control.

If your resource agent supports parameters, you can add them with the `crm_resource` command. For example,

```
# crm_resource --resource Public-IP --set-parameter ip --parameter-value 192.0.2.2
```

would create an entry in the resource like this:

Example 4.4. An example OCF resource with instance attributes

```
<primitive id="Public-IP" class="ocf" type="IPaddr" provider="heartbeat">
```

```
<instance_attributes id="params-public-ip">
  <nvpair id="public-ip-addr" name="ip" value="192.0.2.2"/>
</instance_attributes>
</primitive>
```

For an OCF resource, the result would be an environment variable called `OCF_RESKEY_ip` with a value of `192.0.2.2`.

The list of instance attributes supported by an OCF resource agent can be found by calling the resource agent with the `meta-data` command. The output contains an XML description of all the supported attributes, their purpose and default values.

Example 4.5. Displaying the metadata for the Dummy resource agent template

```
# export OCF_ROOT=/usr/lib/ocf
# $OCF_ROOT/resource.d/pacemaker/Dummy meta-data
```

```
<?xml version="1.0"?>
<!DOCTYPE resource-agent SYSTEM "ra-api-1.dtd">
<resource-agent name="Dummy" version="1.0">
<version>1.0</version>
```

```
<longdesc>
```

This is a Dummy Resource Agent. It does absolutely nothing except keep track of whether its running or not.

Its purpose in life is for testing and to serve as a template for RA writers.

NB: Please pay attention to the timeouts specified in the actions section below. They should be meaningful for the kind of resource the agent manages. They should be the minimum advised timeouts, but they shouldn't/cannot cover all possible resource instances. So, try to be neither overly generous nor too stingy, but moderate. The minimum timeouts should never be below 10 seconds.

```
</longdesc>
```

```
<shortdesc>Example stateless resource agent</shortdesc>
```

```
<parameters>
```

```
<parameter name="state" unique="1">
```

```
<longdesc>
```

Location to store the resource state in.

```
</longdesc>
```

```
<shortdesc>State file</shortdesc>
```

```
<content type="string" default="/var/run/Dummy-default.state" />
```

```
</parameter>
```

```
<parameter name="fake" unique="0">
```

```
<longdesc>
```

Fake attribute that can be changed to cause a reload

```
</longdesc>
```

```
<shortdesc>Fake attribute that can be changed to cause a reload</shortdesc>
```

```
<content type="string" default="dummy" />
```

```
</parameter>
```

```
<parameter name="op_sleep" unique="1">
```

```
<longdesc>
Number of seconds to sleep during operations. This can be used to test how
the cluster reacts to operation timeouts.
</longdesc>
<shortdesc>Operation sleep duration in seconds.</shortdesc>
<content type="string" default="0" />
</parameter>

</parameters>

<actions>
<action name="start"          timeout="20" />
<action name="stop"          timeout="20" />
<action name="monitor"       timeout="20" interval="10" depth="0"/>
<action name="reload"        timeout="20" />
<action name="migrate_to"     timeout="20" />
<action name="migrate_from"   timeout="20" />
<action name="validate-all"   timeout="20" />
<action name="meta-data"      timeout="5" />
</actions>
</resource-agent>
```

Resource Operations

Operations are actions the cluster can perform on a resource by calling the resource agent. Resource agents must support certain common operations such as start, stop, and monitor, and may implement any others.

Operations may be explicitly configured for two purposes: to override defaults for options (such as timeout) that the cluster will use whenever it initiates the operation, and to run an operation on a recurring basis (for example, to monitor the resource for failure).

Example 4.6. An OCF resource with a non-default start timeout

```
<primitive id="Public-IP" class="ocf" type="IPaddr" provider="heartbeat">
  <operations>
    <op id="Public-IP-start" name="start" timeout="60s"/>
  </operations>
  <instance_attributes id="params-public-ip">
    <nvpair id="public-ip-addr" name="ip" value="192.0.2.2"/>
  </instance_attributes>
</primitive>
```

Pacemaker identifies operations by a combination of name and interval, so this combination must be unique for each resource. That is, you should not configure two operations for the same resource with the same name and interval.

Operation Properties

Operation properties may be specified directly in the `op` element as XML attributes, or in a separate `meta_attributes` block as `nvpair` elements. XML attributes take precedence over `nvpair` elements if both are specified.

Table 4.3. Properties of an Operation

Field	Default	Description
id		A unique name for the operation.
name		The action to perform. This can be any action supported by the agent; common values include <code>monitor</code> , <code>start</code> , and <code>stop</code> .
interval	0	How frequently (in seconds) to perform the operation. A value of 0 means "when needed". A positive value defines a <i>recurring action</i> , which is typically used with <code>monitor</code> .
timeout		How long to wait before declaring the action has failed
on-fail	restart (<i>except for stop operations, which default to fence when STONITH is enabled and block otherwise</i>)	<p>The action to take if this action ever fails. Allowed values:</p> <ul style="list-style-type: none"> <code>ignore</code>: Pretend the resource did not fail. <code>block</code>: Don't perform any further operations on the resource. <code>stop</code>: Stop the resource and do not start it elsewhere. <code>restart</code>: Stop the resource and start it again (possibly on a different node). <code>fence</code>: STONITH the node on which the resource failed. <code>standby</code>: Move <i>all</i> resources away from the node on which the resource failed.
enabled	TRUE	If false, ignore this operation definition. This is typically used to pause a particular recurring <code>monitor</code> operation; for instance, it can complement the respective resource being unmanaged (<code>is-managed=false</code>), as this alone will not block any configured monitoring. Disabling the operation does not suppress all actions of the given type. Allowed values: <code>true</code> , <code>false</code> .
record-pending	TRUE	If <code>true</code> , the intention to perform the operation is recorded so that GUIs and CLI tools can indicate that an operation is in progress. This is best set as an <i>operation default</i> (see the section called "Setting Global Defaults for Operations"). Allowed values: <code>true</code> , <code>false</code> .
role		Run the operation only on node(s) that the cluster thinks should be in the specified role. This only makes sense for recurring <code>monitor</code> operations. Allowed (case-sensitive) values: <code>Stopped</code> ,

Field	Default	Description
		Started, and in the case of promotable clone resources, Slave and Master.

Monitoring Resources for Failure

When Pacemaker first starts a resource, it runs one-time monitor operations (referred to as *probes*) to ensure the resource is running where it's supposed to be, and not running where it's not supposed to be. (This behavior can be affected by the `resource-discovery` location constraint property.)

Other than those initial probes, Pacemaker will *not* (by default) check that the resource continues to stay healthy.⁶ You must configure monitor operations explicitly to perform these checks.

Example 4.7. An OCF resource with a recurring health check

```
<primitive id="Public-IP" class="ocf" type="IPaddr" provider="heartbeat">
  <operations>
    <op id="Public-IP-start" name="start" timeout="60s"/>
    <op id="Public-IP-monitor" name="monitor" interval="60s"/>
  </operations>
  <instance_attributes id="params-public-ip">
    <nvpair id="public-ip-addr" name="ip" value="192.0.2.2"/>
  </instance_attributes>
</primitive>
```

By default, a monitor operation will ensure that the resource is running where it is supposed to. The `target-role` property can be used for further checking.

For example, if a resource has one monitor operation with `interval=10` `role=Started` and a second monitor operation with `interval=11` `role=Stopped`, the cluster will run the first monitor on any nodes it thinks *should* be running the resource, and the second monitor on any nodes that it thinks *should not* be running the resource (for the truly paranoid, who want to know when an administrator manually starts a service by mistake).

Note

Currently, monitors with `role=Stopped` are not implemented for clone resources.

Monitoring Resources When Administration is Disabled

Recurring monitor operations behave differently under various administrative settings:

- When a resource is unmanaged (by setting `is-managed=false`): No monitors will be stopped.

If the unmanaged resource is stopped on a node where the cluster thinks it should be running, the cluster will detect and report that it is not, but it will not consider the monitor failed, and will not try to start the resource until it is managed again.

Starting the unmanaged resource on a different node is strongly discouraged and will at least cause the cluster to consider the resource failed, and may require the resource's `target-role` to be set to `Stopped` then `Started` to be recovered.

⁶ Currently, anyway. Automatic monitoring operations may be added in a future version of Pacemaker.

- When a node is put into standby: All resources will be moved away from the node, and all monitor operations will be stopped on the node, except those specifying role as Stopped (which will be newly initiated if appropriate).
- When the cluster is put into maintenance mode: All resources will be marked as unmanaged. All monitor operations will be stopped, except those specifying role as Stopped (which will be newly initiated if appropriate). As with single unmanaged resources, starting a resource on a node other than where the cluster expects it to be will cause problems.

Setting Global Defaults for Operations

You can change the global default values for operation properties in a given cluster. These are defined in an `op_defaults` section of the CIB's configuration section, and can be set with `crm_attribute`. For example,

```
# crm_attribute --type op_defaults --name timeout --update 20s
```

would default each operation's `timeout` to 20 seconds. If an operation's definition also includes a value for `timeout`, then that value would be used for that operation instead.

When Implicit Operations Take a Long Time

The cluster will always perform a number of implicit operations: `start`, `stop` and a non-recurring monitor operation used at startup to check whether the resource is already active. If one of these is taking too long, then you can create an entry for them and specify a longer timeout.

Example 4.8. An OCF resource with custom timeouts for its implicit actions

```
<primitive id="Public-IP" class="ocf" type="IPaddr" provider="heartbeat">
  <operations>
    <op id="public-ip-startup" name="monitor" interval="0" timeout="90s"/>
    <op id="public-ip-start" name="start" interval="0" timeout="180s"/>
    <op id="public-ip-stop" name="stop" interval="0" timeout="15min"/>
  </operations>
  <instance_attributes id="params-public-ip">
    <nvpair id="public-ip-addr" name="ip" value="192.0.2.2"/>
  </instance_attributes>
</primitive>
```

Multiple Monitor Operations

Provided no two operations (for a single resource) have the same name and interval, you can have as many monitor operations as you like. In this way, you can do a superficial health check every minute and progressively more intense ones at higher intervals.

To tell the resource agent what kind of check to perform, you need to provide each monitor with a different value for a common parameter. The OCF standard creates a special parameter called `OCF_CHECK_LEVEL` for this purpose and dictates that it is "made available to the resource agent without the normal `OCF_RESKEY` prefix".

Whatever name you choose, you can specify it by adding an `instance_attributes` block to the `op` tag. It is up to each resource agent to look for the parameter and decide how to use it.

Example 4.9. An OCF resource with two recurring health checks, performing different levels of checks specified via OCF_CHECK_LEVEL.

```
<primitive id="Public-IP" class="ocf" type="IPaddr" provider="heartbeat">
  <operations>
    <op id="public-ip-health-60" name="monitor" interval="60">
      <instance_attributes id="params-public-ip-depth-60">
        <nvpair id="public-ip-depth-60" name="OCF_CHECK_LEVEL" value="10"/>
      </instance_attributes>
    </op>
    <op id="public-ip-health-300" name="monitor" interval="300">
      <instance_attributes id="params-public-ip-depth-300">
        <nvpair id="public-ip-depth-300" name="OCF_CHECK_LEVEL" value="20"/>
      </instance_attributes>
    </op>
  </operations>
  <instance_attributes id="params-public-ip">
    <nvpair id="public-ip-level" name="ip" value="192.0.2.2"/>
  </instance_attributes>
</primitive>
```

Disabling a Monitor Operation

The easiest way to stop a recurring monitor is to just delete it. However, there can be times when you only want to disable it temporarily. In such cases, simply add `enabled=false` to the operation's definition.

Example 4.10. Example of an OCF resource with a disabled health check

```
<primitive id="Public-IP" class="ocf" type="IPaddr" provider="heartbeat">
  <operations>
    <op id="public-ip-check" name="monitor" interval="60s" enabled="false"/>
  </operations>
  <instance_attributes id="params-public-ip">
    <nvpair id="public-ip-addr" name="ip" value="192.0.2.2"/>
  </instance_attributes>
</primitive>
```

This can be achieved from the command line by executing:

```
# cibadmin --modify --xml-text '<op id="public-ip-check" enabled="false"/>'
```

Once you've done whatever you needed to do, you can then re-enable it with

```
# cibadmin --modify --xml-text '<op id="public-ip-check" enabled="true"/>'
```

Chapter 5. Resource Constraints

Scores

Scores of all kinds are integral to how the cluster works. Practically everything from moving a resource to deciding which resource to stop in a degraded cluster is achieved by manipulating scores in some way.

Scores are calculated per resource and node. Any node with a negative score for a resource can't run that resource. The cluster places a resource on the node with the highest score for it.

Infinity Math

Pacemaker implements `INFINITY` (or equivalently, `+INFINITY`) internally as a score of 1,000,000. Addition and subtraction with it follow these three basic rules:

- Any value + `INFINITY` = `INFINITY`
- Any value - `INFINITY` = `-INFINITY`
- `INFINITY` - `INFINITY` = `-INFINITY`

Note

What if you want to use a score higher than 1,000,000? Typically this possibility arises when someone wants to base the score on some external metric that might go above 1,000,000.

The short answer is you can't.

The long answer is it is sometimes possible work around this limitation creatively. You may be able to set the score to some computed value based on the external metric rather than use the metric directly. For nodes, you can store the metric as a node attribute, and query the attribute when computing the score (possibly as part of a custom resource agent).

Deciding Which Nodes a Resource Can Run On

Location constraints tell the cluster which nodes a resource can run on.

There are two alternative strategies. One way is to say that, by default, resources can run anywhere, and then the location constraints specify nodes that are not allowed (an *opt-out* cluster). The other way is to start with nothing able to run anywhere, and use location constraints to selectively enable allowed nodes (an *opt-in* cluster).

Whether you should choose opt-in or opt-out depends on your personal preference and the make-up of your cluster. If most of your resources can run on most of the nodes, then an opt-out arrangement is likely to result in a simpler configuration. On the other-hand, if most resources can only run on a small subset of nodes, an opt-in configuration might be simpler.

Location Properties

Table 5.1. Attributes of a `rsc_location` Element

Attribute	Default	Description
<code>id</code>		A unique name for the constraint (required)
<code>rsc</code>		The name of the resource to which this constraint applies. A location constraint must either have a <code>rsc</code> , have a <code>rsc-pattern</code> , or contain at least one resource set.
<code>rsc-pattern</code>		A pattern matching the names of resources to which this constraint applies. The syntax is the same as POSIX [http://pubs.opengroup.org/onlinepubs/9699919799/basedefs/V1_chap09.html#tag_09_04] extended regular expressions, with the addition of an initial <code>!</code> indicating that resources <i>not</i> matching the pattern are selected. If the regular expression contains submatches, and the constraint is governed by a rule, the submatches can be referenced as <code>%0</code> through <code>%9</code> in the rule's <code>score-attribute</code> or a rule expression's <code>attribute</code> . A location constraint must either have a <code>rsc</code> , have a <code>rsc-pattern</code> , or contain at least one resource set.
<code>node</code>		The name of the node to which this constraint applies. A location constraint must either have a <code>node</code> and <code>score</code> , or contain at least one rule.
<code>score</code>		Positive values indicate a preference for running the affected resource(s) on node — the higher the value, the stronger the preference. Negative values indicate the resource(s) should avoid this node (a value of <code>-INFINITY</code> changes "should" to "must"). A location constraint must either have a <code>node</code> and <code>score</code> , or contain at least one rule.
<code>resource-discovery</code>	<code>always</code>	<p>Whether Pacemaker should perform resource discovery (that is, check whether the resource is already running) for this resource on this node. This should normally be left as the default, so that rogue instances of a service can be stopped when they are running where they are not supposed to be. However, there are two situations where disabling resource discovery is a good idea: when a service is not installed on a node, discovery might return an error (properly written OCF agents will not, so this is usually only seen with other agent types); and when Pacemaker Remote is used to scale a cluster to hundreds of nodes, limiting resource discovery to allowed nodes can significantly boost performance.</p> <ul style="list-style-type: none"> • <code>always</code>: Always perform resource discovery for the specified resource on this node. • <code>never</code>: Never perform resource discovery for the specified resource on this node. This option should generally be used with a <code>-INFINITY</code> score, although that is not strictly required. • <code>exclusive</code>: Perform resource discovery for the specified resource only on this node (and other nodes similarly marked as <code>exclusive</code>). Multiple location constraints using <code>exclusive</code> discovery for the same

Attribute	Default	Description
		resource across different nodes creates a subset of nodes resource-discovery is exclusive to. If a resource is marked for <i>exclusive</i> discovery on one or more nodes, that resource is only allowed to be placed within that subset of nodes.

Warning

Setting resource-discovery to *never* or *exclusive* removes Pacemaker's ability to detect and stop unwanted instances of a service running where it's not supposed to be. It is up to the system administrator (you!) to make sure that the service can *never* be active on nodes without resource-discovery (such as by leaving the relevant software uninstalled).

Asymmetrical "Opt-In" Clusters

To create an opt-in cluster, start by preventing resources from running anywhere by default:

```
# crm_attribute --name symmetric-cluster --update false
```

Then start enabling nodes. The following fragment says that the web server prefers **sles-1**, the database prefers **sles-2** and both can fail over to **sles-3** if their most preferred node fails.

Example 5.1. Opt-in location constraints for two resources

```
<constraints>
  <rsc_location id="loc-1" rsc="Webserver" node="sles-1" score="200"/>
  <rsc_location id="loc-2" rsc="Webserver" node="sles-3" score="0"/>
  <rsc_location id="loc-3" rsc="Database" node="sles-2" score="200"/>
  <rsc_location id="loc-4" rsc="Database" node="sles-3" score="0"/>
</constraints>
```

Symmetrical "Opt-Out" Clusters

To create an opt-out cluster, start by allowing resources to run anywhere by default:

```
# crm_attribute --name symmetric-cluster --update true
```

Then start disabling nodes. The following fragment is the equivalent of the above opt-in configuration.

Example 5.2. Opt-out location constraints for two resources

```
<constraints>
  <rsc_location id="loc-1" rsc="Webserver" node="sles-1" score="200"/>
  <rsc_location id="loc-2-do-not-run" rsc="Webserver" node="sles-2" score="-INFI
  <rsc_location id="loc-3-do-not-run" rsc="Database" node="sles-1" score="-INFIN
  <rsc_location id="loc-4" rsc="Database" node="sles-2" score="200"/>
</constraints>
```

What if Two Nodes Have the Same Score

If two nodes have the same score, then the cluster will choose one. This choice may seem random and may not be what was intended, however the cluster was not given enough information to know any better.

Example 5.3. Constraints where a resource prefers two nodes equally

```
<constraints>
  <rsc_location id="loc-1" rsc="Webserver" node="sles-1" score="INFINITY"/>
  <rsc_location id="loc-2" rsc="Webserver" node="sles-2" score="INFINITY"/>
  <rsc_location id="loc-3" rsc="Database" node="sles-1" score="500"/>
  <rsc_location id="loc-4" rsc="Database" node="sles-2" score="300"/>
  <rsc_location id="loc-5" rsc="Database" node="sles-2" score="200"/>
</constraints>
```

In the example above, assuming no other constraints and an inactive cluster, Webserver would probably be placed on sles-1 and Database on sles-2. It would likely have placed Webserver based on the node's uname and Database based on the desire to spread the resource load evenly across the cluster. However other factors can also be involved in more complex configurations.

Specifying the Order in which Resources Should Start/Stop

Ordering constraints tell the cluster the order in which certain resource actions should occur.

Important

Ordering constraints affect *only* the ordering of resource actions; they do *not* require that the resources be placed on the same node. If you want resources to be started on the same node *and* in a specific order, you need both an ordering constraint *and* a colocation constraint (see the section called “Placing Resources Relative to other Resources”), or alternatively, a group (see the section called “Groups - A Syntactic Shortcut”).

Ordering Properties

Table 5.2. Attributes of a rsc_order Element

Field	Default	Description
id		A unique name for the constraint
first		Name of the resource that the then resource depends on
then		Name of the dependent resource
first-action	start	The action that the first resource must complete before then-action can be initiated for the then resource. Allowed values: start, stop, promote, demote.
then-action	value of first-action	The action that the then resource can execute only after the first-action on the first resource has completed. Allowed values: start, stop, promote, demote.

Field	Default	Description
kind	Mandatory	<p>How to enforce the constraint. Allowed values:</p> <ul style="list-style-type: none"> Mandatory: then-action will never be initiated for the then resource unless and until first-action successfully completes for the first resource. Optional: The constraint applies only if both specified resource actions are scheduled in the same transition (that is, in response to the same cluster state). This means that then-action is allowed on the then resource regardless of the state of the first resource, but if both actions happen to be scheduled at the same time, they will be ordered. Serialize: Ensure that the specified actions are never performed concurrently for the specified resources. First-action and then-action can be executed in either order, but one must complete before the other can be initiated. An example use case is when resource start-up puts a high load on the host.
symmetrical	TRUE for Mandatory and Optional kinds. FALSE for Serialize kind.	If true, the reverse of the constraint applies for the opposite action (for example, if B starts after A starts, then B stops before A stops). Serialize orders cannot be symmetrical.

Promote and demote apply to the master role of promotable resources.

Optional and mandatory ordering

Here is an example of ordering constraints where Database *must* start before Webserver, and IP *should* start before Webserver if they both need to be started:

Example 5.4. Optional and mandatory ordering constraints

```
<constraints>
<rsc_order id="order-1" first="IP" then="Webserver" kind="Optional"/>
<rsc_order id="order-2" first="Database" then="Webserver" kind="Mandatory" />
</constraints>
```

Because the above example lets symmetrical default to TRUE, Webserver must be stopped before Database can be stopped, and Webserver should be stopped before IP if they both need to be stopped.

Placing Resources Relative to other Resources

Colocation constraints tell the cluster that the location of one resource depends on the location of another one.

Colocation has an important side-effect: it affects the order in which resources are assigned to a node. Think about it: You can't place A relative to B unless you know where B is.¹

So when you are creating colocation constraints, it is important to consider whether you should colocate A with B, or B with A.

Another thing to keep in mind is that, assuming A is colocated with B, the cluster will take into account A's preferences when deciding which node to choose for B.

For a detailed look at exactly how this occurs, see Colocation Explained [http://clusterlabs.org/doc/Colocation_Explained.pdf].

Important

Colocation constraints affect *only* the placement of resources; they do *not* require that the resources be started in a particular order. If you want resources to be started on the same node *and* in a specific order, you need both an ordering constraint (see the section called “Specifying the Order in which Resources Should Start/Stop”) *and* a colocation constraint, or alternatively, a group (see the section called “Groups - A Syntactic Shortcut”).

Colocation Properties

Table 5.3. Attributes of a `rsc_colocation` Constraint

Field	Default	Description
<code>id</code>		A unique name for the constraint (required).
<code>rsc</code>		The name of a resource that should be located relative to <code>with-rsc</code> (required).
<code>with-rsc</code>		The name of the resource used as the colocation target. The cluster will decide where to put this resource first and then decide where to put <code>rsc</code> (required).
<code>node-attribute</code>	<code>#uname</code>	The node attribute that must be the same on the node running <code>rsc</code> and the node running <code>with-rsc</code> for the constraint to be satisfied. (For details, see the section called “Colocation by Node Attribute”.)
<code>score</code>		Positive values indicate the resources should run on the same node. Negative values indicate the resources should run on different nodes. Values of <code>+/- +INFINITY</code> change “should” to “must”.

Mandatory Placement

Mandatory placement occurs when the constraint's score is `+INFINITY` or `-INFINITY`. In such cases, if the constraint can't be satisfied, then the `rsc` resource is not permitted to run. For `score=INFINITY`, this includes cases where the `with-rsc` resource is not active.

If you need resource A to always run on the same machine as resource B, you would add the following constraint:

¹ While the human brain is sophisticated enough to read the constraint in any order and choose the correct one depending on the situation, the cluster is not quite so smart. Yet.

Example 5.5. Mandatory colocation constraint for two resources

```
<rsc_colocation id="colocate" rsc="A" with-rsc="B" score="INFINITY"/>
```

Remember, because `INFINITY` was used, if B can't run on any of the cluster nodes (for whatever reason) then A will not be allowed to run. Whether A is running or not has no effect on B.

Alternatively, you may want the opposite — that A *cannot* run on the same machine as B. In this case, use `score="-INFINITY"`.

Example 5.6. Mandatory anti-colocation constraint for two resources

```
<rsc_colocation id="anti-colocate" rsc="A" with-rsc="B" score="-INFINITY"/>
```

Again, by specifying `-INFINITY`, the constraint is binding. So if the only place left to run is where B already is, then A may not run anywhere.

As with `INFINITY`, B can run even if A is stopped. However, in this case A also can run if B is stopped, because it still meets the constraint of A and B not running on the same node.

Advisory Placement

If mandatory placement is about "must" and "must not", then advisory placement is the "I'd prefer if" alternative. For constraints with scores greater than `-INFINITY` and less than `INFINITY`, the cluster will try to accommodate your wishes but may ignore them if the alternative is to stop some of the cluster resources.

As in life, where if enough people prefer something it effectively becomes mandatory, advisory colocation constraints can combine with other elements of the configuration to behave as if they were mandatory.

Example 5.7. Advisory colocation constraint for two resources

```
<rsc_colocation id="colocate-maybe" rsc="A" with-rsc="B" score="500"/>
```

Colocation by Node Attribute

The `node-attribute` property of a colocation constraints allows you to express the requirement, "these resources must be on similar nodes".

As an example, imagine that you have two Storage Area Networks (SANs) that are not controlled by the cluster, and each node is connected to one or the other. You may have two resources `r1` and `r2` such that `r2` needs to use the same SAN as `r1`, but doesn't necessarily have to be on the same exact node. In such a case, you could define a node attribute named `san`, with the value `san1` or `san2` on each node as appropriate. Then, you could colocate `r2` with `r1` using `node-attribute` set to `san`.

Resource Sets

Resource sets allow multiple resources to be affected by a single constraint.

Example 5.8. A set of 3 resources

```
<resource_set id="resource-set-example">  
  <resource_ref id="A"/>  
  <resource_ref id="B"/>
```

```
<resource_ref id="C"/>
</resource_set>
```

Resource sets are valid inside `rsc_location`, `rsc_order` (see the section called “Ordering Sets of Resources”), `rsc_colocation` (see the section called “Colocating Sets of Resources”), and `rsc_ticket` (see the section called “Configuring Ticket Dependencies”) constraints.

A resource set has a number of properties that can be set, though not all have an effect in all contexts.

Table 5.4. Attributes of a `resource_set` Element

Field	Default	Description
<code>id</code>		A unique name for the set
<code>sequential</code>	true	Whether the members of the set must be acted on in order. Meaningful within <code>rsc_order</code> and <code>rsc_colocation</code> .
<code>require-all</code>	true	Whether all members of the set must be active before continuing. With the current implementation, the cluster may continue even if only one member of the set is started, but if more than one member of the set is starting at the same time, the cluster will still wait until all of those have started before continuing (this may change in future versions). Meaningful within <code>rsc_order</code> .
<code>role</code>		Limit the effect of the constraint to the specified role. Meaningful within <code>rsc_location</code> , <code>rsc_colocation</code> and <code>rsc_ticket</code> .
<code>action</code>		Limit the effect of the constraint to the specified action. Meaningful within <code>rsc_order</code> .
<code>score</code>		<i>Advanced use only.</i> Use a specific score for this set within the constraint.

Ordering Sets of Resources

A common situation is for an administrator to create a chain of ordered resources, such as:

Example 5.9. A chain of ordered resources

```
<constraints>
  <rsc_order id="order-1" first="A" then="B" />
  <rsc_order id="order-2" first="B" then="C" />
  <rsc_order id="order-3" first="C" then="D" />
</constraints>
```

Figure 5.1. Visual representation of the four resources' start order for the above constraints

Ordered Set

To simplify this situation, resource sets (see the section called “Resource Sets”) can be used within ordering constraints:

Example 5.10. A chain of ordered resources expressed as a set

```
<constraints>
  <rsc_order id="order-1">
    <resource_set id="ordered-set-example" sequential="true">
      <resource_ref id="A"/>
      <resource_ref id="B"/>
      <resource_ref id="C"/>
      <resource_ref id="D"/>
    </resource_set>
  </rsc_order>
</constraints>
```

While the set-based format is not less verbose, it is significantly easier to get right and maintain.

Important

If you use a higher-level tool, pay attention to how it exposes this functionality. Depending on the tool, creating a set $A \rightarrow B$ may be equivalent to $A \rightarrow B$, or $B \rightarrow A$.

Ordering Multiple Sets

The syntax can be expanded to allow sets of resources to be ordered relative to each other, where the members of each individual set may be ordered or unordered (controlled by the `sequential` property). In the example below, A and B can both start in parallel, as can C and D, however C and D can only start once *both* A *and* B are active.

Example 5.11. Ordered sets of unordered resources

```
<constraints>
  <rsc_order id="order-1">
    <resource_set id="ordered-set-1" sequential="false">
      <resource_ref id="A"/>
      <resource_ref id="B"/>
    </resource_set>
    <resource_set id="ordered-set-2" sequential="false">
      <resource_ref id="C"/>
      <resource_ref id="D"/>
    </resource_set>
  </rsc_order>
</constraints>
```

Figure 5.2. Visual representation of the start order for two ordered sets of unordered resources

Of course either set—or both sets—of resources can also be internally ordered (by setting `sequential="true"`) and there is no limit to the number of sets that can be specified.

Example 5.12. Advanced use of set ordering - Three ordered sets, two of which are internally unordered

```
<constraints>
```

```
<rsc_order id="order-1">
  <resource_set id="ordered-set-1" sequential="false">
    <resource_ref id="A"/>
    <resource_ref id="B"/>
  </resource_set>
  <resource_set id="ordered-set-2" sequential="true">
    <resource_ref id="C"/>
    <resource_ref id="D"/>
  </resource_set>
  <resource_set id="ordered-set-3" sequential="false">
    <resource_ref id="E"/>
    <resource_ref id="F"/>
  </resource_set>
</rsc_order>
</constraints>
```

Figure 5.3. Visual representation of the start order for the three sets defined above

Important

An ordered set with `sequential=false` makes sense only if there is another set in the constraint. Otherwise, the constraint has no effect.

Resource Set OR Logic

The unordered set logic discussed so far has all been "AND" logic. To illustrate this take the 3 resource set figure in the previous section. Those sets can be expressed, (A and B) then (C) then (D) then (E and F).

Say for example we want to change the first set, (A and B), to use "OR" logic so the sets look like this: (A or B) then (C) then (D) then (E and F). This functionality can be achieved through the use of the `require-all` option. This option defaults to TRUE which is why the "AND" logic is used by default. Setting `require-all=false` means only one resource in the set needs to be started before continuing on to the next set.

Example 5.13. Resource Set "OR" logic: Three ordered sets, where the first set is internally unordered with "OR" logic

```
<constraints>
  <rsc_order id="order-1">
    <resource_set id="ordered-set-1" sequential="false" require-all="false">
      <resource_ref id="A"/>
      <resource_ref id="B"/>
    </resource_set>
    <resource_set id="ordered-set-2" sequential="true">
      <resource_ref id="C"/>
      <resource_ref id="D"/>
    </resource_set>
    <resource_set id="ordered-set-3" sequential="false">
      <resource_ref id="E"/>
      <resource_ref id="F"/>
    </resource_set>
  </rsc_order>
</constraints>
```

```
</rsc_order>
</constraints>
```

Important

An ordered set with `require-all=false` makes sense only in conjunction with `sequential=false`. Think of it like this: `sequential=false` modifies the set to be an unordered set using "AND" logic by default, and adding `require-all=false` flips the unordered set's "AND" logic to "OR" logic.

Colocating Sets of Resources

Another common situation is for an administrator to create a set of colocated resources.

The simplest way to do this is to define a resource group (see the section called “Groups - A Syntactic Shortcut”), but that cannot always accurately express the desired relationships. For example, maybe the resources do not need to be ordered.

Another way would be to define each relationship as an individual constraint, but that causes a difficult-to-follow constraint explosion as the number of resources and combinations grow.

Example 5.14. Colocation chain as individual constraints, where A is placed first, then B, then C, then D

```
<constraints>
  <rsc_colocation id="coloc-1" rsc="D" with-rsc="C" score="INFINITY"/>
  <rsc_colocation id="coloc-2" rsc="C" with-rsc="B" score="INFINITY"/>
  <rsc_colocation id="coloc-3" rsc="B" with-rsc="A" score="INFINITY"/>
</constraints>
```

To express complicated relationships with a simplified syntax², resource sets can be used within colocation constraints.

Example 5.15. Equivalent colocation chain expressed using `resource_set`

```
<constraints>
  <rsc_colocation id="coloc-1" score="INFINITY" >
    <resource_set id="colocated-set-example" sequential="true">
      <resource_ref id="A"/>
      <resource_ref id="B"/>
      <resource_ref id="C"/>
      <resource_ref id="D"/>
    </resource_set>
  </rsc_colocation>
</constraints>
```

Note

Within a `resource_set`, the resources are listed in the order they are *placed*, which is the reverse of the order in which they are *colocated*. In the above example, resource A is placed before resource B, which is the same as saying resource B is colocated with resource A.

² which is not the same as saying easy to follow

As with individual constraints, a resource that can't be active prevents any resource that must be colocated with it from being active. In both of the two previous examples, if B is unable to run, then both C and by inference D must remain stopped.

Important

If you use a higher-level tool, pay attention to how it exposes this functionality. Depending on the tool, creating a set A B may be equivalent to A with B, or B with A.

Resource sets can also be used to tell the cluster that entire *sets* of resources must be colocated relative to each other, while the individual members within any one set may or may not be colocated relative to each other (determined by the set's `sequential` property).

In the following example, resources B, C, and D will each be colocated with A (which will be placed first). A must be able to run in order for any of the resources to run, but any of B, C, or D may be stopped without affecting any of the others.

Example 5.16. Using colocated sets to specify a shared dependency

```
<constraints>
  <rsc_colocation id="coloc-1" score="INFINITY" >
    <resource_set id="colocated-set-2" sequential="false">
      <resource_ref id="B"/>
      <resource_ref id="C"/>
      <resource_ref id="D"/>
    </resource_set>
    <resource_set id="colocated-set-1" sequential="true">
      <resource_ref id="A"/>
    </resource_set>
  </rsc_colocation>
</constraints>
```

Note

Pay close attention to the order in which resources and sets are listed. While the members of any one sequential set are placed first to last (i.e., the colocation dependency is last with first), multiple sets are placed last to first (i.e. the colocation dependency is first with last).

Important

A colocated set with `sequential="false"` makes sense only if there is another set in the constraint. Otherwise, the constraint has no effect.

There is no inherent limit to the number and size of the sets used. The only thing that matters is that in order for any member of one set in the constraint to be active, all members of sets listed after it must also be active (and naturally on the same node); and if a set has `sequential="true"`, then in order for one member of that set to be active, all members listed before it must also be active.

If desired, you can restrict the dependency to instances of promotable clone resources that are in a specific role, using the set's `role` property.

Example 5.17. Colocation in which the members of the middle set have no interdependencies, and the last set listed applies only to instances in the master role

```
<constraints>
```

```
<rsc_colocation id="coloc-1" score="INFINITY" >
  <resource_set id="colocated-set-1" sequential="true">
    <resource_ref id="F"/>
    <resource_ref id="G"/>
  </resource_set>
  <resource_set id="colocated-set-2" sequential="false">
    <resource_ref id="C"/>
    <resource_ref id="D"/>
    <resource_ref id="E"/>
  </resource_set>
  <resource_set id="colocated-set-3" sequential="true" role="Master">
    <resource_ref id="A"/>
    <resource_ref id="B"/>
  </resource_set>
</rsc_colocation>
</constraints>
```

Figure 5.4. Visual representation of the above example (resources are placed from left to right)

Note

Unlike ordered sets, colocated sets do not use the `require-all` option.

Chapter 6. Fencing

What Is Fencing?

Fencing is the ability to make a node unable to run resources, even when that node is unresponsive to cluster commands.

Fencing is also known as *STONITH*, an acronym for "Shoot The Other Node In The Head", since the most common fencing method is cutting power to the node. Another method is "fabric fencing", cutting the node's access to some capability required to run resources (such as network access or a shared disk).

Why Is Fencing Necessary?

Fencing protects your data from being corrupted by malfunctioning nodes or unintentional concurrent access to shared resources.

Fencing protects against the "split brain" failure scenario, where cluster nodes have lost the ability to reliably communicate with each other but are still able to run resources. If the cluster just assumed that uncommunicative nodes were down, then multiple instances of a resource could be started on different nodes.

The effect of split brain depends on the resource type. For example, an IP address brought up on two hosts on a network will cause packets to randomly be sent to one or the other host, rendering the IP useless. For a database or clustered file system, the effect could be much more severe, causing data corruption or divergence.

Fencing also is used when a resource cannot otherwise be stopped. If a failed resource fails to stop, it cannot be recovered elsewhere. Fencing the resource's node is the only way to ensure the resource is recoverable.

Users may also configure the `on-fail` property of any resource operation to `fencing`, in which case the cluster will fence the resource's node if the operation fails.

Fence Devices

A *fence device* (or *fencing device*) is a special type of resource that provides the means to fence a node.

Examples of fencing devices include intelligent power switches and IPMI devices that accept SNMP commands to cut power to a node, and iSCSI controllers that allow SCSI reservations to be used to cut a node's access to a shared disk.

Since fencing devices will be used to recover from loss of networking connectivity to other nodes, it is essential that they do not rely on the same network as the cluster itself, otherwise that network becomes a single point of failure.

Since loss of a node due to power outage is indistinguishable from loss of network connectivity to that node, it is also essential that at least one fence device for a node does not share power with that node. For example, an on-board IPMI controller that shares power with its host should not be used as the sole fencing device for that host.

Since fencing is used to isolate malfunctioning nodes, no fence device should rely on its target functioning properly. This includes, for example, devices that `ssh` into a node and issue a shutdown command (such devices might be suitable for testing, but never for production).

Fence Agents

A *fence agent* (or *fencing agent*) is a `stonith`-class resource agent.

The fence agent standard provides commands (such as `off` and `reboot`) that the cluster can use to fence nodes. As with other resource agent classes, this allows a layer of abstraction so that Pacemaker doesn't need any knowledge about specific fencing technologies — that knowledge is isolated in the agent.

When a Fence Device Can Be Used

Fencing devices do not actually "run" like most services. Typically, they just provide an interface for sending commands to an external device.

Additionally, fencing may be initiated by Pacemaker, by other cluster-aware software such as DRBD or DLM, or manually by an administrator, at any point in the cluster life cycle, including before any resources have been started.

To accommodate this, Pacemaker does not require the fence device resource to be "started" in order to be used. Whether a fence device is started or not determines whether a node runs any recurring monitor for the device, and gives the node a slight preference for being chosen to execute fencing using that device.

By default, any node can execute any fencing device. If a fence device is disabled by setting its `target-role` to `Stopped`, then no node can use that device. If mandatory location constraints prevent a specific node from "running" a fence device, then that node will never be chosen to execute fencing using the device. A node may fence itself, but the cluster will choose that only if no other nodes can do the fencing.

A common configuration scenario is to have one fence device per target node. In such a case, users often configure anti-location constraints so that the target node does not monitor its own device. The best practice is to make the constraint optional (i.e. a finite negative score rather than `-INFINITY`), so that the node can fence itself if no other nodes can.

Limitations of Fencing Resources

Fencing resources have certain limitations that other resource classes don't:

- They may have only one set of meta-attributes and one set of instance attributes.
- If rules are used to determine fencing resource options, these may only be evaluated when first read, meaning that later changes to the rules will have no effect. Therefore, it is better to avoid confusion and not use rules at all with fencing resources.

These limitations could be revisited if there is significant user demand.

Special Options for Fencing Resources

The table below lists special instance attributes that may be set for any fencing resource (*not* meta-attributes, even though they are interpreted by pacemaker rather than the fence agent). These are also listed in the man page for `pacemaker-fenced`.

Table 6.1. Additional Properties of Fencing Resources

Field	Type	Default	Description
stonith-timeout	NA	NA	Older versions used this to override the default period to wait for a STONITH (reboot, on, off) action to complete for this device. It has been replaced by the <code>pcmk_reboot_timeout</code> and <code>pcmk_off_timeout</code> properties.
provides	string		Any special capability provided by the fence device. Currently, only one such capability is meaningful: <code>unfencing</code> (see the section called “Unfencing”).
pcmk_host_map	string		A mapping of host names to ports numbers for devices that do not support host names. Example: <code>node1:1;node2:2,3</code> tells the cluster to use port 1 for node1 and ports 2 and 3 for node2 . If <code>pcmk_host_check</code> is explicitly set to <code>static-list</code> , either this or <code>pcmk_host_list</code> must be set.
pcmk_host_list	string		A list of machines controlled by this device. If <code>pcmk_host_check</code> is explicitly set to <code>static-list</code> , either this or <code>pcmk_host_map</code> must be set.
pcmk_host_check	string	A value appropriate to other configuration options and device capabilities (see note below)	How to determine which machines are controlled by the device. Allowed values: <ul style="list-style-type: none"> • <code>dynamic-list</code>: query the device via the "list" command • <code>static-list</code>: check the <code>pcmk_host_list</code> or <code>pcmk_host_map</code> attribute • <code>status</code>: query the device via the "status" command • <code>none</code>: assume every device can fence every machine
pcmk_delay_max	time	0s	Enable a random delay of up to the time specified before executing fencing actions. This is sometimes used in two-node clusters to ensure that the nodes don't fence each other at the same time. The overall delay introduced by pacemaker is derived from this random

Field	Type	Default	Description
			delay value adding a static delay so that the sum is kept below the maximum delay.
pcmk_delay_base	time	0s	Enable a static delay before executing fencing actions. This can be used e.g. in two-node clusters to ensure that the nodes don't fence each other, by having separate fencing resources with different values. The node that is fenced with the shorter delay will lose a fencing race. The overall delay introduced by pacemaker is derived from this value plus a random delay such that the sum is kept below the maximum delay.
pcmk_action_limit	integer	1	The maximum number of actions that can be performed in parallel on this device, if the cluster option <code>concurrent-fencing</code> is <code>true</code> . -1 is unlimited.
pcmk_host_argument	string	port	<i>Advanced use only.</i> Which parameter should be supplied to the resource agent to identify the node to be fenced. Some devices do not support the standard <code>port</code> parameter or may provide additional ones. Use this to specify an alternate, device-specific parameter. A value of <code>none</code> tells the cluster not to supply any additional parameters.
pcmk_reboot_action	string	reboot	<i>Advanced use only.</i> The command to send to the resource agent in order to reboot a node. Some devices do not support the standard commands or may provide additional ones. Use this to specify an alternate, device-specific command.
pcmk_reboot_timeout	time	60s	<i>Advanced use only.</i> Specify an alternate timeout to use for <code>reboot</code> actions instead of the value of <code>stonith-timeout</code> . Some devices need much more or less time to complete than normal. Use this to specify an alternate, device-specific timeout.

Field	Type	Default	Description
<code>pcmk_reboot_retries</code>	integer	2	<i>Advanced use only.</i> The maximum number of times to retry the <code>reboot</code> command within the timeout period. Some devices do not support multiple connections, and operations may fail if the device is busy with another task, so Pacemaker will automatically retry the operation, if there is time remaining. Use this option to alter the number of times Pacemaker retries before giving up.
<code>pcmk_off_action</code>	string	off	<i>Advanced use only.</i> The command to send to the resource agent in order to shut down a node. Some devices do not support the standard commands or may provide additional ones. Use this to specify an alternate, device-specific command.
<code>pcmk_off_timeout</code>	time	60s	<i>Advanced use only.</i> Specify an alternate timeout to use for <code>off</code> actions instead of the value of <code>stonith-timeout</code> . Some devices need much more or less time to complete than normal. Use this to specify an alternate, device-specific timeout.
<code>pcmk_off_retries</code>	integer	2	<i>Advanced use only.</i> The maximum number of times to retry the <code>off</code> command within the timeout period. Some devices do not support multiple connections, and operations may fail if the device is busy with another task, so Pacemaker will automatically retry the operation, if there is time remaining. Use this option to alter the number of times Pacemaker retries before giving up.
<code>pcmk_list_action</code>	string	list	<i>Advanced use only.</i> The command to send to the resource agent in order to list nodes. Some devices do not support the standard commands or may provide additional ones. Use this to specify an alternate, device-specific command.
<code>pcmk_list_timeout</code>	time	60s	<i>Advanced use only.</i> Specify an alternate timeout to use for <code>list</code> actions instead of the value of <code>stonith-timeout</code> . Some devices need much more or less time to complete than normal. Use this to specify an alternate, device-specific timeout.

Field	Type	Default	Description
<code>pcmk_list_retries</code>	integer	2	<i>Advanced use only.</i> The maximum number of times to retry the <code>list</code> command within the timeout period. Some devices do not support multiple connections, and operations may fail if the device is busy with another task, so Pacemaker will automatically retry the operation, if there is time remaining. Use this option to alter the number of times Pacemaker retries before giving up.
<code>pcmk_monitor_action</code>	string	monitor	<i>Advanced use only.</i> The command to send to the resource agent in order to report extended status. Some devices do not support the standard commands or may provide additional ones. Use this to specify an alternate, device-specific command.
<code>pcmk_monitor_timeout</code>	time	60s	<i>Advanced use only.</i> Specify an alternate timeout to use for <code>monitor</code> actions instead of the value of <code>stonith-timeout</code> . Some devices need much more or less time to complete than normal. Use this to specify an alternate, device-specific timeout.
<code>pcmk_monitor_retries</code>	integer	2	<i>Advanced use only.</i> The maximum number of times to retry the <code>monitor</code> command within the timeout period. Some devices do not support multiple connections, and operations may fail if the device is busy with another task, so Pacemaker will automatically retry the operation, if there is time remaining. Use this option to alter the number of times Pacemaker retries before giving up.
<code>pcmk_status_action</code>	string	status	<i>Advanced use only.</i> The command to send to the resource agent in order to report status. Some devices do not support the standard commands or may provide additional ones. Use this to specify an alternate, device-specific command.
<code>pcmk_status_timeout</code>	time	60s	<i>Advanced use only.</i> Specify an alternate timeout to use for <code>status</code> actions instead of the value of <code>stonith-timeout</code> . Some devices need much more or less time to complete than normal. Use this to specify an alternate, device-specific timeout.

Field	Type	Default	Description
<code>pcmk_status_retries</code>	integer	2	<i>Advanced use only.</i> The maximum number of times to retry the <code>status</code> command within the timeout period. Some devices do not support multiple connections, and operations may fail if the device is busy with another task, so Pacemaker will automatically retry the operation, if there is time remaining. Use this option to alter the number of times Pacemaker retries before giving up.

Note

The default value for `pcmk_host_check` is `static-list` if either `pcmk_host_list` or `pcmk_host_map` is configured. If neither of those are configured, the default is `dynamic-list` if the fence device supports the `list` action, or `status` if the fence device supports the `status` action but not the `list` action. If none of those conditions apply, the default is `none`.

Unfencing

With fabric fencing (such as cutting network or shared disk access rather than power), it is expected that the cluster will fence the node, and then a system administrator must manually investigate what went wrong, correct any issues found, then reboot (or restart the cluster services on) the node.

Once the node reboots and rejoins the cluster, some fabric fencing devices require an explicit command to restore the node's access. This capability is called *unfencing* and is typically implemented as the fence agent's `on` command.

If any cluster resource has `requires` set to `unfencing`, then that resource will not be probed or started on a node until that node has been unfenced.

Fence Devices Dependent on Other Resources

In some cases, a fence device may require some other cluster resource (such as an IP address) to be active in order to function properly.

This is obviously undesirable in general: fencing may be required when the depended-on resource is not active, or fencing may be required because the node running the depended-on resource is no longer responding.

However, this may be acceptable under certain conditions:

- The dependent fence device should not be able to target any node that is allowed to run the depended-on resource.
- The depended-on resource should not be disabled during production operation.
- The `concurrent-fencing` cluster property should be set to `true`. Otherwise, if both the node running the depended-on resource and some node targeted by the dependent fence device need to be fenced, the fencing of the node running the depended-on resource might be ordered first, making the second fencing impossible and blocking further recovery. With concurrent fencing, the dependent fence

device might fail at first due to the depended-on resource being unavailable, but it will be retried and eventually succeed once the resource is brought back up.

Even under those conditions, there is one unlikely problem scenario. The DC always schedules fencing of itself after any other fencing needed, to avoid unnecessary repeated DC elections. If the dependent fence device targets the DC, and both the DC and a different node running the depended-on resource need to be fenced, the DC fencing will always fail and block further recovery. Note, however, that losing a DC node entirely causes some other node to become DC and schedule the fencing, so this is only a risk when a stop or other operation with `on-fail` set to `fencing` fails on the DC.

Configuring Fencing

1. Find the correct driver:

```
# stonith_admin --list-installed
```

2. Find the required parameters associated with the device (replacing `$AGENT_NAME` with the name obtained from the previous step):

```
# stonith_admin --metadata --agent $AGENT_NAME
```

3. Create a file called `stonith.xml` containing a primitive resource with a class of `stonith`, a type equal to the agent name obtained earlier, and a parameter for each of the values returned in the previous step.

4. If the device does not know how to fence nodes based on their `uname`, you may also need to set the special `pcmk_host_map` parameter. See `man pacemaker-fenced` for details.

5. If the device does not support the `list` command, you may also need to set the special `pcmk_host_list` and/or `pcmk_host_check` parameters. See `man pacemaker-fenced` for details.

6. If the device does not expect the victim to be specified with the `port` parameter, you may also need to set the special `pcmk_host_argument` parameter. See `man pacemaker-fenced` for details.

7. Upload it into the CIB using `cibadmin`:

```
# cibadmin -C -o resources --xml-file stonith.xml
```

8. Set `stonith-enabled` to true:

```
# crm_attribute -t crm_config -n stonith-enabled -v true
```

9. Once the `stonith` resource is running, you can test it by executing the following (although you might want to stop the cluster on that machine first):

```
# stonith_admin --reboot nodename
```

Example Fencing Configuration

Assume we have a chassis containing four nodes and an IPMI device active on 192.0.2.1. We would choose the `fence_ipmilan` driver, and obtain the following list of parameters:

Example 6.1. Obtaining a list of Fence Agent Parameters

```
# stonith_admin --metadata -a fence_ipmilan
```

```
<resource-agent name="fence_ipmilan" shortdesc="Fence agent for IPMI over LAN">
  <symlink name="fence_ilo3" shortdesc="Fence agent for HP iLO3"/>
  <symlink name="fence_ilo4" shortdesc="Fence agent for HP iLO4"/>
  <symlink name="fence_idrac" shortdesc="Fence agent for Dell iDRAC"/>
  <symlink name="fence_imm" shortdesc="Fence agent for IBM Integrated Management M
  <longdesc>
</longdesc>
<vendor-url>
</vendor-url>
<parameters>
  <parameter name="auth" unique="0" required="0">
    <getopt mixed="-A"/>
    <content type="string"/>
    <shortdesc>
  </shortdesc>
  </parameter>
  <parameter name="ipaddr" unique="0" required="1">
    <getopt mixed="-a"/>
    <content type="string"/>
    <shortdesc>
  </shortdesc>
  </parameter>
  <parameter name="passwd" unique="0" required="0">
    <getopt mixed="-p"/>
    <content type="string"/>
    <shortdesc>
  </shortdesc>
  </parameter>
  <parameter name="passwd_script" unique="0" required="0">
    <getopt mixed="-S"/>
    <content type="string"/>
    <shortdesc>
  </shortdesc>
  </parameter>
  <parameter name="lanplus" unique="0" required="0">
    <getopt mixed="-P"/>
    <content type="boolean"/>
    <shortdesc>
  </shortdesc>
  </parameter>
  <parameter name="login" unique="0" required="0">
    <getopt mixed="-l"/>
    <content type="string"/>
    <shortdesc>
  </shortdesc>
  </parameter>
  <parameter name="action" unique="0" required="0">
    <getopt mixed="-o"/>
    <content type="string" default="reboot"/>
    <shortdesc>
  </shortdesc>
  </parameter>
  <parameter name="timeout" unique="0" required="0">
    <getopt mixed="-t"/>
```

```
    <content type="string"/>
    <shortdesc>
  </shortdesc>
</parameter>
<parameter name="cipher" unique="0" required="0">
  <getopt mixed="-C"/>
  <content type="string"/>
  <shortdesc>
</shortdesc>
</parameter>
<parameter name="method" unique="0" required="0">
  <getopt mixed="-M"/>
  <content type="string" default="onoff"/>
  <shortdesc>
</shortdesc>
</parameter>
<parameter name="power_wait" unique="0" required="0">
  <getopt mixed="-T"/>
  <content type="string" default="2"/>
  <shortdesc>
</shortdesc>
</parameter>
<parameter name="delay" unique="0" required="0">
  <getopt mixed="-f"/>
  <content type="string"/>
  <shortdesc>
</shortdesc>
</parameter>
<parameter name="privlvl" unique="0" required="0">
  <getopt mixed="-L"/>
  <content type="string"/>
  <shortdesc>
</shortdesc>
</parameter>
<parameter name="verbose" unique="0" required="0">
  <getopt mixed="-v"/>
  <content type="boolean"/>
  <shortdesc>
</shortdesc>
</parameter>
</parameters>
<actions>
  <action name="on"/>
  <action name="off"/>
  <action name="reboot"/>
  <action name="status"/>
  <action name="diag"/>
  <action name="list"/>
  <action name="monitor"/>
  <action name="metadata"/>
  <action name="stop" timeout="20s"/>
  <action name="start" timeout="20s"/>
</actions>
</resource-agent>
```

Based on that, we would create a fencing resource fragment that might look like this:

Example 6.2. An IPMI-based Fencing Resource

```
<primitive id="Fencing" class="stonith" type="fence_ipmilan" >
  <instance_attributes id="Fencing-params" >
    <nvpair id="Fencing-passwd" name="passwd" value="testuser" />
    <nvpair id="Fencing-login" name="login" value="abc123" />
    <nvpair id="Fencing-ipaddr" name="ipaddr" value="192.0.2.1" />
    <nvpair id="Fencing-pcmk_host_list" name="pcm_k_host_list" value="pcm_k-1 pcm_k-2" />
  </instance_attributes>
  <operations>
    <op id="Fencing-monitor-10m" interval="10m" name="monitor" timeout="300s" />
  </operations>
</primitive>
```

Finally, we need to enable fencing:

```
# crm_attribute -t crm_config -n stonith-enabled -v true
```

Fencing Topologies

Pacemaker supports fencing nodes with multiple devices through a feature called *fencing topologies*. Fencing topologies may be used to provide alternative devices in case one fails, or to require multiple devices to all be executed successfully in order to consider the node successfully fenced, or even a combination of the two.

Create the individual devices as you normally would, then define one or more `fencing-level` entries in the `fencing-topology` section of the configuration.

- Each fencing level is attempted in order of ascending `index`. Allowed values are 1 through 9.
- If a device fails, processing terminates for the current level. No further devices in that level are exercised, and the next level is attempted instead.
- If the operation succeeds for all the listed devices in a level, the level is deemed to have passed.
- The operation is finished when a level has passed (success), or all levels have been attempted (failed).
- If the operation failed, the next step is determined by the scheduler and/or the controller.

Some possible uses of topologies include:

- Try on-board IPMI, then an intelligent power switch if that fails
- Try fabric fencing of both disk and network, then fall back to power fencing if either fails
- Wait up to a certain time for a kernel dump to complete, then cut power to the node

Table 6.2. Properties of Fencing Levels

Field	Description
id	A unique name for the level
target	The name of a single node to which this level applies

Field	Description
target-pattern	An extended regular expression (as defined in POSIX [http://pubs.opengroup.org/onlinepubs/9699919799/basedefs/V1_chap09.html#tag_09_04]) matching the names of nodes to which this level applies
target-attribute	The name of a node attribute that is set (to target-value) for nodes to which this level applies
target-value	The node attribute value (of target-attribute) that is set for nodes to which this level applies
index	The order in which to attempt the levels. Levels are attempted in ascending order <i>until one succeeds</i> . Valid values are 1 through 9.
devices	A comma-separated list of devices that must all be tried for this level

Example 6.3. Fencing topology with different devices for different nodes

```
<cib crm_feature_set="3.0.6" validate-with="pacemaker-1.2" admin_epoch="1" epoch=
  <configuration>
    ...
    <fencing-topology>
      <!-- For pcmk-1, try poison-pill and fail back to power -->
      <fencing-level id="f-p1.1" target="pcmk-1" index="1" devices="poison-pill"/>
      <fencing-level id="f-p1.2" target="pcmk-1" index="2" devices="power"/>

      <!-- For pcmk-2, try disk and network, and fail back to power -->
      <fencing-level id="f-p2.1" target="pcmk-2" index="1" devices="disk,network"/>
      <fencing-level id="f-p2.2" target="pcmk-2" index="2" devices="power"/>
    </fencing-topology>
    ...
  </configuration>
  <status/>
</cib>
```

Example Dual-Layer, Dual-Device Fencing Topologies

The following example illustrates an advanced use of fencing-topology in a cluster with the following properties:

- 3 nodes (2 active prod-mysql nodes, 1 prod_mysql-rep in standby for quorum purposes)
- the active nodes have an IPMI-controlled power board reached at 192.0.2.1 and 192.0.2.2
- the active nodes also have two independent PSUs (Power Supply Units) connected to two independent PDUs (Power Distribution Units) reached at 198.51.100.1 (port 10 and port 11) and 203.0.113.1 (port 10 and port 11)
- the first fencing method uses the fence_ipmi agent
- the second fencing method uses the fence_apc_snmp agent targetting 2 fencing devices (one per PSU, either port 10 or 11)
- fencing is only implemented for the active nodes and has location constraints
- fencing topology is set to try IPMI fencing first then default to a "sure-kill" dual PDU fencing


```

    <nvpair id="fence_prod-mysql1_apc1-instance_attributes-login" name="login" val
    <nvpair id="fence_prod-mysql1_apc1-instance_attributes-passwd" name="passwd" v
    <nvpair id="fence_prod-mysql1_apc1-instance_attributes-pcmk_host_list" name="p
  </instance_attributes>
</primitive>
<primitive class="stonith" id="fence_prod-mysql1_apc2" type="fence_apc_snmp">
  <instance_attributes id="fence_prod-mysql1_apc2-instance_attributes">
    <nvpair id="fence_prod-mysql1_apc2-instance_attributes-ipaddr" name="ipaddr" v
    <nvpair id="fence_prod-mysql1_apc2-instance_attributes-action" name="action" v
    <nvpair id="fence_prod-mysql1_apc2-instance_attributes-port" name="port" value
    <nvpair id="fence_prod-mysql1_apc2-instance_attributes-login" name="login" val
    <nvpair id="fence_prod-mysql1_apc2-instance_attributes-passwd" name="passwd" v
    <nvpair id="fence_prod-mysql1_apc2-instance_attributes-pcmk_host_list" name="p
  </instance_attributes>
</primitive>
<primitive class="stonith" id="fence_prod-mysql2_apc1" type="fence_apc_snmp">
  <instance_attributes id="fence_prod-mysql2_apc1-instance_attributes">
    <nvpair id="fence_prod-mysql2_apc1-instance_attributes-ipaddr" name="ipaddr" v
    <nvpair id="fence_prod-mysql2_apc1-instance_attributes-action" name="action" v
    <nvpair id="fence_prod-mysql2_apc1-instance_attributes-port" name="port" value
    <nvpair id="fence_prod-mysql2_apc1-instance_attributes-login" name="login" val
    <nvpair id="fence_prod-mysql2_apc1-instance_attributes-passwd" name="passwd" v
    <nvpair id="fence_prod-mysql2_apc1-instance_attributes-pcmk_host_list" name="p
  </instance_attributes>
</primitive>
<primitive class="stonith" id="fence_prod-mysql2_apc2" type="fence_apc_snmp">
  <instance_attributes id="fence_prod-mysql2_apc2-instance_attributes">
    <nvpair id="fence_prod-mysql2_apc2-instance_attributes-ipaddr" name="ipaddr" v
    <nvpair id="fence_prod-mysql2_apc2-instance_attributes-action" name="action" v
    <nvpair id="fence_prod-mysql2_apc2-instance_attributes-port" name="port" value
    <nvpair id="fence_prod-mysql2_apc2-instance_attributes-login" name="login" val
    <nvpair id="fence_prod-mysql2_apc2-instance_attributes-passwd" name="passwd" v
    <nvpair id="fence_prod-mysql2_apc2-instance_attributes-pcmk_host_list" name="p
  </instance_attributes>
</primitive>

```

Location Constraints. To prevent STONITH from trying to run a fencing agent on the same node it is supposed to fence, constraints are placed on all the fencing primitives:

```

<constraints>
  <rscl_location id="l_fence_prod-mysql1_ipmi" node="prod-mysql1" rsc="fence_prod-m
  <rscl_location id="l_fence_prod-mysql2_ipmi" node="prod-mysql2" rsc="fence_prod-m
  <rscl_location id="l_fence_prod-mysql1_apc2" node="prod-mysql1" rsc="fence_prod-m
  <rscl_location id="l_fence_prod-mysql1_apc1" node="prod-mysql1" rsc="fence_prod-m
  <rscl_location id="l_fence_prod-mysql2_apc1" node="prod-mysql2" rsc="fence_prod-m
  <rscl_location id="l_fence_prod-mysql2_apc2" node="prod-mysql2" rsc="fence_prod-m
</constraints>

```

Fencing topology. Now that all the fencing resources are defined, it's time to create the right topology. We want to first fence using IPMI and if that does not work, fence both PDUs to effectively and surely kill the node.

```

<fencing-topology>
  <fencing-level devices="fence_prod-mysql1_ipmi" id="fencing-2" index="1" target=
  <fencing-level devices="fence_prod-mysql1_apc1,fence_prod-mysql1_apc2" id="fenci

```

```

    <fencing-level devices="fence_prod-mysql2_ipmi" id="fencing-0" index="1" target=
    <fencing-level devices="fence_prod-mysql2_apc1,fence_prod-mysql2_apc2" id="fenci
</fencing-topology>

```

Please note, in fencing-topology, the lowest index value determines the priority of the first fencing method.

Final configuration. Put together, the configuration looks like this:

```

<cib admin_epoch="0" crm_feature_set="3.0.7" epoch="292" have-quorum="1" num_updat
<configuration>
  <crm_config>
    <cluster_property_set id="cib-bootstrap-options">
      <nvpair id="cib-bootstrap-options-stonith-enabled" name="stonith-enabled"
      <nvpair id="cib-bootstrap-options-stonith-action" name="stonith-action" va
      <nvpair id="cib-bootstrap-options-expected-quorum-votes" name="expected-qu
      ...
    </cluster_property_set>
  </crm_config>
  <nodes>
    <node id="prod-mysql1" uname="prod-mysql1">
    <node id="prod-mysql2" uname="prod-mysql2"/>
    <node id="prod-mysql-repl" uname="prod-mysql-repl"/>
      <instance_attributes id="prod-mysql-repl">
        <nvpair id="prod-mysql-repl-standby" name="standby" value="on"/>
      </instance_attributes>
    </node>
  </nodes>
  <resources>
    <primitive class="stonith" id="fence_prod-mysql1_ipmi" type="fence_ipmilan">
      <instance_attributes id="fence_prod-mysql1_ipmi-instance_attributes">
        <nvpair id="fence_prod-mysql1_ipmi-instance_attributes-ipaddr" name="ipa
        <nvpair id="fence_prod-mysql1_ipmi-instance_attributes-action" name="act
        <nvpair id="fence_prod-mysql1_ipmi-instance_attributes-login" name="logi
        <nvpair id="fence_prod-mysql1_ipmi-instance_attributes-passwd" name="pas
        <nvpair id="fence_prod-mysql1_ipmi-instance_attributes-verbose" name="ve
        <nvpair id="fence_prod-mysql1_ipmi-instance_attributes-pcmk_host_list" n
        <nvpair id="fence_prod-mysql1_ipmi-instance_attributes-lanplus" name="la
      </instance_attributes>
    </primitive>
    <primitive class="stonith" id="fence_prod-mysql2_ipmi" type="fence_ipmilan">
      <instance_attributes id="fence_prod-mysql2_ipmi-instance_attributes">
        <nvpair id="fence_prod-mysql2_ipmi-instance_attributes-ipaddr" name="ipa
        <nvpair id="fence_prod-mysql2_ipmi-instance_attributes-action" name="act
        <nvpair id="fence_prod-mysql2_ipmi-instance_attributes-login" name="logi
        <nvpair id="fence_prod-mysql2_ipmi-instance_attributes-passwd" name="pas
        <nvpair id="fence_prod-mysql2_ipmi-instance_attributes-verbose" name="ve
        <nvpair id="fence_prod-mysql2_ipmi-instance_attributes-pcmk_host_list" n
        <nvpair id="fence_prod-mysql2_ipmi-instance_attributes-lanplus" name="la
      </instance_attributes>
    </primitive>
    <primitive class="stonith" id="fence_prod-mysql1_apc1" type="fence_apc_snmp">
      <instance_attributes id="fence_prod-mysql1_apc1-instance_attributes">
        <nvpair id="fence_prod-mysql1_apc1-instance_attributes-ipaddr" name="ipa

```

```

        <nvpair id="fence_prod-mysql1_apc1-instance_attributes-action" name="act
        <nvpair id="fence_prod-mysql1_apc1-instance_attributes-port" name="port"
        <nvpair id="fence_prod-mysql1_apc1-instance_attributes-login" name="logi
        <nvpair id="fence_prod-mysql1_apc1-instance_attributes-passwd" name="pas
        <nvpair id="fence_prod-mysql1_apc1-instance_attributes-pcmk_host_list" n
    </instance_attributes>
</primitive>
<primitive class="stonith" id="fence_prod-mysql1_apc2" type="fence_apc_snmp"
    <instance_attributes id="fence_prod-mysql1_apc2-instance_attributes">
        <nvpair id="fence_prod-mysql1_apc2-instance_attributes-ipaddr" name="ipa
        <nvpair id="fence_prod-mysql1_apc2-instance_attributes-action" name="act
        <nvpair id="fence_prod-mysql1_apc2-instance_attributes-port" name="port"
        <nvpair id="fence_prod-mysql1_apc2-instance_attributes-login" name="logi
        <nvpair id="fence_prod-mysql1_apc2-instance_attributes-passwd" name="pas
        <nvpair id="fence_prod-mysql1_apc2-instance_attributes-pcmk_host_list" n
    </instance_attributes>
</primitive>
<primitive class="stonith" id="fence_prod-mysql2_apc1" type="fence_apc_snmp"
    <instance_attributes id="fence_prod-mysql2_apc1-instance_attributes">
        <nvpair id="fence_prod-mysql2_apc1-instance_attributes-ipaddr" name="ipa
        <nvpair id="fence_prod-mysql2_apc1-instance_attributes-action" name="act
        <nvpair id="fence_prod-mysql2_apc1-instance_attributes-port" name="port"
        <nvpair id="fence_prod-mysql2_apc1-instance_attributes-login" name="logi
        <nvpair id="fence_prod-mysql2_apc1-instance_attributes-passwd" name="pas
        <nvpair id="fence_prod-mysql2_apc1-instance_attributes-pcmk_host_list" n
    </instance_attributes>
</primitive>
<primitive class="stonith" id="fence_prod-mysql2_apc2" type="fence_apc_snmp"
    <instance_attributes id="fence_prod-mysql2_apc2-instance_attributes">
        <nvpair id="fence_prod-mysql2_apc2-instance_attributes-ipaddr" name="ipa
        <nvpair id="fence_prod-mysql2_apc2-instance_attributes-action" name="act
        <nvpair id="fence_prod-mysql2_apc2-instance_attributes-port" name="port"
        <nvpair id="fence_prod-mysql2_apc2-instance_attributes-login" name="logi
        <nvpair id="fence_prod-mysql2_apc2-instance_attributes-passwd" name="pas
        <nvpair id="fence_prod-mysql2_apc2-instance_attributes-pcmk_host_list" n
    </instance_attributes>
</primitive>
</resources>
<constraints>
    <rsc_location id="l_fence_prod-mysql1_ipmi" node="prod-mysql1" rsc="fence_pr
    <rsc_location id="l_fence_prod-mysql2_ipmi" node="prod-mysql2" rsc="fence_pr
    <rsc_location id="l_fence_prod-mysql1_apc2" node="prod-mysql1" rsc="fence_pr
    <rsc_location id="l_fence_prod-mysql1_apc1" node="prod-mysql1" rsc="fence_pr
    <rsc_location id="l_fence_prod-mysql2_apc1" node="prod-mysql2" rsc="fence_pr
    <rsc_location id="l_fence_prod-mysql2_apc2" node="prod-mysql2" rsc="fence_pr
</constraints>
<fencing-topology>
    <fencing-level devices="fence_prod-mysql1_ipmi" id="fencing-2" index="1" tar
    <fencing-level devices="fence_prod-mysql1_apc1,fence_prod-mysql1_apc2" id="f
    <fencing-level devices="fence_prod-mysql2_ipmi" id="fencing-0" index="1" tar
    <fencing-level devices="fence_prod-mysql2_apc1,fence_prod-mysql2_apc2" id="f
</fencing-topology>
...
</configuration>

```

</cib>

Remapping Reboots

When the cluster needs to reboot a node, whether because `stonith-action` is `reboot` or because a reboot was manually requested (such as by `stonith_admin --reboot`), it will remap that to other commands in two cases:

1. If the chosen fencing device does not support the `reboot` command, the cluster will ask it to perform `off` instead.
2. If a fencing topology level with multiple devices must be executed, the cluster will ask all the devices to perform `off`, then ask the devices to perform `on`.

To understand the second case, consider the example of a node with redundant power supplies connected to intelligent power switches. Rebooting one switch and then the other would have no effect on the node. Turning both switches off, and then on, actually reboots the node.

In such a case, the fencing operation will be treated as successful as long as the `off` commands succeed, because then it is safe for the cluster to recover any resources that were on the node. Timeouts and errors in the `on` phase will be logged but ignored.

When a reboot operation is remapped, any action-specific timeout for the remapped action will be used (for example, `pcmk_off_timeout` will be used when executing the `off` command, not `pcmk_reboot_timeout`).

Chapter 7. Alerts

Alerts may be configured to take some external action when a cluster event occurs (node failure, resource starting or stopping, etc.).

Alert Agents

As with resource agents, the cluster calls an external program (an *alert agent*) to handle alerts. The cluster passes information about the event to the agent via environment variables. Agents can do anything desired with this information (send an e-mail, log to a file, update a monitoring system, etc.).

Example 7.1. Simple alert configuration

```
<configuration>
  <alerts>
    <alert id="my-alert" path="/path/to/my-script.sh" />
  </alerts>
</configuration>
```

In the example above, the cluster will call `my-script.sh` for each event.

Multiple alert agents may be configured; the cluster will call all of them for each event.

Alert agents will be called only on cluster nodes. They will be called for events involving Pacemaker Remote nodes, but they will never be called *on* those nodes.

Alert Recipients

Usually alerts are directed towards a recipient. Thus each alert may be additionally configured with one or more recipients. The cluster will call the agent separately for each recipient.

Example 7.2. Alert configuration with recipient

```
<configuration>
  <alerts>
    <alert id="my-alert" path="/path/to/my-script.sh">
      <recipient id="my-alert-recipient" value="some-address"/>
    </alert>
  </alerts>
</configuration>
```

In the above example, the cluster will call `my-script.sh` for each event, passing the recipient `some-address` as an environment variable.

The recipient may be anything the alert agent can recognize — an IP address, an e-mail address, a file name, whatever the particular agent supports.

Alert Meta-Attributes

As with resource agents, meta-attributes can be configured for alert agents to affect how Pacemaker calls them.

Table 7.1. Meta-Attributes of an Alert

Meta-Attribute	Default	Description
timestamp-format	%H:%M:%S.%06N	Format the cluster will use when sending the event's timestamp to the agent. This is a string as used with the <code>date(1)</code> command.
timeout	30s	If the alert agent does not complete within this amount of time, it will be terminated.

Meta-attributes can be configured per alert agent and/or per recipient.

Example 7.3. Alert configuration with meta-attributes

```
<configuration>
  <alerts>
    <alert id="my-alert" path="/path/to/my-script.sh">
      <meta_attributes id="my-alert-attributes">
        <nvpair id="my-alert-attributes-timeout" name="timeout"
          value="15s"/>
      </meta_attributes>
      <recipient id="my-alert-recipient1" value="someuser@example.com">
        <meta_attributes id="my-alert-recipient1-attributes">
          <nvpair id="my-alert-recipient1-timestamp-format"
            name="timestamp-format" value="%D %H:%M"/>
        </meta_attributes>
      </recipient>
      <recipient id="my-alert-recipient2" value="otheruser@example.com">
        <meta_attributes id="my-alert-recipient2-attributes">
          <nvpair id="my-alert-recipient2-timestamp-format"
            name="timestamp-format" value="%c"/>
        </meta_attributes>
      </recipient>
    </alert>
  </alerts>
</configuration>
```

In the above example, the `my-script.sh` will get called twice for each event, with each call using a 15-second timeout. One call will be passed the recipient `someuser@example.com` and a timestamp in the format `%D %H:%M`, while the other call will be passed the recipient `otheruser@example.com` and a timestamp in the format `%c`.

Alert Instance Attributes

As with resource agents, agent-specific configuration values may be configured as instance attributes. These will be passed to the agent as additional environment variables. The number, names and allowed values of these instance attributes are completely up to the particular agent.

Example 7.4. Alert configuration with instance attributes

```
<configuration>
  <alerts>
    <alert id="my-alert" path="/path/to/my-script.sh">
      <meta_attributes id="my-alert-attributes">
        <nvpair id="my-alert-attributes-timeout" name="timeout"
```



```
        value="15s"/>
    </meta_attributes>
    <instance_attributes id="my-alert-options">
        <nvpair id="my-alert-options-debug" name="debug" value="false"/>
    </instance_attributes>
    <recipient id="my-alert-recipient1" value="someuser@example.com"/>
</alert>
</alerts>
</configuration>
```

Alert Filters

By default, an alert agent will be called for node events, fencing events, and resource events. An agent may choose to ignore certain types of events, but there is still the overhead of calling it for those events. To eliminate that overhead, you may select which types of events the agent should receive.

Example 7.5. Alert configuration to receive only node events and fencing events

```
<configuration>
  <alerts>
    <alert id="my-alert" path="/path/to/my-script.sh">
      <select>
        <select_nodes />
        <select_fencing />
      </select>
      <recipient id="my-alert-recipient1" value="someuser@example.com"/>
    </alert>
  </alerts>
</configuration>
```

The possible options within `<select>` are `<select_nodes>`, `<select_fencing>`, `<select_resources>`, and `<select_attributes>`.

With `<select_attributes>` (the only event type not enabled by default), the agent will receive alerts when a node attribute changes. If you wish the agent to be called only when certain attributes change, you can configure that as well.

Example 7.6. Alert configuration to be called when certain node attributes change

```
<configuration>
  <alerts>
    <alert id="my-alert" path="/path/to/my-script.sh">
      <select>
        <select_attributes>
          <attribute id="alert-standby" name="standby" />
          <attribute id="alert-shutdown" name="shutdown" />
        </select_attributes>
      </select>
      <recipient id="my-alert-recipient1" value="someuser@example.com"/>
    </alert>
  </alerts>
</configuration>
```

Node attribute alerts are currently considered experimental. Alerts may be limited to attributes set via `attrd_updater`, and agents may be called multiple times with the same attribute value.

Using the Sample Alert Agents

Pacemaker provides several sample alert agents, installed in `/usr/share/pacemaker/alerts` by default.

While these sample scripts may be copied and used as-is, they are provided mainly as templates to be edited to suit your purposes. See their source code for the full set of instance attributes they support.

Example 7.7. Sending cluster events as SNMP traps

```
<configuration>
  <alerts>
    <alert id="snmp_alert" path="/path/to/alert_snmp.sh">
      <instance_attributes id="config_for_alert_snmp">
        <nvpair id="trap_node_states" name="trap_node_states" value="all"/>
      </instance_attributes>
      <meta_attributes id="config_for_timestamp">
        <nvpair id="ts_fmt" name="timestamp-format"
          value="%Y-%m-%d,%H:%M:%S.%01N"/>
      </meta_attributes>
      <recipient id="snmp_destination" value="192.168.1.2"/>
    </alert>
  </alerts>
</configuration>
```

Example 7.8. Sending cluster events as e-mails

```
<configuration>
  <alerts>
    <alert id="smtp_alert" path="/path/to/alert_smtp.sh">
      <instance_attributes id="config_for_alert_smtp">
        <nvpair id="email_sender" name="email_sender"
          value="donotreply@example.com"/>
      </instance_attributes>
      <recipient id="smtp_destination" value="admin@example.com"/>
    </alert>
  </alerts>
</configuration>
```

Writing an Alert Agent

Table 7.2. Environment variables passed to alert agents

Environment Variable	Description
CRM_alert_kind	The type of alert (node, fencing, resource, or attribute)
CRM_alert_version	The version of Pacemaker sending the alert
CRM_alert_recipient	The configured recipient
CRM_alert_node_sequence	A sequence number increased whenever an alert is being issued on the local node, which can be used to reference the order in which alerts have been issued by Pacemaker. An alert for an

Environment Variable	Description
	event that happened later in time reliably has a higher sequence number than alerts for earlier events. Be aware that this number has no cluster-wide meaning.
CRM_alert_timestamp	A timestamp created prior to executing the agent, in the format specified by the <code>timestamp-format</code> meta-attribute. This allows the agent to have a reliable, high-precision time of when the event occurred, regardless of when the agent itself was invoked (which could potentially be delayed due to system load, etc.).
CRM_alert_timestamp_epoch	The same time as <code>CRM_alert_timestamp</code> , expressed as the integer number of seconds since January 1, 1970. This (along with <code>CRM_alert_timestamp_usec</code>) can be useful for alert agents that need to format time in a specific way rather than let the user configure it.
CRM_alert_timestamp_usec	The same time as <code>CRM_alert_timestamp</code> , expressed as the integer number of microseconds since <code>CRM_alert_timestamp_epoch</code> .
CRM_alert_node	Name of affected node
CRM_alert_desc	Detail about event. For node alerts, this is the node's current state (member or lost). For fencing alerts, this is a summary of the requested fencing operation, including origin, target, and fencing operation error code, if any. For resource alerts, this is a readable string equivalent of <code>CRM_alert_status</code> .
CRM_alert_nodeid	ID of node whose status changed (provided with node alerts only)
CRM_alert_task	The requested fencing or resource operation (provided with fencing and resource alerts only)
CRM_alert_rc	The numerical return code of the fencing or resource operation (provided with fencing and resource alerts only)
CRM_alert_rsc	The name of the affected resource (resource alerts only)
CRM_alert_interval	The interval of the resource operation (resource alerts only)
CRM_alert_target_rc	The expected numerical return code of the operation (resource alerts only)
CRM_alert_status	A numerical code used by Pacemaker to represent the operation result (resource alerts only)
CRM_alert_exec_time	The (wall-clock) time, in milliseconds, that it took to execute the action. If the action timed out, <code>CRM_alert_status</code> will be 2, <code>CRM_alert_desc</code> will be "Timed Out", and this value will be the action timeout. May not be supported on all platforms. (resource alerts only)
CRM_alert_attribute_name	The name of the node attribute that changed (attribute alerts only)
CRM_alert_attribute_value	The new value of the node attribute that changed (attribute alerts only)

Special concerns when writing alert agents:

- Alert agents may be called with no recipient (if none is configured), so the agent must be able to handle this situation, even if it only exits in that case. (Users may modify the configuration in stages, and add a recipient later.)
- If more than one recipient is configured for an alert, the alert agent will be called once per recipient. If an agent is not able to run concurrently, it should be configured with only a single recipient. The agent is free, however, to interpret the recipient as a list.
- When a cluster event occurs, all alerts are fired off at the same time as separate processes. Depending on how many alerts and recipients are configured, and on what is done within the alert agents, a significant load burst may occur. The agent could be written to take this into consideration, for example by queueing resource-intensive actions into some other instance, instead of directly executing them.
- Alert agents are run as the `hacluster` user, which has a minimal set of permissions. If an agent requires additional privileges, it is recommended to configure `sudo` to allow the agent to run the necessary commands as another user with the appropriate privileges.
- As always, take care to validate and sanitize user-configured parameters, such as `CRM_alert_timestamp` (whose content is specified by the user-configured `timestamp-format`), `CRM_alert_recipient`, and all instance attributes. Mostly this is needed simply to protect against configuration errors, but if some user can modify the CIB without having `hacluster`-level access to the cluster nodes, it is a potential security concern as well, to avoid the possibility of code injection.

Note

The alerts interface is designed to be backward compatible with the external scripts interface used by the `ocf:pacemaker:ClusterMon` resource, which is now deprecated. To preserve this compatibility, the environment variables passed to alert agents are available prepended with `CRM_notify_` as well as `CRM_alert_`. One break in compatibility is that `ClusterMon` ran external scripts as the `root` user, while alert agents are run as the `hacluster` user.

Chapter 8. Rules

Rules can be used to make your configuration more dynamic, allowing values to change depending on the time or the value of a node attribute. Examples of things rules are useful for:

- Set a higher value for `resource-stickiness` during working hours, to minimize downtime, and a lower value on weekends, to allow resources to move to their most preferred locations when people aren't around to notice.
- Automatically place the cluster into maintenance mode during a scheduled maintenance window.
- Assign certain nodes and resources to a particular department via custom node attributes and meta-attributes, and add a single location constraint that restricts the department's resources to run only on those nodes.

Each constraint type or property set that supports rules may contain one or more `rule` elements specifying conditions under which the constraint or properties take effect. Examples later in this section will make this clearer.

Rule Properties

Table 8.1. Attributes of a rule Element

Attribute	Default	Description
<code>id</code>		A unique name for the rule (required)
<code>role</code>	<code>Started</code>	The rule is in effect only when the resource is in the specified role. Allowed values are <code>Started</code> , <code>Slave</code> , and <code>Master</code> . A rule with <code>role="Master"</code> cannot determine the initial location of a clone instance and will only affect which of the active instances will be promoted.
<code>score</code>		If this rule is used in a location constraint and evaluates to true, apply this score to the constraint. Only one of <code>score</code> and <code>score-attribute</code> may be used.
<code>score-attribute</code>		If this rule is used in a location constraint and evaluates to true, use the value of this node attribute as the score to apply to the constraint. Only one of <code>score</code> and <code>score-attribute</code> may be used.
<code>boolean-op</code>	<code>and</code>	If this rule contains more than one condition, a value of <code>and</code> specifies that the rule evaluates to true only if all conditions are true, and a value of <code>or</code> specifies that the rule evaluates to true if any condition is true.

A rule element must contain one or more conditions. A condition may be an `expression` element, a `date_expression` element, or another rule element.

Node Attribute Expressions

Expressions are rule conditions based on the values of node attributes.

Table 8.2. Attributes of an expression Element

Field	Default	Description
id		A unique name for the expression (required)
attribute		The node attribute to test (required)
type	string	How the node attributes should be compared. Allowed values are <code>string</code> , <code>integer</code> , and <code>version</code> .
operation		<p>The comparison to perform (required). Allowed values:</p> <ul style="list-style-type: none"> • <code>lt</code>: True if the node attribute value is less than the comparison value • <code>gt</code>: True if the node attribute value is greater than the comparison value • <code>lte</code>: True if the node attribute value is less than or equal to the comparison value • <code>gte</code>: True if the node attribute value is greater than or equal to the comparison value • <code>eq</code>: True if the node attribute value is equal to the comparison value • <code>ne</code>: True if the node attribute value is not equal to the comparison value • <code>defined</code>: True if the node has the named attribute • <code>not_defined</code>: True if the node does not have the named attribute
value		User-supplied value for comparison (required for operations other than <code>defined</code> and <code>not_defined</code>)
value-source	literal	<p>How the value is derived. Allowed values:</p> <ul style="list-style-type: none"> • <code>literal</code>: value is a literal string to compare against • <code>param</code>: value is the name of a resource parameter to compare against (only valid in location constraints) • <code>meta</code>: value is the name of a resource meta-attribute to compare against (only valid in location constraints)

In addition to custom node attributes defined by the administrator, the cluster defines special, built-in node attributes for each node that can also be used in rule expressions.

Table 8.3. Built-in Node Attributes

Name	Value
#uname	Node name
#id	Node ID

Name	Value
#kind	Node type. Possible values are <code>cluster</code> , <code>remote</code> , and <code>container</code> . Kind is <code>remote</code> for Pacemaker Remote nodes created with the <code>ocf:pacemaker:remote</code> resource, and <code>container</code> for Pacemaker Remote guest nodes and bundle nodes
#is_dc	"true" if this node is a Designated Controller (DC), "false" otherwise
#cluster-name	The value of the <code>cluster-name</code> cluster property, if set
#site-name	The value of the <code>site-name</code> node attribute, if set, otherwise identical to <code>#cluster-name</code>
#role	The role the relevant promotable clone resource has on this node. Valid only within a rule for a location constraint for a promotable clone resource.

Date/Time Expressions

Date/time expressions are rule conditions based (as the name suggests) on the current date and time.

A `date_expression` element may optionally contain a `date_spec` or `duration` element depending on the context.

Table 8.4. Attributes of a `date_expression` Element

Field	Description
<code>id</code>	A unique name for the expression (required)
<code>start</code>	A date/time conforming to the ISO8601 [http://en.wikipedia.org/wiki/ISO_8601] specification. May be used when operation is <code>in_range</code> (in which case at least one of <code>start</code> or <code>end</code> must be specified) or <code>gt</code> (in which case <code>start</code> is required).
<code>end</code>	A date/time conforming to the ISO8601 [http://en.wikipedia.org/wiki/ISO_8601] specification. May be used when operation is <code>in_range</code> (in which case at least one of <code>start</code> or <code>end</code> must be specified) or <code>lt</code> (in which case <code>end</code> is required).
<code>operation</code>	Compares the current date/time with the start and/or end date, depending on the context. Allowed values: <ul style="list-style-type: none"> <code>gt</code>: True if the current date/time is after <code>start</code> <code>lt</code>: True if the current date/time is before <code>end</code> <code>in_range</code>: True if the current date/time is after <code>start</code> (if specified) and before either <code>end</code> (if specified) or <code>start</code> plus the value of the <code>duration</code> element (if one is contained in the <code>date_expression</code>) <code>date_spec</code>: True if the current date/time matches the specification given in the contained <code>date_spec</code> element (described below)

Note

There is no `eq`, `neq`, `gte`, or `lte` operation, since they would be valid only for a single second.

Date Specifications

A `date_spec` element is used to create a cron-like expression relating to time. Each field can contain a single number or range. Any field not supplied is ignored.

Table 8.5. Attributes of a `date_spec` Element

Field	Description
<code>id</code>	A unique name for the object (required)
<code>hours</code>	Allowed values: 0-23 (where 0 is midnight and 23 is 11 p.m.)
<code>monthdays</code>	Allowed values: 1-31 (depending on month and year)
<code>weekdays</code>	Allowed values: 1-7 (where 1 is Monday and 7 is Sunday)
<code>yeardays</code>	Allowed values: 1-366 (depending on the year)
<code>months</code>	Allowed values: 1-12
<code>weeks</code>	Allowed values: 1-53 (depending on weekyear)
<code>years</code>	Year according to the Gregorian calendar
<code>weekyears</code>	Year in which the week started; for example, 1 January 2005 can be specified in ISO 8601 as <i>2005-001 Ordinal</i> , <i>2005-01-01 Gregorian</i> or <i>2004-W53-6 Weekly</i> and thus would match <code>years="2005"</code> or <code>weekyears="2004"</code>
<code>moon</code>	Allowed values are 0-7 (where 0 is the new moon and 4 is full moon). Seriously, you can use this. This was implemented to demonstrate the ease with which new comparisons could be added.

For example, `monthdays="1"` matches the first day of every month, and `hours="09-17"` matches the hours between 9 a.m. and 5 p.m. (inclusive).

At this time, multiple ranges (e.g. `weekdays="1, 2"` or `weekdays="1-2, 5-6"`) are not supported.

Note

Pacemaker can calculate when evaluation of a `date_expression` with an operation of `gt`, `lt`, or `in_range` will next change, and schedule a cluster re-check for that time. However, it does not do this for `date_spec`. Instead, it evaluates the `date_spec` whenever a cluster re-check naturally happens via a cluster event or the `cluster-recheck-interval` cluster option. For example, if you have a `date_spec` enabling a resource from 9 a.m. to 5 p.m., and `cluster-recheck-interval` has been set to 5 minutes, then sometime between 9 a.m. and 9:05 a.m. the cluster would notice that it needs to start the resource, and sometime between 5 p.m. and 5:05 p.m. it would realize that it needs to stop the resource. The timing of the actual start and stop actions will further depend on factors such as any other actions the cluster may need to perform first, and the load of the machine.

Durations

A duration is used to calculate a value for end when one is not supplied to `in_range` operations. It contains one or more attributes each containing a single number. Any attribute not supplied is ignored.

Table 8.6. Attributes of a duration Element

Field	Description
id	A unique name for this duration element (required)
seconds	This many seconds will be added to the total duration
minutes	This many minutes will be added to the total duration
hours	This many hours will be added to the total duration
weeks	This many weeks will be added to the total duration
months	This many months will be added to the total duration
years	This many years will be added to the total duration

Example Time-Based Expressions

A small sample of how time-based expressions can be used:

Example 8.1. True if now is any time in the year 2005

```
<rule id="rule1" score="INFINITY">
  <date_expression id="date_expr1" start="2005-001" operation="in_range">
    <duration id="duration1" years="1"/>
  </date_expression>
</rule>
```

Example 8.2. Equivalent expression

```
<rule id="rule2" score="INFINITY">
  <date_expression id="date_expr2" operation="date_spec">
    <date_spec id="date_spec2" years="2005"/>
  </date_expression>
</rule>
```

Example 8.3. 9am-5pm Monday-Friday

```
<rule id="rule3" score="INFINITY">
  <date_expression id="date_expr3" operation="date_spec">
    <date_spec id="date_spec3" hours="9-16" weekdays="1-5"/>
  </date_expression>
</rule>
```

Please note that the 16 matches up to 16:59:59, as the numeric value (hour) still matches!

Example 8.4. 9am-6pm Monday through Friday or anytime Saturday

```
<rule id="rule4" score="INFINITY" boolean-op="or">
  <date_expression id="date_expr4-1" operation="date_spec">
    <date_spec id="date_spec4-1" hours="9-16" weekdays="1-5"/>
  </date_expression>
```

```
<date_expression id="date_expr4-2" operation="date_spec">
  <date_spec id="date_spec4-2" weekdays="6"/>
</date_expression>
</rule>
```

Example 8.5. 9am-5pm or 9pm-12am Monday through Friday

```
<rule id="rule5" score="INFINITY" boolean-op="and">
  <rule id="rule5-nested1" score="INFINITY" boolean-op="or">
    <date_expression id="date_expr5-1" operation="date_spec">
      <date_spec id="date_spec5-1" hours="9-16"/>
    </date_expression>
    <date_expression id="date_expr5-2" operation="date_spec">
      <date_spec id="date_spec5-2" hours="21-23"/>
    </date_expression>
  </rule>
  <date_expression id="date_expr5-3" operation="date_spec">
    <date_spec id="date_spec5-3" weekdays="1-5"/>
  </date_expression>
</rule>
```

Example 8.6. Mondays in March 2005

```
<rule id="rule6" score="INFINITY" boolean-op="and">
  <date_expression id="date_expr6-1" operation="date_spec">
    <date_spec id="date_spec6" weekdays="1"/>
  </date_expression>
  <date_expression id="date_expr6-2" operation="in_range"
    start="2005-03-01" end="2005-04-01"/>
</rule>
```

Note

Because no time is specified with the above dates, 00:00:00 is implied. This means that the range includes all of 2005-03-01 but none of 2005-04-01. You may wish to write `end="2005-03-31T23:59:59"` to avoid confusion.

Example 8.7. A full moon on Friday the 13th

```
<rule id="rule7" score="INFINITY" boolean-op="and">
  <date_expression id="date_expr7" operation="date_spec">
    <date_spec id="date_spec7" weekdays="5" monthdays="13" moon="4"/>
  </date_expression>
</rule>
```

Using Rules to Determine Resource Location

A location constraint may contain one or more top-level rules. The cluster will act as if there is a separate location constraint for each rule that evaluates as true.

Consider the following simple location constraint:

Example 8.8. Prevent resource "webserver" from running on node3

```
<rsc_location id="ban-apache-on-node3" rsc="webserver"
              score="-INFINITY" node="node3" />
```

The constraint can be more verbosely written using a rule:

Example 8.9. Prevent resource "webserver" from running on node3 using rule

```
<rsc_location id="ban-apache-on-node3" rsc="webserver">
  <rule id="ban-apache-rule" score="-INFINITY">
    <expression id="ban-apache-expr" attribute="#uname"
               operation="eq" value="node3" />
  </rule>
</rsc_location>
```

The advantage of using the expanded form is that one could add more expressions (for example, limiting the constraint to certain days of the week), or activate the constraint by some node attribute other than node name.

Location Rules Based on Other Node Properties

The expanded form allows us to match on node properties other than its name. If we rated each machine's CPU power such that the cluster had the following nodes section:

Example 8.10. A sample nodes section for use with score-attribute

```
<nodes>
  <node id="uuid1" uname="c001n01" type="normal">
    <instance_attributes id="uuid1-custom_attrs">
      <nvpair id="uuid1-cpu_mips" name="cpu_mips" value="1234" />
    </instance_attributes>
  </node>
  <node id="uuid2" uname="c001n02" type="normal">
    <instance_attributes id="uuid2-custom_attrs">
      <nvpair id="uuid2-cpu_mips" name="cpu_mips" value="5678" />
    </instance_attributes>
  </node>
</nodes>
```

then we could prevent resources from running on underpowered machines with this rule:

```
<rule id="need-more-power-rule" score="-INFINITY">
  <expression id="need-more-power-expr" attribute="cpu_mips"
             operation="lt" value="3000" />
</rule>
```

Using score-attribute Instead of score

When using `score-attribute` instead of `score`, each node matched by the rule has its score adjusted differently, according to its value for the named node attribute. Thus, in the previous example, if a rule used `score-attribute="cpu_mips"`, `c001n01` would have its preference to run the resource increased by 1234 whereas `c001n02` would have its preference increased by 5678.

Using Rules to Define Options

Rules may be used to control a variety of options:

- Cluster options (`cluster_property_set` elements)
- Node attributes (as `instance_attributes` or `utilization` elements inside a `node` element)
- Resource options (as `utilization`, `meta_attributes`, or `instance_attributes` elements inside a resource definition element or `op`, `rsc_defaults`, `op_defaults`, or `template` element)
- Operation properties (`meta_attributes` inside an `op` or `op_defaults` element)

Using Rules to Control Resource Options

Often some cluster nodes will be different from their peers. Sometimes, these differences — e.g. the location of a binary or the names of network interfaces — require resources to be configured differently depending on the machine they're hosted on.

By defining multiple `instance_attributes` objects for the resource and adding a rule to each, we can easily handle these special cases.

In the example below, `mySpecialRsc` will use `eth1` and port 9999 when run on `node1`, `eth2` and port 8888 on `node2` and default to `eth0` and port 9999 for all other nodes.

Example 8.11. Defining different resource options based on the node name

```
<primitive id="mySpecialRsc" class="ocf" type="Special" provider="me">
  <instance_attributes id="special-node1" score="3">
    <rule id="node1-special-case" score="INFINITY" >
      <expression id="node1-special-case-expr" attribute="#uname"
        operation="eq" value="node1"/>
    </rule>
    <nvpair id="node1-interface" name="interface" value="eth1"/>
  </instance_attributes>
  <instance_attributes id="special-node2" score="2" >
    <rule id="node2-special-case" score="INFINITY">
      <expression id="node2-special-case-expr" attribute="#uname"
        operation="eq" value="node2"/>
    </rule>
    <nvpair id="node2-interface" name="interface" value="eth2"/>
    <nvpair id="node2-port" name="port" value="8888"/>
  </instance_attributes>
  <instance_attributes id="defaults" score="1" >
    <nvpair id="default-interface" name="interface" value="eth0"/>
    <nvpair id="default-port" name="port" value="9999"/>
  </instance_attributes>
</primitive>
```

The order in which `instance_attributes` objects are evaluated is determined by their score (highest to lowest). If not supplied, score defaults to zero, and objects with an equal score are processed in listed order. If the `instance_attributes` object has no rule or a rule that evaluates to `true`, then for any parameter the resource does not yet have a value for, the resource will use the parameter values defined by the `instance_attributes`.

For example, given the configuration above, if the resource is placed on node1:

1. `special-node1` has the highest score (3) and so is evaluated first; its rule evaluates to `true`, so `interface` is set to `eth1`.
2. `special-node2` is evaluated next with score 2, but its rule evaluates to `false`, so it is ignored.
3. `defaults` is evaluated last with score 1, and has no rule, so its values are examined; `interface` is already defined, so the value here is not used, but `port` is not yet defined, so `port` is set to 9999.

Using Rules to Control Resource Defaults

Rules can be used for resource and operation defaults. The following example illustrates how to set a different `resource-stickiness` value during and outside work hours. This allows resources to automatically move back to their most preferred hosts, but at a time that (in theory) does not interfere with business activities.

Example 8.12. Change `resource-stickiness` during working hours

```
<rsc_defaults>
  <meta_attributes id="core-hours" score="2">
    <rule id="core-hour-rule" score="0">
      <date_expression id="nine-to-five-Mon-to-Fri" operation="date_spec">
        <date_spec id="nine-to-five-Mon-to-Fri-spec" hours="9-16" weekdays="1-5">
        </date_spec>
      </date_expression>
    </rule>
    <nvpair id="core-stickiness" name="resource-stickiness" value="INFINITY"/>
  </meta_attributes>
  <meta_attributes id="after-hours" score="1">
    <nvpair id="after-stickiness" name="resource-stickiness" value="0"/>
  </meta_attributes>
</rsc_defaults>
```

Rules may be used similarly in `instance_attributes` or `utilization` blocks.

Any single block may directly contain only a single rule, but that rule may itself contain any number of rules.

Using Rules to Control Cluster Options

Controlling cluster options is achieved in much the same manner as specifying different resource options on different nodes.

The difference is that because they are cluster options, one cannot (or should not, because they won't work) use attribute-based expressions. The following example illustrates how to set `maintenance_mode` during a scheduled maintenance window. This will keep the cluster running but not monitor, start, or stop resources during this time.

Example 8.13. Schedule a maintenance window for 9 to 11 p.m. CDT Sept. 20, 2019

```
<crm_config>
  <cluster_property_set id="cib-bootstrap-options">
```

```
<nvpair id="bootstrap-stonith-enabled" name="stonith-enabled" value="1"/>
</cluster_property_set>
<cluster_property_set id="normal-set" score="10">
  <nvpair id="normal-maintenance-mode" name="maintenance-mode" value="false"/>
</cluster_property_set>
<cluster_property_set id="maintenance-window-set" score="1000">
  <nvpair id="maintenance-nvpair1" name="maintenance-mode" value="true"/>
  <rule id="maintenance-rule1" score="INFINITY">
    <date_expression id="maintenance-date1" operation="in_range"
      start="2019-09-20 21:00:00 -05:00" end="2019-09-20 23:00:00 -05:00"/>
  </rule>
</cluster_property_set>
</crm_config>
```

Important

The `cluster_property_set` with an `id` set to "cib-bootstrap-options" will *always* have the highest priority, regardless of any scores. Therefore, rules in another `cluster_property_set` can never take effect for any properties listed in the bootstrap set.

Chapter 9. Advanced Configuration

Specifying When Recurring Actions are Performed

By default, recurring actions are scheduled relative to when the resource started. So if your resource was last started at 14:32 and you have a backup set to be performed every 24 hours, then the backup will always run in the middle of the business day — hardly desirable.

To specify a date and time that the operation should be relative to, set the operation's `interval-origin`. The cluster uses this point to calculate the correct `start-delay` such that the operation will occur at $origin + (interval * N)$.

So, if the operation's interval is 24h, its `interval-origin` is set to 02:00 and it is currently 14:32, then the cluster would initiate the operation with a start delay of 11 hours and 28 minutes. If the resource is moved to another node before 2am, then the operation is cancelled.

The value specified for `interval` and `interval-origin` can be any date/time conforming to the ISO8601 standard [http://en.wikipedia.org/wiki/ISO_8601]. By way of example, to specify an operation that would run on the first Monday of 2009 and every Monday after that, you would add:

Example 9.1. Specifying a Base for Recurring Action Intervals

```
<op id="my-weekly-action" name="custom-action" interval="P7D" interval-origin="200
```

Handling Resource Failure

By default, Pacemaker will attempt to recover failed resources by restarting them. However, failure recovery is highly configurable.

Failure Counts

Pacemaker tracks resource failures for each combination of node, resource, and operation (start, stop, monitor, etc.).

You can query the fail count for a particular node, resource, and/or operation using the `crm_failcount` command. For example, to see how many times the 10-second monitor for `myrsc` has failed on `node1`, run:

```
# crm_failcount --query -r myrsc -N node1 -n monitor -I 10s
```

If you omit the node, `crm_failcount` will use the local node. If you omit the operation and interval, `crm_failcount` will display the sum of the fail counts for all operations on the resource.

You can use `crm_resource --cleanup` or `crm_failcount --delete` to clear fail counts. For example, to clear the above monitor failures, run:

```
# crm_resource --cleanup -r myrsc -N node1 -n monitor -I 10s
```

If you omit the resource, `crm_resource --cleanup` will clear failures for all resources. If you omit the node, it will clear failures on all nodes. If you omit the operation and interval, it will clear the failures for all operations on the resource.

Note

Even when cleaning up only a single operation, all failed operations will disappear from the status display. This allows us to trigger a re-check of the resource's current status.

Higher-level tools may provide other commands for querying and clearing fail counts.

The `crm_mon` tool shows the current cluster status, including any failed operations. To see the current fail counts for any failed resources, call `crm_mon` with the `--failcounts` option. This shows the fail counts per resource (that is, the sum of any operation fail counts for the resource).

Failure Response

Normally, if a running resource fails, pacemaker will try to stop it and start it again. Pacemaker will choose the best location to start it each time, which may be the same node that it failed on.

However, if a resource fails repeatedly, it is possible that there is an underlying problem on that node, and you might desire trying a different node in such a case. Pacemaker allows you to set your preference via the `migration-threshold` resource meta-attribute.¹

If you define `migration-threshold=N` for a resource, it will be banned from the original node after *N* failures.

Note

The `migration-threshold` is per *resource*, even though fail counts are tracked per *operation*. The operation fail counts are added together to compare against the `migration-threshold`.

By default, fail counts remain until manually cleared by an administrator using `crm_resource --cleanup` or `crm_failcount --delete` (hopefully after first fixing the failure's cause). It is possible to have fail counts expire automatically by setting the `failure-timeout` resource meta-attribute.

Important

A successful operation does not clear past failures. If a recurring monitor operation fails once, succeeds many times, then fails again days later, its fail count is 2. Fail counts are cleared only by manual intervention or failure timeout.

For example, a setting of `migration-threshold=2` and `failure-timeout=60s` would cause the resource to move to a new node after 2 failures, and allow it to move back (depending on stickiness and constraint scores) after one minute.

Note

`failure-timeout` is measured since the most recent failure. That is, older failures do not individually time out and lower the fail count. Instead, all failures are timed out simultaneously (and the fail count is reset to 0) if there is no new failure for the timeout period.

There are two exceptions to the migration threshold concept: when a resource either fails to start or fails to stop.

If the cluster property `start-failure-is-fatal` is set to `true` (which is the default), start failures cause the fail count to be set to `INFINITY` and thus always cause the resource to move immediately.

¹ The naming of this option was perhaps unfortunate as it is easily confused with live migration, the process of moving a resource from one node to another without stopping it. Xen virtual guests are the most common example of resources that can be migrated in this manner.

Stop failures are slightly different and crucial. If a resource fails to stop and STONITH is enabled, then the cluster will fence the node in order to be able to start the resource elsewhere. If STONITH is not enabled, then the cluster has no way to continue and will not try to start the resource elsewhere, but will try to stop it again after the failure timeout.

Moving Resources

Moving Resources Manually

There are primarily two occasions when you would want to move a resource from its current location: when the whole node is under maintenance, and when a single resource needs to be moved.

Standby Mode

Since everything eventually comes down to a score, you could create constraints for every resource to prevent them from running on one node. While pacemaker configuration can seem convoluted at times, not even we would require this of administrators.

Instead, one can set a special node attribute which tells the cluster "don't let anything run here". There is even a helpful tool to help query and set it, called `crm_standby`. To check the standby status of the current machine, run:

```
# crm_standby -G
```

A value of `on` indicates that the node is *not* able to host any resources, while a value of `off` says that it *can*.

You can also check the status of other nodes in the cluster by specifying the `--node` option:

```
# crm_standby -G --node sles-2
```

To change the current node's standby status, use `-v` instead of `-G`:

```
# crm_standby -v on
```

Again, you can change another host's value by supplying a hostname with `--node`.

A cluster node in standby mode will not run resources, but still contributes to quorum, and may fence or be fenced by nodes.

Moving One Resource

When only one resource is required to move, we could do this by creating location constraints. However, once again we provide a user-friendly shortcut as part of the `crm_resource` command, which creates and modifies the extra constraints for you. If `Email` were running on `sles-1` and you wanted it moved to a specific location, the command would look something like:

```
# crm_resource -M -r Email -H sles-2
```

Behind the scenes, the tool will create the following location constraint:

```
<rsc_location rsc="Email" node="sles-2" score="INFINITY"/>
```

It is important to note that subsequent invocations of `crm_resource -M` are not cumulative. So, if you ran these commands

```
# crm_resource -M -r Email -H sles-2
# crm_resource -M -r Email -H sles-3
```

then it is as if you had never performed the first command.

To allow the resource to move back again, use:

```
# crm_resource -U -r Email
```

Note the use of the word *allow*. The resource can move back to its original location but, depending on resource-stickiness, it might stay where it is. To be absolutely certain that it moves back to sles-1, move it there before issuing the call to `crm_resource -U`:

```
# crm_resource -M -r Email -H sles-1
# crm_resource -U -r Email
```

Alternatively, if you only care that the resource should be moved from its current location, try:

```
# crm_resource -B -r Email
```

Which will instead create a negative constraint, like

```
<rsc_location rsc="Email" node="sles-1" score="-INFINITY"/>
```

This will achieve the desired effect, but will also have long-term consequences. As the tool will warn you, the creation of a `-INFINITY` constraint will prevent the resource from running on that node until `crm_resource -U` is used. This includes the situation where every other cluster node is no longer available!

In some cases, such as when `resource-stickiness` is set to `INFINITY`, it is possible that you will end up with the problem described in the section called “What if Two Nodes Have the Same Score”. The tool can detect some of these cases and deals with them by creating both positive and negative constraints. E.g.

Email prefers sles-1 with a score of `-INFINITY`

Email prefers sles-2 with a score of `INFINITY`

which has the same long-term consequences as discussed earlier.

Moving Resources Due to Connectivity Changes

You can configure the cluster to move resources when external connectivity is lost in two steps.

Tell Pacemaker to Monitor Connectivity

First, add an **ocf:pacemaker:ping** resource to the cluster. The **ping** resource uses the system utility of the same name to test whether a list of machines (specified by DNS hostname or IPv4/IPv6 address) are reachable and uses the results to maintain a node attribute called `pingd` by default.²

Note

Older versions of Pacemaker used a different agent **ocf:pacemaker:pingd** which is now deprecated in favor of **ping**. If your version of Pacemaker does not contain the **ping** resource

² The attribute name is customizable, in order to allow multiple ping groups to be defined.

agent, download the latest version from <https://github.com/ClusterLabs/pacemaker/tree/master/extra/resources/ping>

Normally, the ping resource should run on all cluster nodes, which means that you'll need to create a clone. A template for this can be found below along with a description of the most interesting parameters.

Table 9.1. Common Options for a *ping* Resource

Field	Description
dampen	The time to wait (dampening) for further changes to occur. Use this to prevent a resource from bouncing around the cluster when cluster nodes notice the loss of connectivity at slightly different times.
multiplier	The number of connected ping nodes gets multiplied by this value to get a score. Useful when there are multiple ping nodes configured.
host_list	The machines to contact in order to determine the current connectivity status. Allowed values include resolvable DNS host names, IPv4 and IPv6 addresses.

Example 9.2. An example ping cluster resource that checks node connectivity once every minute

```
<clone id="Connected">
  <primitive id="ping" provider="pacemaker" class="ocf" type="ping">
    <instance_attributes id="ping-attrs">
      <nvpair id="pingd-dampen" name="dampen" value="5s"/>
      <nvpair id="pingd-multiplier" name="multiplier" value="1000"/>
      <nvpair id="pingd-hosts" name="host_list" value="my.gateway.com www.bigcorp."/>
    </instance_attributes>
    <operations>
      <op id="ping-monitor-60s" interval="60s" name="monitor"/>
    </operations>
  </primitive>
</clone>
```

Important

You're only half done. The next section deals with telling Pacemaker how to deal with the connectivity status that `ocf:pacemaker:ping` is recording.

Tell Pacemaker How to Interpret the Connectivity Data

Important

Before attempting the following, make sure you understand Chapter 8, Rules.

There are a number of ways to use the connectivity data.

The most common setup is for people to have a single ping target (e.g. the service network's default gateway), to prevent the cluster from running a resource on any unconnected node.

Example 9.3. Don't run a resource on unconnected nodes

```
<rsc_location id="WebServer-no-connectivity" rsc="Webserver">
```

```
<rule id="ping-exclude-rule" score="-INFINITY" >
  <expression id="ping-exclude" attribute="pingd" operation="not_defined"/>
</rule>
</rsc_location>
```

A more complex setup is to have a number of ping targets configured. You can require the cluster to only run resources on nodes that can connect to all (or a minimum subset) of them.

Example 9.4. Run only on nodes connected to three or more ping targets.

```
<primitive id="ping" provider="pacemaker" class="ocf" type="ping">
... <!-- omitting some configuration to highlight important parts -->
  <nvpair id="pingd-multiplier" name="multiplier" value="1000"/>
...
</primitive>
...
<rsc_location id="WebServer-connectivity" rsc="Webserver">
  <rule id="ping-prefer-rule" score="-INFINITY" >
    <expression id="ping-prefer" attribute="pingd" operation="lt" value="3000"/>
  </rule>
</rsc_location>
```

Alternatively, you can tell the cluster only to *prefer* nodes with the best connectivity. Just be sure to set multiplier to a value higher than that of resource-stickiness (and don't set either of them to INFINITY).

Example 9.5. Prefer the node with the most connected ping nodes

```
<rsc_location id="WebServer-connectivity" rsc="Webserver">
  <rule id="ping-prefer-rule" score-attribute="pingd" >
    <expression id="ping-prefer" attribute="pingd" operation="defined"/>
  </rule>
</rsc_location>
```

It is perhaps easier to think of this in terms of the simple constraints that the cluster translates it into. For example, if **sles-1** is connected to all five ping nodes but **sles-2** is only connected to two, then it would be as if you instead had the following constraints in your configuration:

Example 9.6. How the cluster translates the above location constraint

```
<rsc_location id="ping-1" rsc="Webserver" node="sles-1" score="5000"/>
<rsc_location id="ping-2" rsc="Webserver" node="sles-2" score="2000"/>
```

The advantage is that you don't have to manually update any constraints whenever your network connectivity changes.

You can also combine the concepts above into something even more complex. The example below shows how you can prefer the node with the most connected ping nodes provided they have connectivity to at least three (again assuming that multiplier is set to 1000).

Example 9.7. A more complex example of choosing a location based on connectivity

```
<rsc_location id="WebServer-connectivity" rsc="Webserver">
  <rule id="ping-exclude-rule" score="-INFINITY" >
    <expression id="ping-exclude" attribute="pingd" operation="lt" value="3000"/>
  </rule>
</rsc_location>
```

```
</rule>
<rule id="ping-prefer-rule" score-attribute="pingd" >
  <expression id="ping-prefer" attribute="pingd" operation="defined"/>
</rule>
</rsc_location>
```

Migrating Resources

Normally, when the cluster needs to move a resource, it fully restarts the resource (i.e. stops the resource on the current node and starts it on the new node).

However, some types of resources, such as Xen virtual guests, are able to move to another location without loss of state (often referred to as live migration or hot migration). In pacemaker, this is called resource migration. Pacemaker can be configured to migrate a resource when moving it, rather than restarting it.

Not all resources are able to migrate; see the Migration Checklist below, and those that can, won't do so in all situations. Conceptually, there are two requirements from which the other prerequisites follow:

- The resource must be active and healthy at the old location; and
- everything required for the resource to run must be available on both the old and new locations.

The cluster is able to accommodate both *push* and *pull* migration models by requiring the resource agent to support two special actions: `migrate_to` (performed on the current location) and `migrate_from` (performed on the destination).

In push migration, the process on the current location transfers the resource to the new location where it later activated. In this scenario, most of the work would be done in the `migrate_to` action and, if anything, the activation would occur during `migrate_from`.

Conversely for pull, the `migrate_to` action is practically empty and `migrate_from` does most of the work, extracting the relevant resource state from the old location and activating it.

There is no wrong or right way for a resource agent to implement migration, as long as it works.

Migration Checklist

- The resource may not be a clone.
- The resource must use an OCF style agent.
- The resource must not be in a failed or degraded state.
- The resource agent must support `migrate_to` and `migrate_from` actions, and advertise them in its metadata.
- The resource must have the `allow-migrate` meta-attribute set to `true` (which is not the default).

If an otherwise migratable resource depends on another resource via an ordering constraint, there are special situations in which it will be restarted rather than migrated.

For example, if the resource depends on a clone, and at the time the resource needs to be moved, the clone has instances that are stopping and instances that are starting, then the resource will be restarted. The scheduler is not yet able to model this situation correctly and so takes the safer (if less optimal) path.

Also, if a migratable resource depends on a non-migratable resource, and both need to be moved, the migratable resource will be restarted.

Tracking Node Health

A node may be functioning adequately as far as cluster membership is concerned, and yet be "unhealthy" in some respect that makes it an undesirable location for resources. For example, a disk drive may be reporting SMART errors, or the CPU may be highly loaded.

Pacemaker offers a way to automatically move resources off unhealthy nodes.

Node Health Attributes

Pacemaker will treat any node attribute whose name starts with `#health` as an indicator of node health. Node health attributes may have one of the following values:

Table 9.2. Allowed Values for Node Health Attributes

Value	Intended significance
red	This indicator is unhealthy
yellow	This indicator is becoming unhealthy
green	This indicator is healthy
<i>integer</i>	A numeric score to apply to all resources on this node (0 or positive is healthy, negative is unhealthy)

Node Health Strategy

Pacemaker assigns a node health score to each node, as the sum of the values of all its node health attributes. This score will be used as a location constraint applied to this node for all resources.

The `node-health-strategy` cluster option controls how Pacemaker responds to changes in node health attributes, and how it translates `red`, `yellow`, and `green` to scores.

Allowed values are:

Table 9.3. Node Health Strategies

Value	Effect
none	Do not track node health attributes at all.
migrate-on-red	Assign the value of <code>-INFINITY</code> to <code>red</code> , and 0 to <code>yellow</code> and <code>green</code> . This will cause all resources to move off the node if any attribute is <code>red</code> .
only-green	Assign the value of <code>-INFINITY</code> to <code>red</code> and <code>yellow</code> , and 0 to <code>green</code> . This will cause all resources to move off the node if any attribute is <code>red</code> or <code>yellow</code> .
progressive	Assign the value of the <code>node-health-red</code> cluster option to <code>red</code> , the value of <code>node-health-yellow</code> to <code>yellow</code> , and the value of <code>node-health-green</code> to <code>green</code> . Each node is additionally assigned a score of <code>node-health-base</code> (this allows resources to start even if some attributes are <code>yellow</code>). This strategy gives the administrator finer control over how important each value is.
custom	Track node health attributes using the same values as <code>progressive</code> for <code>red</code> , <code>yellow</code> , and <code>green</code> , but do not take them into account. The

Value	Effect
	administrator is expected to implement a policy by defining rules (see Chapter 8, Rules) referencing node health attributes.

Measuring Node Health

Since Pacemaker calculates node health based on node attributes, any method that sets node attributes may be used to measure node health. The most common ways are resource agents or separate daemons.

Pacemaker provides examples that can be used directly or as a basis for custom code. The `ocf:pacemaker:HealthCPU` and `ocf:pacemaker:HealthSMART` resource agents set node health attributes based on CPU and disk parameters. The `ipmiservicelogd` daemon sets node health attributes based on IPMI values (the `ocf:pacemaker:SystemHealth` resource agent can be used to manage the daemon as a cluster resource).

In order to take advantage of this feature - firstly add the resource to your cluster, preferably as a cloned resource to constantly measure health on all nodes:

```
<clone id="resHealthIOWait-clone">
  <primitive class="ocf" id="HealthIOWait" provider="pacemaker" type="HealthIOWait">
    <instance_attributes id="resHealthIOWait-instance_attributes">
      <nvpair id="resHealthIOWait-instance_attributes-red_limit" name="red_limit">
      <nvpair id="resHealthIOWait-instance_attributes-yellow_limit" name="yellow_limit">
    </instance_attributes>
    <operations>
      <op id="resHealthIOWait-monitor-interval-5" interval="5" name="monitor" timeout="5s">
      <op id="resHealthIOWait-start-interval-0s" interval="0s" name="start" timeout="0s">
      <op id="resHealthIOWait-stop-interval-0s" interval="0s" name="stop" timeout="0s">
    </operations>
  </primitive>
</clone>
```

This way `attrd_updater` will set proper status for each node running this resource. Any attribute matching `"#health-[a-zA-z]+"` will force cluster to migrate all resources from unhealthy node and place it on other nodes according to all constraints defined in your cluster.

When the node is no longer faulty you can force the cluster to restart the cloned resource on faulty node and make it available to take resources, in this case since we are using `HealthIOWait` provider:

```
# attrd_updater -n "#health-iowait" -U "green" --node="<nodename>" -d "60s"
```

Reloading Services After a Definition Change

The cluster automatically detects changes to the definition of services it manages. The normal response is to stop the service (using the old definition) and start it again (with the new definition). This works well, but some services are smarter and can be told to use a new set of options without restarting.

To take advantage of this capability, the resource agent must:

1. Accept the `reload` operation and perform any required actions. *The actions here depend completely on your application!*

Example 9.8. The DRBD agent's logic for supporting `reload`

```
case $1 in
```

```

start)
    drbd_start
    ;;
stop)
    drbd_stop
    ;;
reload)
    drbd_reload
    ;;
monitor)
    drbd_monitor
    ;;
*)
    drbd_usage
    exit $OCF_ERR_UNIMPLEMENTED
    ;;
esac
exit $?

```

2. Advertise the reload operation in the actions section of its metadata

Example 9.9. The DRBD Agent Advertising Support for the reload Operation

```

<?xml version="1.0"?>
  <!DOCTYPE resource-agent SYSTEM "ra-api-1.dtd">
  <resource-agent name="drbd">
    <version>1.1</version>

    <longdesc>
      Master/Slave OCF Resource Agent for DRBD
    </longdesc>

    ...

    <actions>
      <action name="start"    timeout="240" />
      <action name="reload"   timeout="240" />
      <action name="promote"  timeout="90" />
      <action name="demote"   timeout="90" />
      <action name="notify"   timeout="90" />
      <action name="stop"     timeout="100" />
      <action name="meta-data" timeout="5" />
      <action name="validate-all" timeout="30" />
    </actions>
  </resource-agent>

```

3. Advertise one or more parameters that can take effect using reload.

Any parameter with the unique set to 0 is eligible to be used in this way.

Example 9.10. Parameter that can be changed using reload

```

<parameter name="drbdconf" unique="0">
  <longdesc>Full path to the drbd.conf file.</longdesc>

```



```
<shortdesc>Path to drbd.conf</shortdesc>
<content type="string" default="{OCF_RESKEY_drbdconf_default}"/>
</parameter>
```

Once these requirements are satisfied, the cluster will automatically know to reload the resource (instead of restarting) when a non-unique field changes.

Note

Metadata will not be re-read unless the resource needs to be started. This may mean that the resource will be restarted the first time, even though you changed a parameter with `unique=0`.

Note

If both a unique and non-unique field are changed simultaneously, the resource will still be restarted.

Chapter 10. Advanced Resource Types

Groups - A Syntactic Shortcut

One of the most common elements of a cluster is a set of resources that need to be located together, start sequentially, and stop in the reverse order. To simplify this configuration, we support the concept of groups.

Example 10.1. A group of two primitive resources

```
<group id="shortcut">
  <primitive id="Public-IP" class="ocf" type="IPAddr" provider="heartbeat">
    <instance_attributes id="params-public-ip">
      <nvpair id="public-ip-addr" name="ip" value="192.0.2.2"/>
    </instance_attributes>
  </primitive>
  <primitive id="Email" class="lsb" type="exim"/>
</group>
```

Although the example above contains only two resources, there is no limit to the number of resources a group can contain. The example is also sufficient to explain the fundamental properties of a group:

- Resources are started in the order they appear in (Public-IP first, then Email)
- Resources are stopped in the reverse order to which they appear in (Email first, then Public-IP)

If a resource in the group can't run anywhere, then nothing after that is allowed to run, too.

- If Public-IP can't run anywhere, neither can Email;
- but if Email can't run anywhere, this does not affect Public-IP in any way

The group above is logically equivalent to writing:

Example 10.2. How the cluster sees a group resource

```
<configuration>
  <resources>
    <primitive id="Public-IP" class="ocf" type="IPAddr" provider="heartbeat">
      <instance_attributes id="params-public-ip">
        <nvpair id="public-ip-addr" name="ip" value="192.0.2.2"/>
      </instance_attributes>
    </primitive>
    <primitive id="Email" class="lsb" type="exim"/>
  </resources>
  <constraints>
    <rsc_colocation id="xxx" rsc="Email" with-rsc="Public-IP" score="INFINITY"/>
    <rsc_order id="yyy" first="Public-IP" then="Email"/>
  </constraints>
</configuration>
```

Obviously as the group grows bigger, the reduced configuration effort can become significant.

Another (typical) example of a group is a DRBD volume, the filesystem mount, an IP address, and an application that uses them.

Group Properties

Table 10.1. Properties of a Group Resource

Field	Description
id	A unique name for the group

Group Options

Groups inherit the `priority`, `target-role`, and `is-managed` properties from primitive resources. See the section called “Resource Options” for information about those properties.

Group Instance Attributes

Groups have no instance attributes. However, any that are set for the group object will be inherited by the group’s children.

Group Contents

Groups may only contain a collection of cluster resources (see the section called “Resource Properties”). To refer to a child of a group resource, just use the child’s `id` instead of the group’s.

Group Constraints

Although it is possible to reference a group’s children in constraints, it is usually preferable to reference the group itself.

Example 10.3. Some constraints involving groups

```
<constraints>
  <rsc_location id="group-prefers-node1" rsc="shortcut" node="node1" score="500"
  <rsc_colocation id="webserver-with-group" rsc="Webserver" with-rsc="shortcut"/>
  <rsc_order id="start-group-then-webserver" first="Webserver" then="shortcut"/>
</constraints>
```

Group Stickiness

Stickiness, the measure of how much a resource wants to stay where it is, is additive in groups. Every active resource of the group will contribute its stickiness value to the group’s total. So if the default `resource-stickiness` is 100, and a group has seven members, five of which are active, then the group as a whole will prefer its current location with a score of 500.

Clones - Resources That Can Have Multiple Active Instances

Clone resources are resources that can have more than one copy active at the same time. This allows you, for example, to run a copy of a daemon on every node. You can clone any primitive or group resource.¹

Anonymous versus Unique Clones

A clone resource is configured to be either *anonymous* or *globally unique*.

Anonymous clones are the simplest. These behave completely identically everywhere they are running. Because of this, there can be only one instance of an anonymous clone active per node.

The instances of globally unique clones are distinct entities. All instances are launched identically, but one instance of the clone is not identical to any other instance, whether running on the same node or a different node. As an example, a cloned IP address can use special kernel functionality such that each instance handles a subset of requests for the same IP address.

Promotable clones

If a clone is *promotable*, its instances can perform a special role that Pacemaker will manage via the `promote` and `demote` actions of the resource agent.

Services that support such a special role have various terms for the special role and the default role: primary and secondary, master and replica, controller and worker, etc. Pacemaker uses the terms *master* and *slave*,² but is agnostic to what the service calls them or what they do.

All that Pacemaker cares about is that an instance comes up in the default role when started, and the resource agent supports the `promote` and `demote` actions to manage entering and exiting the special role.

Clone Properties

Table 10.2. Properties of a Clone Resource

Field	Description
<code>id</code>	A unique name for the clone

Clone Options

Options inherited from primitive resources: `priority`, `target-role`, `is-managed`

Table 10.3. Clone-specific configuration options

Field	Default	Description
<code>globally-unique</code>	false	If true, each clone instance performs a distinct function
<code>clone-max</code>	number of nodes in cluster	The maximum number of clone instances that can be started across the entire cluster

¹ Of course, the service must support running multiple instances.

² These are historical terms that will eventually be replaced, but the extensive use of them and the need for backward compatibility makes it a long process. You may see examples using a `master` tag instead of a `clone` tag with the `promotable` meta-attribute set to true; the `master` tag is supported, but deprecated, and will be removed in a future version. You may also see such services referred to as *multi-state* or *stateful*; these means the same thing as *promotable*.

Field	Default	Description
clone-node-max	1	If globally-unique is true, the maximum number of clone instances that can be started on a single node
clone-min	0	Require at least this number of clone instances to be runnable before allowing resources depending on the clone to be runnable. A value of 0 means require all clone instances to be runnable.
notify	false	Call the resource agent's notify action for all active instances, before and after starting or stopping any clone instance. The resource agent must support this action. Allowed values: false, true
ordered	false	If true, clone instances must be started sequentially instead of in parallel. Allowed values: false, true
interleave	false	When this clone is ordered relative to another clone, if this option is false (the default), the ordering is relative to <i>all</i> instances of the other clone, whereas if this option is true, the ordering is relative only to instances on the same node. Allowed values: false, true
promotable	false	If true, clone instances can perform a special role that Pacemaker will manage via the resource agent's promote and demote actions. The resource agent must support these actions. Allowed values: false, true
promoted-max	1	If promotable is true, the number of instances that can be promoted at one time across the entire cluster
promoted-node-max	1	If promotable is true and globally-unique is false, the number of clone instances can be promoted at one time on a single node

For backward compatibility, master-max and master-node-max are accepted as aliases for promoted-max and promoted-node-max, but are deprecated since 2.0.0, and support for them will be removed in a future version.

Clone Contents

Clones must contain exactly one primitive or group resource.

Example 10.4. A clone that runs a web server on all nodes

```
<clone id="apache-clone">
  <primitive id="apache" class="lsb" type="apache">
    <operations>
      <op id="apache-monitor" name="monitor" interval="30"/>
    </operations>
  </primitive>
</clone>
```

Warning

You should never reference the name of a clone's child (the primitive or group resource being cloned). If you think you need to do this, you probably need to re-evaluate your design.

Clone Instance Attributes

Clones have no instance attributes; however, any that are set here will be inherited by the clone's child.

Clone Constraints

In most cases, a clone will have a single instance on each active cluster node. If this is not the case, you can indicate which nodes the cluster should preferentially assign copies to with resource location constraints. These constraints are written no differently from those for primitive resources except that the clone's `id` is used.

Example 10.5. Some constraints involving clones

```
<constraints>
  <rsc_location id="clone-prefers-node1" rsc="apache-clone" node="node1" score="1" />
  <rsc_colocation id="stats-with-clone" rsc="apache-stats" with="apache-clone" />
  <rsc_order id="start-clone-then-stats" first="apache-clone" then="apache-stats" />
</constraints>
```

Ordering constraints behave slightly differently for clones. In the example above, `apache-stats` will wait until all copies of `apache-clone` that need to be started have done so before being started itself. Only if *no* copies can be started will `apache-stats` be prevented from being active. Additionally, the clone will wait for `apache-stats` to be stopped before stopping itself.

Colocation of a primitive or group resource with a clone means that the resource can run on any node with an active instance of the clone. The cluster will choose an instance based on where the clone is running and the resource's own location preferences.

Colocation between clones is also possible. If one clone A is colocated with another clone B, the set of allowed locations for A is limited to nodes on which B is (or will be) active. Placement is then performed normally.

Promotable Clone Constraints

For promotable clone resources, the `first-action` and/or `then-action` fields for ordering constraints may be set to `promote` or `demote` to constrain the master role, and colocation constraints may contain `rsc-role` and/or `with-rsc-role` fields.

Table 10.4. Additional colocation constraint options for promotable clone resources

Field	Default	Description
<code>rsc-role</code>	Started	An additional attribute of colocation constraints that specifies the role that <code>rsc</code> must be in. Allowed values: Started, Master, Slave.
<code>with-rsc-role</code>	Started	An additional attribute of colocation constraints that specifies the role that <code>with-rsc</code> must be in. Allowed values: Started, Master, Slave.

Example 10.6. Constraints involving promotable clone resources

```
<constraints>
```

```

<rsc_location id="db-prefers-node1" rsc="database" node="node1" score="500"/>
<rsc_colocation id="backup-with-db-slave" rsc="backup"
  with-rsc="database" with-rsc-role="Slave"/>
<rsc_colocation id="myapp-with-db-master" rsc="myApp"
  with-rsc="database" with-rsc-role="Master"/>
<rsc_order id="start-db-before-backup" first="database" then="backup"/>
<rsc_order id="promote-db-then-app" first="database" first-action="promote"
  then="myApp" then-action="start"/>
</constraints>

```

In the example above, myApp will wait until one of the database copies has been started and promoted to master before being started itself on the same node. Only if no copies can be promoted will myApp be prevented from being active. Additionally, the cluster will wait for myApp to be stopped before demoting the database.

Colocation of a primitive or group resource with a promotable clone resource means that it can run on any node with an active instance of the promotable clone resource that has the specified role (master or slave). In the example above, the cluster will choose a location based on where database is running as a master, and if there are multiple master instances it will also factor in myApp's own location preferences when deciding which location to choose.

Colocation with regular clones and other promotable clone resources is also possible. In such cases, the set of allowed locations for the rsc clone is (after role filtering) limited to nodes on which the with-rsc promotable clone resource is (or will be) in the specified role. Placement is then performed as normal.

Using Promotable Clone Resources in Colocation Sets

Table 10.5. Additional colocation set options relevant to promotable clone resources

Field	Default	Description
role	Started	The role that <i>all members</i> of the set must be in. Allowed values: Started, Master, Slave.

In the following example B's master must be located on the same node as A's master. Additionally resources C and D must be located on the same node as A's and B's masters.

Example 10.7. Colocate C and D with A's and B's master instances

```

<constraints>
  <rsc_colocation id="coloc-1" score="INFINITY" >
    <resource_set id="colocated-set-example-1" sequential="true" role="Master">
      <resource_ref id="A"/>
      <resource_ref id="B"/>
    </resource_set>
    <resource_set id="colocated-set-example-2" sequential="true">
      <resource_ref id="C"/>
      <resource_ref id="D"/>
    </resource_set>
  </rsc_colocation>
</constraints>

```

Using Promotable Clone Resources in Ordered Sets

Table 10.6. Additional ordered set options relevant to promotable clone resources

Field	Default	Description
action	value of first-action	An additional attribute of ordering constraint sets that specifies the action that applies to <i>all members</i> of the set. Allowed values: start, stop, promote, demote.

Example 10.8. Start C and D after first promoting A and B

```
<constraints>
  <rsc_order id="order-1" score="INFINITY" >
    <resource_set id="ordered-set-1" sequential="true" action="promote">
      <resource_ref id="A"/>
      <resource_ref id="B"/>
    </resource_set>
    <resource_set id="ordered-set-2" sequential="true" action="start">
      <resource_ref id="C"/>
      <resource_ref id="D"/>
    </resource_set>
  </rsc_order>
</constraints>
```

In the above example, B cannot be promoted to a master role until A has been promoted. Additionally, resources C and D must wait until A and B have been promoted before they can start.

Clone Stickiness

To achieve a stable allocation pattern, clones are slightly sticky by default. If no value for `resource-stickiness` is provided, the clone will use a value of 1. Being a small value, it causes minimal disturbance to the score calculations of other resources but is enough to prevent Pacemaker from needlessly moving copies around the cluster.

Note

For globally unique clones, this may result in multiple instances of the clone staying on a single node, even after another eligible node becomes active (for example, after being put into standby mode then made active again). If you do not want this behavior, specify a `resource-stickiness` of 0 for the clone temporarily and let the cluster adjust, then set it back to 1 if you want the default behavior to apply again.

Important

If `resource-stickiness` is set in the `rsc_defaults` section, it will apply to clone instances as well. This means an explicit `resource-stickiness` of 0 in `rsc_defaults` works differently from the implicit default used when `resource-stickiness` is not specified.

Clone Resource Agent Requirements

Any resource can be used as an anonymous clone, as it requires no additional support from the resource agent. Whether it makes sense to do so depends on your resource and its resource agent.

Resource Agent Requirements for Globally Unique Clones

Globally unique clones require additional support in the resource agent. In particular, it must only respond with `{OCF_SUCCESS}` if the node has that exact instance active. All other probes for instances of the clone should result in `{OCF_NOT_RUNNING}` (or one of the other OCF error codes if they are failed).

Individual instances of a clone are identified by appending a colon and a numerical offset, e.g. `apache : 2`.

Resource agents can find out how many copies there are by examining the `OCF_RESKEY_CRM_meta_clone_max` environment variable and which instance it is by examining `OCF_RESKEY_CRM_meta_clone`.

The resource agent must not make any assumptions (based on `OCF_RESKEY_CRM_meta_clone`) about which numerical instances are active. In particular, the list of active copies will not always be an unbroken sequence, nor always start at 0.

Resource Agent Requirements for Promotable Clones

Promotable clone resources require two extra actions, `demote` and `promote`, which are responsible for changing the state of the resource. Like `start` and `stop`, they should return `{OCF_SUCCESS}` if they completed successfully or a relevant error code if they did not.

The states can mean whatever you wish, but when the resource is started, it must come up in the mode called `slave`. From there the cluster will decide which instances to promote to `master`.

In addition to the clone requirements for monitor actions, agents must also *accurately* report which state they are in. The cluster relies on the agent to report its status (including role) accurately and does not indicate to the agent what role it currently believes it to be in.

Table 10.7. Role implications of OCF return codes

Monitor Return Code	Description
OCF_NOT_RUNNING	Stopped
OCF_SUCCESS	Running (Slave)
OCF_RUNNING_MASTER	Running (Master)
OCF_FAILED_MASTER	Failed (Master)
Other	Failed (Slave)

Clone Notifications

If the clone has the `notify` meta-attribute set to `true`, and the resource agent supports the `notify` action, Pacemaker will call the action when appropriate, passing a number of extra variables which, when combined with additional context, can be used to calculate the current state of the cluster and what is about to happen to it.

Table 10.8. Environment variables supplied with Clone notify actions

Variable	Description
OCF_RESKEY_CRM_meta_notify_type	Allowed values: <code>pre</code> , <code>post</code>

Variable	Description
OCF_RESKEY_CRM_meta_notify_operation	Allowed values: start, stop
OCF_RESKEY_CRM_meta_notify_start_resource	Resources to be started
OCF_RESKEY_CRM_meta_notify_stop_resource	Resources to be stopped
OCF_RESKEY_CRM_meta_notify_active_resource	Resources that are running
OCF_RESKEY_CRM_meta_notify_inactive_resource	Resources that are not running
OCF_RESKEY_CRM_meta_notify_start_uname	Nodes on which resources will be started
OCF_RESKEY_CRM_meta_notify_stop_uname	Nodes on which resources will be stopped
OCF_RESKEY_CRM_meta_notify_active_uname	Nodes on which resources are running

The variables come in pairs, such as `OCF_RESKEY_CRM_meta_notify_start_resource` and `OCF_RESKEY_CRM_meta_notify_start_uname`, and should be treated as an array of whitespace-separated elements.

`OCF_RESKEY_CRM_meta_notify_inactive_resource` is an exception, as the matching `uname` variable does not exist since inactive resources are not running on any node.

Thus, in order to indicate that `clone:0` will be started on `sles-1`, `clone:2` will be started on `sles-3`, and `clone:3` will be started on `sles-2`, the cluster would set:

Example 10.9. Notification variables

```
OCF_RESKEY_CRM_meta_notify_start_resource="clone:0 clone:2 clone:3"
OCF_RESKEY_CRM_meta_notify_start_uname="sles-1 sles-3 sles-2"
```

Note

Pacemaker will log but otherwise ignore failures of notify actions.

Interpretation of Notification Variables

Pre-notification (stop):

- Active resources: `$OCF_RESKEY_CRM_meta_notify_active_resource`
- Inactive resources: `$OCF_RESKEY_CRM_meta_notify_inactive_resource`
- Resources to be started: `$OCF_RESKEY_CRM_meta_notify_start_resource`
- Resources to be stopped: `$OCF_RESKEY_CRM_meta_notify_stop_resource`

Post-notification (stop) / Pre-notification (start):

- Active resources
 - `$OCF_RESKEY_CRM_meta_notify_active_resource`

- minus `$OCF_RESKEY_CRM_meta_notify_stop_resource`
- Inactive resources
 - `$OCF_RESKEY_CRM_meta_notify_inactive_resource`
 - plus `$OCF_RESKEY_CRM_meta_notify_stop_resource`
- Resources that were started: `$OCF_RESKEY_CRM_meta_notify_start_resource`
- Resources that were stopped: `$OCF_RESKEY_CRM_meta_notify_stop_resource`

Post-notification (start):

- Active resources:
 - `$OCF_RESKEY_CRM_meta_notify_active_resource`
 - minus `$OCF_RESKEY_CRM_meta_notify_stop_resource`
 - plus `$OCF_RESKEY_CRM_meta_notify_start_resource`
- Inactive resources:
 - `$OCF_RESKEY_CRM_meta_notify_inactive_resource`
 - plus `$OCF_RESKEY_CRM_meta_notify_stop_resource`
 - minus `$OCF_RESKEY_CRM_meta_notify_start_resource`
- Resources that were started: `$OCF_RESKEY_CRM_meta_notify_start_resource`
- Resources that were stopped: `$OCF_RESKEY_CRM_meta_notify_stop_resource`

Extra Notifications for Promotable Clones

Table 10.9. Extra environment variables supplied for promotable clones

Variable	Description
<code>OCF_RESKEY_CRM_meta_notify_master_resource</code>	Resources that are running in Master mode
<code>OCF_RESKEY_CRM_meta_notify_slave_resource</code>	Resources that are running in Slave mode
<code>OCF_RESKEY_CRM_meta_notify_promote_resource</code>	Resources to be promoted
<code>OCF_RESKEY_CRM_meta_notify_demote_resource</code>	Resources to be demoted
<code>OCF_RESKEY_CRM_meta_notify_promote_uname</code>	Nodes on which resources will be promoted
<code>OCF_RESKEY_CRM_meta_notify_demote_uname</code>	Nodes on which resources will be demoted
<code>OCF_RESKEY_CRM_meta_notify_master_uname</code>	Nodes on which resources are running in Master mode
<code>OCF_RESKEY_CRM_meta_notify_slave_uname</code>	Nodes on which resources are running in Slave mode

Interpretation of Promotable Notification Variables

Pre-notification (demote):

- Active resources: `$OCF_RESKEY_CRM_meta_notify_active_resource`
- Master resources: `$OCF_RESKEY_CRM_meta_notify_master_resource`
- Slave resources: `$OCF_RESKEY_CRM_meta_notify_slave_resource`
- Inactive resources: `$OCF_RESKEY_CRM_meta_notify_inactive_resource`
- Resources to be started: `$OCF_RESKEY_CRM_meta_notify_start_resource`
- Resources to be promoted: `$OCF_RESKEY_CRM_meta_notify_promote_resource`
- Resources to be demoted: `$OCF_RESKEY_CRM_meta_notify_demote_resource`
- Resources to be stopped: `$OCF_RESKEY_CRM_meta_notify_stop_resource`

Post-notification (demote) / Pre-notification (stop):

- Active resources: `$OCF_RESKEY_CRM_meta_notify_active_resource`
- Master resources:
 - `$OCF_RESKEY_CRM_meta_notify_master_resource`
 - minus `$OCF_RESKEY_CRM_meta_notify_demote_resource`
- Slave resources: `$OCF_RESKEY_CRM_meta_notify_slave_resource`
- Inactive resources: `$OCF_RESKEY_CRM_meta_notify_inactive_resource`
- Resources to be started: `$OCF_RESKEY_CRM_meta_notify_start_resource`
- Resources to be promoted: `$OCF_RESKEY_CRM_meta_notify_promote_resource`
- Resources to be demoted: `$OCF_RESKEY_CRM_meta_notify_demote_resource`
- Resources to be stopped: `$OCF_RESKEY_CRM_meta_notify_stop_resource`
- Resources that were demoted: `$OCF_RESKEY_CRM_meta_notify_demote_resource`

Post-notification (stop) / Pre-notification (start)

- Active resources:
 - `$OCF_RESKEY_CRM_meta_notify_active_resource`
 - minus `$OCF_RESKEY_CRM_meta_notify_stop_resource`
- Master resources:
 - `$OCF_RESKEY_CRM_meta_notify_master_resource`
 - minus `$OCF_RESKEY_CRM_meta_notify_demote_resource`

- Slave resources:
 - `$OCF_RESKEY_CRM_meta_notify_slave_resource`
 - minus `$OCF_RESKEY_CRM_meta_notify_stop_resource`
- Inactive resources:
 - `$OCF_RESKEY_CRM_meta_notify_inactive_resource`
 - plus `$OCF_RESKEY_CRM_meta_notify_stop_resource`
- Resources to be started: `$OCF_RESKEY_CRM_meta_notify_start_resource`
- Resources to be promoted: `$OCF_RESKEY_CRM_meta_notify_promote_resource`
- Resources to be demoted: `$OCF_RESKEY_CRM_meta_notify_demote_resource`
- Resources to be stopped: `$OCF_RESKEY_CRM_meta_notify_stop_resource`
- Resources that were demoted: `$OCF_RESKEY_CRM_meta_notify_demote_resource`
- Resources that were stopped: `$OCF_RESKEY_CRM_meta_notify_stop_resource`

Post-notification (start) / Pre-notification (promote)

- Active resources:
 - `$OCF_RESKEY_CRM_meta_notify_active_resource`
 - minus `$OCF_RESKEY_CRM_meta_notify_stop_resource`
 - plus `$OCF_RESKEY_CRM_meta_notify_start_resource`
- Master resources:
 - `$OCF_RESKEY_CRM_meta_notify_master_resource`
 - minus `$OCF_RESKEY_CRM_meta_notify_demote_resource`
- Slave resources:
 - `$OCF_RESKEY_CRM_meta_notify_slave_resource`
 - minus `$OCF_RESKEY_CRM_meta_notify_stop_resource`
 - plus `$OCF_RESKEY_CRM_meta_notify_start_resource`
- Inactive resources:
 - `$OCF_RESKEY_CRM_meta_notify_inactive_resource`
 - plus `$OCF_RESKEY_CRM_meta_notify_stop_resource`
 - minus `$OCF_RESKEY_CRM_meta_notify_start_resource`
- Resources to be started: `$OCF_RESKEY_CRM_meta_notify_start_resource`
- Resources to be promoted: `$OCF_RESKEY_CRM_meta_notify_promote_resource`

- Resources to be demoted: `$OCF_RESKEY_CRM_meta_notify_demote_resource`
- Resources to be stopped: `$OCF_RESKEY_CRM_meta_notify_stop_resource`
- Resources that were started: `$OCF_RESKEY_CRM_meta_notify_start_resource`
- Resources that were demoted: `$OCF_RESKEY_CRM_meta_notify_demote_resource`
- Resources that were stopped: `$OCF_RESKEY_CRM_meta_notify_stop_resource`

Post-notification (promote)

- Active resources:
 - `$OCF_RESKEY_CRM_meta_notify_active_resource`
 - minus `$OCF_RESKEY_CRM_meta_notify_stop_resource`
 - plus `$OCF_RESKEY_CRM_meta_notify_start_resource`
- Master resources:
 - `$OCF_RESKEY_CRM_meta_notify_master_resource`
 - minus `$OCF_RESKEY_CRM_meta_notify_demote_resource`
 - plus `$OCF_RESKEY_CRM_meta_notify_promote_resource`
- Slave resources:
 - `$OCF_RESKEY_CRM_meta_notify_slave_resource`
 - minus `$OCF_RESKEY_CRM_meta_notify_stop_resource`
 - plus `$OCF_RESKEY_CRM_meta_notify_start_resource`
 - minus `$OCF_RESKEY_CRM_meta_notify_promote_resource`
- Inactive resources:
 - `$OCF_RESKEY_CRM_meta_notify_inactive_resource`
 - plus `$OCF_RESKEY_CRM_meta_notify_stop_resource`
 - minus `$OCF_RESKEY_CRM_meta_notify_start_resource`
- Resources to be started: `$OCF_RESKEY_CRM_meta_notify_start_resource`
- Resources to be promoted: `$OCF_RESKEY_CRM_meta_notify_promote_resource`
- Resources to be demoted: `$OCF_RESKEY_CRM_meta_notify_demote_resource`
- Resources to be stopped: `$OCF_RESKEY_CRM_meta_notify_stop_resource`
- Resources that were started: `$OCF_RESKEY_CRM_meta_notify_start_resource`
- Resources that were promoted: `$OCF_RESKEY_CRM_meta_notify_promote_resource`

- Resources that were demoted: `$OCF_RESKEY_CRM_meta_notify_demote_resource`
- Resources that were stopped: `$OCF_RESKEY_CRM_meta_notify_stop_resource`

Monitoring Promotable Clone Resources

The usual monitor actions are insufficient to monitor a promotable clone resource, because Pacemaker needs to verify not only that the resource is active, but also that its actual role matches its intended one.

Define two monitoring actions: the usual one will cover the slave role, and an additional one with `role="master"` will cover the master role.

Example 10.10. Monitoring both states of a promotable clone resource

```
<clone id="myMasterRsc">
  <meta_attributes id="myMasterRsc-meta">
    <nvpair name="promotable" value="true"/>
  </meta_attributes>
  <primitive id="myRsc" class="ocf" type="myApp" provider="myCorp">
    <operations>
      <op id="public-ip-slave-check" name="monitor" interval="60"/>
      <op id="public-ip-master-check" name="monitor" interval="61" role="Master"/>
    </operations>
  </primitive>
</clone>
```

Important

It is crucial that *every* monitor operation has a different interval! Pacemaker currently differentiates between operations only by resource and interval; so if (for example) a promotable clone resource had the same monitor interval for both roles, Pacemaker would ignore the role when checking the status — which would cause unexpected return codes, and therefore unnecessary complications.

Determining Which Instance is Promoted

Pacemaker can choose a promotable clone instance to be promoted in one of two ways:

- Promotion scores: These are node attributes set via the `crm_master` utility, which generally would be called by the resource agent's start action if it supports promotable clones. This tool automatically detects both the resource and host, and should be used to set a preference for being promoted. Based on this, `promoted-max`, and `promoted-node-max`, the instance(s) with the highest preference will be promoted.
- Constraints: Location constraints can indicate which nodes are most preferred as masters.

Example 10.11. Explicitly preferring node1 to be promoted to master

```
<rsc_location id="master-location" rsc="myMasterRsc">
  <rule id="master-rule" score="100" role="Master">
    <expression id="master-exp" attribute="#uname" operation="eq" value="node1"/>
  </rule>
</rsc_location>
```

Bundles - Isolated Environments

Pacemaker supports a special syntax for launching a container [https://en.wikipedia.org/wiki/Operating-system-level_virtualization] with any infrastructure it requires: the *bundle*.

Pacemaker bundles support Docker [<https://www.docker.com/>], podman [<https://podman.io/>], and rkt [<https://coreos.com/rkt/>] container technologies.³

Example 10.12. A bundle for a containerized web server

```
<bundle id="httpd-bundle">
  <podman image="pcmk:http" replicas="3"/>
  <network ip-range-start="192.168.122.131"
    host-netmask="24"
    host-interface="eth0">
    <port-mapping id="httpd-port" port="80"/>
  </network>
  <storage>
    <storage-mapping id="httpd-syslog"
      source-dir="/dev/log"
      target-dir="/dev/log"
      options="rw"/>
    <storage-mapping id="httpd-root"
      source-dir="/srv/html"
      target-dir="/var/www/html"
      options="rw,Z"/>
    <storage-mapping id="httpd-logs"
      source-dir-root="/var/log/pacemaker/bundles"
      target-dir="/etc/httpd/logs"
      options="rw,Z"/>
  </storage>
  <primitive class="ocf" id="httpd" provider="heartbeat" type="apache"/>
</bundle>
```

Bundle Prerequisites

Before configuring a bundle in Pacemaker, the user must install the appropriate container launch technology (Docker, podman, or rkt), and supply a fully configured container image, on every node allowed to run the bundle.

Pacemaker will create an implicit resource of type `ocf:heartbeat:docker`, `ocf:heartbeat:podman`, or `ocf:heartbeat:rkt` to manage a bundle's container. The user must ensure that the appropriate resource agent is installed on every node allowed to run the bundle.

Bundle Properties

³ Docker is a trademark of Docker, Inc. No endorsement by or association with Docker, Inc. is implied.

Table 10.10. XML Attributes of a bundle Element

Attribute	Description
id	A unique name for the bundle (required)
description	Arbitrary text (not used by Pacemaker)

A bundle must contain exactly one `docker`, `podman`, or `rkt` element.

Bundle Container Properties

Table 10.11. XML attributes of a `docker`, `podman`, or `rkt` Element

Attribute	Default	Description
image		Container image tag (required)
replicas	Value of <code>promoted-max</code> if that is positive, else 1	A positive integer specifying the number of container instances to launch
replicas-per-host	1	A positive integer specifying the number of container instances allowed to run on a single node
promoted-max	0	A non-negative integer that, if positive, indicates that the containerized service should be treated as a promotable service, with this many replicas allowed to run the service in the master role
network		If specified, this will be passed to the <code>docker run</code> , <code>podman run</code> , or <code>rkt run</code> command as the network setting for the container.
run-command	<code>/usr/sbin/pacemaker-remoted</code> if bundle contains a primitive, otherwise none	This command will be run inside the container when launching it ("PID 1"). If the bundle contains a primitive, this command <i>must</i> start <code>pacemaker-remoted</code> (but could, for example, be a script that does other stuff, too).
options		Extra command-line options to pass to the <code>docker run</code> , <code>podman run</code> , or <code>rkt run</code> command

Note

Considerations when using cluster configurations or container images from Pacemaker 1.1:

- If the container image has a pre-2.0.0 version of Pacemaker, set `run-command` to `/usr/sbin/pacemaker_remoted` (note the underbar instead of dash).
- `masters` is accepted as an alias for `promoted-max`, but is deprecated since 2.0.0, and support for it will be removed in a future version.

Bundle Network Properties

A bundle may optionally contain one `<network>` element.

Table 10.12. XML attributes of a network Element

Attribute	Default	Description
<code>add-host</code>	TRUE	If TRUE, and <code>ip-range-start</code> is used, Pacemaker will automatically ensure that <code>/etc/hosts</code> inside the containers has entries for each replica name and its assigned IP.
<code>ip-range-start</code>		If specified, Pacemaker will create an implicit <code>ocf:heartbeat:IPaddr2</code> resource for each container instance, starting with this IP address, using up to <code>replicas</code> sequential addresses. These addresses can be used from the host's network to reach the service inside the container, though it is not visible within the container itself. Only IPv4 addresses are currently supported.
<code>host-netmask</code>	32	If <code>ip-range-start</code> is specified, the IP addresses are created with this CIDR netmask (as a number of bits).
<code>host-interface</code>		If <code>ip-range-start</code> is specified, the IP addresses are created on this host interface (by default, it will be determined from the IP address).
<code>control-port</code>	3121	If the bundle contains a primitive, the cluster will use this integer TCP port for communication with Pacemaker Remote inside the container. Changing this is useful when the container is unable to listen on the default port, for example, when the container uses the host's network rather than <code>ip-range-start</code> (in which case <code>replicas-per-host</code> must be 1), or when the bundle may run on a Pacemaker Remote node that is already listening on the default port. Any <code>PCMK_remote_port</code> environment variable set on the host or in the container is ignored for bundle connections.

Note

Replicas are named by the bundle id plus a dash and an integer counter starting with zero. For example, if a bundle named `httpd-bundle` has `replicas=2`, its containers will be named `httpd-bundle-0` and `httpd-bundle-1`.

Additionally, a network element may optionally contain one or more `port-mapping` elements.

Table 10.13. Attributes of a port-mapping Element

Attribute	Default	Description
<code>id</code>		A unique name for the port mapping (required)
<code>port</code>		If this is specified, connections to this TCP port number on the host network (on the container's

Attribute	Default	Description
		assigned IP address, if <code>ip-range-start</code> is specified) will be forwarded to the container network. Exactly one of <code>port</code> or <code>range</code> must be specified in a <code>port-mapping</code> .
<code>internal-port</code>	value of <code>port</code>	If <code>port</code> and this are specified, connections to <code>port</code> on the host's network will be forwarded to this port on the container network.
<code>range</code>		If this is specified, connections to these TCP port numbers (expressed as <i>first_port-last_port</i>) on the host network (on the container's assigned IP address, if <code>ip-range-start</code> is specified) will be forwarded to the same ports in the container network. Exactly one of <code>port</code> or <code>range</code> must be specified in a <code>port-mapping</code> .

Note

If the bundle contains a `primitive`, Pacemaker will automatically map the `control-port`, so it is not necessary to specify that port in a `port-mapping`.

Bundle Storage Properties

A bundle may optionally contain one storage element. A storage element has no properties of its own, but may contain one or more `storage-mapping` elements.

Table 10.14. Attributes of a storage-mapping Element

Attribute	Default	Description
<code>id</code>		A unique name for the storage mapping (required)
<code>source-dir</code>		The absolute path on the host's filesystem that will be mapped into the container. Exactly one of <code>source-dir</code> and <code>source-dir-root</code> must be specified in a <code>storage-mapping</code> .
<code>source-dir-root</code>		The start of a path on the host's filesystem that will be mapped into the container, using a different subdirectory on the host for each container instance. The subdirectory will be named the same as the replica name. Exactly one of <code>source-dir</code> and <code>source-dir-root</code> must be specified in a <code>storage-mapping</code> .
<code>target-dir</code>		The path name within the container where the host storage will be mapped (required)
<code>options</code>		A comma-separated list of file system mount options to use when mapping the storage

Note

Pacemaker does not define the behavior if the source directory does not already exist on the host. However, it is expected that the container technology and/or its resource agent will create the source directory in that case.

Note

If the bundle contains a `primitive`, Pacemaker will automatically map the equivalent of `source-dir=/etc/pacemaker/authkey` `target-dir=/etc/pacemaker/authkey` and `source-dir-root=/var/log/pacemaker/bundles` `target-dir=/var/log` into the container, so it is not necessary to specify those paths in a `storage-mapping`.

Important

The `PCMK_authkey_location` environment variable must not be set to anything other than the default of `/etc/pacemaker/authkey` on any node in the cluster.

Important

If SELinux is used in enforcing mode on the host, you must ensure the container is allowed to use any storage you mount into it. For Docker and podman bundles, adding "Z" to the mount options will create a container-specific label for the mount that allows the container access.

Bundle Primitive

A bundle may optionally contain one primitive resource. The primitive may have operations, instance attributes, and meta-attributes defined, as usual.

If a bundle contains a primitive resource, the container image must include the Pacemaker Remote daemon, and at least one of `ip-range-start` or `control-port` must be configured in the bundle. Pacemaker will create an implicit `ocf:pacemaker:remote` resource for the connection, launch Pacemaker Remote within the container, and monitor and manage the primitive resource via Pacemaker Remote.

If the bundle has more than one container instance (replica), the primitive resource will function as an implicit clone — a promotable clone if the bundle has `masters` greater than zero.

Note

If you want to pass environment variables to a bundle's Pacemaker Remote connection or primitive, you have two options:

- Environment variables whose value is the same regardless of the underlying host may be set using the container element's `options` attribute.
- If you want variables to have host-specific values, you can use the `storage-mapping` element to map a file on the host as `/etc/pacemaker/pcmk-init.env` in the container. Pacemaker Remote will parse this file as a shell-like format, with variables set as `NAME=VALUE`, ignoring blank lines and comments starting with `"#"`.

Important

When a bundle has a `primitive`, Pacemaker on all cluster nodes must be able to contact Pacemaker Remote inside the bundle's containers.

- The containers must have an accessible network (for example, `network` should not be set to `"none"` with a `primitive`).

- The default, using a distinct network space inside the container, works in combination with `ip-range-start`. Any firewall must allow access from all cluster nodes to the `control-port` on the container IPs.
- If the container shares the host's network space (for example, by setting `network` to "host"), a unique `control-port` should be specified for each bundle. Any firewall must allow access from all cluster nodes to the `control-port` on all cluster and remote node IPs.

Bundle Node Attributes

If the bundle has a `primitive`, the primitive's resource agent may want to set node attributes such as promotion scores. However, with containers, it is not apparent which node should get the attribute.

If the container uses shared storage that is the same no matter which node the container is hosted on, then it is appropriate to use the promotion score on the bundle node itself.

On the other hand, if the container uses storage exported from the underlying host, then it may be more appropriate to use the promotion score on the underlying host.

Since this depends on the particular situation, the `container-attribute-target` resource meta-attribute allows the user to specify which approach to use. If it is set to `host`, then user-defined node attributes will be checked on the underlying host. If it is anything else, the local node (in this case the bundle node) is used as usual.

This only applies to user-defined attributes; the cluster will always check the local node for cluster-defined attributes such as `#uname`.

If `container-attribute-target` is `host`, the cluster will pass additional environment variables to the primitive's resource agent that allow it to set node attributes appropriately: `CRM_meta_container_attribute_target` (identical to the meta-attribute value) and `CRM_meta_physical_host` (the name of the underlying host).

Note

When called by a resource agent, the `attrd_updater` and `crm_attribute` commands will automatically check those environment variables and set attributes appropriately.

Bundle Meta-Attributes

Any meta-attribute set on a bundle will be inherited by the bundle's primitive and any resources implicitly created by Pacemaker for the bundle.

This includes options such as `priority`, `target-role`, and `is-managed`. See the section called "Resource Options" for more information.

Limitations of Bundles

Restarting pacemaker while a bundle is unmanaged or the cluster is in maintenance mode may cause the bundle to fail.

Bundles may not be explicitly cloned or included in groups. This includes the bundle's primitive and any resources implicitly created by Pacemaker for the bundle. (If `replicas` is greater than 1, the bundle will behave like a clone implicitly.)

Bundles do not have instance attributes, utilization attributes, or operations, though a bundle's primitive may have them.

A bundle with a primitive can run on a Pacemaker Remote node only if the bundle uses a distinct `control-port`.

Chapter 11. Reusing Parts of the Configuration

Pacemaker provides multiple ways to simplify the configuration XML by reusing parts of it in multiple places.

Besides simplifying the XML, this also allows you to manipulate multiple configuration elements with a single reference.

Reusing Resource Definitions

If you want to create lots of resources with similar configurations, defining a *resource template* simplifies the task. Once defined, it can be referenced in primitives or in certain types of constraints.

Configuring Resources with Templates

The primitives referencing the template will inherit all meta-attributes, instance attributes, utilization attributes and operations defined in the template. And you can define specific attributes and operations for any of the primitives. If any of these are defined in both the template and the primitive, the values defined in the primitive will take precedence over the ones defined in the template.

Hence, resource templates help to reduce the amount of configuration work. If any changes are needed, they can be done to the template definition and will take effect globally in all resource definitions referencing that template.

Resource templates have a syntax similar to that of primitives.

Example 11.1. Resource template for a migratable Xen virtual machine

```
<template id="vm-template" class="ocf" provider="heartbeat" type="Xen">
  <meta_attributes id="vm-template-meta_attributes">
    <nvpair id="vm-template-meta_attributes-allow-migrate" name="allow-migrate" va
  </meta_attributes>
  <utilization id="vm-template-utilization">
    <nvpair id="vm-template-utilization-memory" name="memory" value="512"/>
  </utilization>
  <operations>
    <op id="vm-template-monitor-15s" interval="15s" name="monitor" timeout="60s"/>
    <op id="vm-template-start-0" interval="0" name="start" timeout="60s"/>
  </operations>
</template>
```

Once you define a resource template, you can use it in primitives by specifying the `template` property.

Example 11.2. Xen primitive resource using a resource template

```
<primitive id="vm1" template="vm-template">
  <instance_attributes id="vm1-instance_attributes">
    <nvpair id="vm1-instance_attributes-name" name="name" value="vm1"/>
    <nvpair id="vm1-instance_attributes-xmfile" name="xmfile" value="/etc/xen/shar
  </instance_attributes>
```

```
</primitive>
```

In the example above, the new primitive `vm1` will inherit everything from `vm-template`. For example, the equivalent of the above two examples would be:

Example 11.3. Equivalent Xen primitive resource not using a resource template

```
<primitive id="vm1" class="ocf" provider="heartbeat" type="Xen">
  <meta_attributes id="vm-template-meta_attributes">
    <nvpair id="vm-template-meta_attributes-allow-migrate" name="allow-migrate" value="false"/>
  </meta_attributes>
  <utilization id="vm-template-utilization">
    <nvpair id="vm-template-utilization-memory" name="memory" value="512"/>
  </utilization>
  <operations>
    <op id="vm-template-monitor-15s" interval="15s" name="monitor" timeout="60s"/>
    <op id="vm-template-start-0" interval="0" name="start" timeout="60s"/>
  </operations>
  <instance_attributes id="vm1-instance_attributes">
    <nvpair id="vm1-instance_attributes-name" name="name" value="vm1"/>
    <nvpair id="vm1-instance_attributes-xmfile" name="xmfile" value="/etc/xen/shared" />
  </instance_attributes>
</primitive>
```

If you want to overwrite some attributes or operations, add them to the particular primitive's definition.

Example 11.4. Xen resource overriding template values

```
<primitive id="vm2" template="vm-template">
  <meta_attributes id="vm2-meta_attributes">
    <nvpair id="vm2-meta_attributes-allow-migrate" name="allow-migrate" value="false"/>
  </meta_attributes>
  <utilization id="vm2-utilization">
    <nvpair id="vm2-utilization-memory" name="memory" value="1024"/>
  </utilization>
  <instance_attributes id="vm2-instance_attributes">
    <nvpair id="vm2-instance_attributes-name" name="name" value="vm2"/>
    <nvpair id="vm2-instance_attributes-xmfile" name="xmfile" value="/etc/xen/shared" />
  </instance_attributes>
  <operations>
    <op id="vm2-monitor-30s" interval="30s" name="monitor" timeout="120s"/>
    <op id="vm2-stop-0" interval="0" name="stop" timeout="60s"/>
  </operations>
</primitive>
```

In the example above, the new primitive `vm2` has special attribute values. Its monitor operation has a longer timeout and interval, and the primitive has an additional stop operation.

To see the resulting definition of a resource, run:

```
# crm_resource --query-xml --resource vm2
```

To see the raw definition of a resource in the CIB, run:

```
# crm_resource --query-xml-raw --resource vm2
```


Using Templates in Constraints

A resource template can be referenced in the following types of constraints:

- order constraints (see the section called “Specifying the Order in which Resources Should Start/Stop”)
- colocation constraints (see the section called “Placing Resources Relative to other Resources”)
- rsc_ticket constraints (for multi-site clusters as described in the section called “Configuring Ticket Dependencies”)

Resource templates referenced in constraints stand for all primitives which are derived from that template. This means, the constraint applies to all primitive resources referencing the resource template. Referencing resource templates in constraints is an alternative to resource sets and can simplify the cluster configuration considerably.

For example, given the example templates earlier in this section:

```
<rsc_colocation id="vm-template-colo-base-rsc" rsc="vm-template" rsc-role="Started"
```

would colocate all VMs with base-rsc and is the equivalent of the following constraint configuration:

```
<rsc_colocation id="vm-colo-base-rsc" score="INFINITY">
  <resource_set id="vm-colo-base-rsc-0" sequential="false" role="Started">
    <resource_ref id="vm1" />
    <resource_ref id="vm2" />
  </resource_set>
  <resource_set id="vm-colo-base-rsc-1">
    <resource_ref id="base-rsc" />
  </resource_set>
</rsc_colocation>
```

Note

In a colocation constraint, only one template may be referenced from either `rsc` or `with-rsc`; the other reference must be a regular resource.

Using Templates in Resource Sets

Resource templates can also be referenced in resource sets.

For example, given the example templates earlier in this section, then:

```
<rsc_order id="order1" score="INFINITY">
  <resource_set id="order1-0">
    <resource_ref id="base-rsc" />
    <resource_ref id="vm-template" />
    <resource_ref id="top-rsc" />
  </resource_set>
</rsc_order>
```

is the equivalent of the following constraint using a sequential resource set:

```
<rsc_order id="order1" score="INFINITY">
  <resource_set id="order1-0">
```

```
<resource_ref id="base-rsc"/>
<resource_ref id="vm1"/>
<resource_ref id="vm2"/>
<resource_ref id="top-rsc"/>
</resource_set>
</rsc_order>
```

Or, if the resources referencing the template can run in parallel, then:

```
<rsc_order id="order2" score="INFINITY">
  <resource_set id="order2-0">
    <resource_ref id="base-rsc"/>
  </resource_set>
  <resource_set id="order2-1" sequential="false">
    <resource_ref id="vm-template"/>
  </resource_set>
  <resource_set id="order2-2">
    <resource_ref id="top-rsc"/>
  </resource_set>
</rsc_order>
```

is the equivalent of the following constraint configuration:

```
<rsc_order id="order2" score="INFINITY">
  <resource_set id="order2-0">
    <resource_ref id="base-rsc"/>
  </resource_set>
  <resource_set id="order2-1" sequential="false">
    <resource_ref id="vm1"/>
    <resource_ref id="vm2"/>
  </resource_set>
  <resource_set id="order2-2">
    <resource_ref id="top-rsc"/>
  </resource_set>
</rsc_order>
```

Reusing Rules, Options and Sets of Operations

Sometimes a number of constraints need to use the same set of rules, and resources need to set the same options and parameters. To simplify this situation, you can refer to an existing object using an `id-ref` instead of an `id`.

So if for one resource you have

```
<rsc_location id="WebServer-connectivity" rsc="Webserver">
  <rule id="ping-prefer-rule" score-attribute="pingd" >
    <expression id="ping-prefer" attribute="pingd" operation="defined"/>
  </rule>
</rsc_location>
```

Then instead of duplicating the rule for all your other resources, you can instead specify:

Example 11.5. Referencing rules from other constraints

```
<rsc_location id="WebDB-connectivity" rsc="WebDB">
```

```
<rule id-ref="ping-prefer-rule"/>
</rsc_location>
```

Important

The cluster will insist that the `rule` exists somewhere. Attempting to add a reference to a non-existing rule will cause a validation failure, as will attempting to remove a rule that is referenced elsewhere.

The same principle applies for `meta_attributes` and `instance_attributes` as illustrated in the example below:

Example 11.6. Referencing attributes, options, and operations from other resources

```
<primitive id="mySpecialRsc" class="ocf" type="Special" provider="me">
  <instance_attributes id="mySpecialRsc-attrs" score="1" >
    <nvpair id="default-interface" name="interface" value="eth0"/>
    <nvpair id="default-port" name="port" value="9999"/>
  </instance_attributes>
  <meta_attributes id="mySpecialRsc-options">
    <nvpair id="failure-timeout" name="failure-timeout" value="5m"/>
    <nvpair id="migration-threshold" name="migration-threshold" value="1"/>
    <nvpair id="stickiness" name="resource-stickiness" value="0"/>
  </meta_attributes>
  <operations id="health-checks">
    <op id="health-check" name="monitor" interval="60s"/>
    <op id="health-check" name="monitor" interval="30min"/>
  </operations>
</primitive>
<primitive id="myOther1Rsc" class="ocf" type="Other" provider="me">
  <instance_attributes id-ref="mySpecialRsc-attrs"/>
  <meta_attributes id-ref="mySpecialRsc-options"/>
  <operations id-ref="health-checks"/>
</primitive>
```

`id-ref` can similarly be used with `resource_set` (in any constraint type), `nvpair`, and `operations`.

Tagging Configuration Elements

Pacemaker allows you to *tag* any configuration element that has an XML ID.

The main purpose of tagging is to support higher-level user interface tools; Pacemaker itself only uses tags within constraints. Therefore, what you can do with tags mostly depends on the tools you use.

Configuring Tags

A tag is simply a named list of XML IDs.

Example 11.7. Tag referencing three resources

```
<tags>
  <tag id="all-vms">
```

```
<obj_ref id="vm1" />
<obj_ref id="vm2" />
<obj_ref id="vm3" />
</tag>
</tags>
```

What you can do with this new tag depends on what your higher-level tools support. For example, a tool might allow you to enable or disable all of the tagged resources at once, or show the status of just the tagged resources.

A single configuration element can be listed in any number of tags.

Using Tags in Constraints and Resource Sets

Pacemaker itself only uses tags in constraints. If you supply a tag name instead of a resource name in any constraint, the constraint will apply to all resources listed in that tag.

Example 11.8. Constraint using a tag

```
<rsc_order id="order1" first="storage" then="all-vms" kind="Mandatory" />
```

In the example above, assuming the `all-vms` tag is defined as in the previous example, the constraint will behave the same as:

Example 11.9. Equivalent constraints without tags

```
<rsc_order id="order1-1" first="storage" then="vm1" kind="Mandatory" />
<rsc_order id="order1-2" first="storage" then="vm2" kind="Mandatory" />
<rsc_order id="order1-3" first="storage" then="vm2" kind="Mandatory" />
```

A tag may be used directly in the constraint, or indirectly by being listed in a resource set used in the constraint. When used in a resource set, an expanded tag will honor the set's `sequential` property.

Chapter 12. Utilization and Placement Strategy

Pacemaker decides where to place a resource according to the resource allocation scores on every node. The resource will be allocated to the node where the resource has the highest score.

If the resource allocation scores on all the nodes are equal, by the default placement strategy, Pacemaker will choose a node with the least number of allocated resources for balancing the load. If the number of resources on each node is equal, the first eligible node listed in the CIB will be chosen to run the resource.

Often, in real-world situations, different resources use significantly different proportions of a node's capacities (memory, I/O, etc.). We cannot balance the load ideally just according to the number of resources allocated to a node. Besides, if resources are placed such that their combined requirements exceed the provided capacity, they may fail to start completely or run with degraded performance.

To take these factors into account, Pacemaker allows you to configure:

1. The capacity a certain node provides.
2. The capacity a certain resource requires.
3. An overall strategy for placement of resources.

Utilization attributes

To configure the capacity that a node provides or a resource requires, you can use *utilization attributes* in node and resource objects. You can name utilization attributes according to your preferences and define as many name/value pairs as your configuration needs. However, the attributes' values must be integers.

Example 12.1. Specifying CPU and RAM capacities of two nodes

```
<node id="node1" type="normal" uname="node1">
  <utilization id="node1-utilization">
    <nvpair id="node1-utilization-cpu" name="cpu" value="2"/>
    <nvpair id="node1-utilization-memory" name="memory" value="2048"/>
  </utilization>
</node>
<node id="node2" type="normal" uname="node2">
  <utilization id="node2-utilization">
    <nvpair id="node2-utilization-cpu" name="cpu" value="4"/>
    <nvpair id="node2-utilization-memory" name="memory" value="4096"/>
  </utilization>
</node>
```

Example 12.2. Specifying CPU and RAM consumed by several resources

```
<primitive id="rsc-small" class="ocf" provider="pacemaker" type="Dummy">
  <utilization id="rsc-small-utilization">
    <nvpair id="rsc-small-utilization-cpu" name="cpu" value="1"/>
  </utilization>
</primitive>
```

```
<nvpair id="rsc-small-utilization-memory" name="memory" value="1024"/>
</utilization>
</primitive>
<primitive id="rsc-medium" class="ocf" provider="pacemaker" type="Dummy">
  <utilization id="rsc-medium-utilization">
    <nvpair id="rsc-medium-utilization-cpu" name="cpu" value="2"/>
    <nvpair id="rsc-medium-utilization-memory" name="memory" value="2048"/>
  </utilization>
</primitive>
<primitive id="rsc-large" class="ocf" provider="pacemaker" type="Dummy">
  <utilization id="rsc-large-utilization">
    <nvpair id="rsc-large-utilization-cpu" name="cpu" value="3"/>
    <nvpair id="rsc-large-utilization-memory" name="memory" value="3072"/>
  </utilization>
</primitive>
```

A node is considered eligible for a resource if it has sufficient free capacity to satisfy the resource's requirements. The nature of the required or provided capacities is completely irrelevant to Pacemaker — it just makes sure that all capacity requirements of a resource are satisfied before placing a resource to a node.

Placement Strategy

After you have configured the capacities your nodes provide and the capacities your resources require, you need to set the `placement-strategy` in the global cluster options, otherwise the capacity configurations have *no effect*.

Four values are available for the `placement-strategy`:

default	Utilization values are not taken into account at all. Resources are allocated according to allocation scores. If scores are equal, resources are evenly distributed across nodes.
utilization	Utilization values are taken into account <i>only</i> when deciding whether a node is considered eligible (i.e. whether it has sufficient free capacity to satisfy the resource's requirements). Load-balancing is still done based on the number of resources allocated to a node.
balanced	Utilization values are taken into account when deciding whether a node is eligible to serve a resource <i>and</i> when load-balancing, so an attempt is made to spread the resources in a way that optimizes resource performance.
minimal	Utilization values are taken into account <i>only</i> when deciding whether a node is eligible to serve a resource. For load-balancing, an attempt is made to concentrate the resources on as few nodes as possible, thereby enabling possible power savings on the remaining nodes.

Set `placement-strategy` with `crm_attribute`:

```
# crm_attribute --name placement-strategy --update balanced
```

Now Pacemaker will ensure the load from your resources will be distributed evenly throughout the cluster, without the need for convoluted sets of colocation constraints.

Allocation Details

Which node is preferred to get consumed first when allocating resources?

- The node with the highest node weight gets consumed first. Node weight is a score maintained by the cluster to represent node health.
- If multiple nodes have the same node weight:
 - If `placement-strategy` is `default` or `utilization`, the node that has the least number of allocated resources gets consumed first.
 - If their numbers of allocated resources are equal, the first eligible node listed in the CIB gets consumed first.
 - If `placement-strategy` is `balanced`, the node that has the most free capacity gets consumed first.
 - If the free capacities of the nodes are equal, the node that has the least number of allocated resources gets consumed first.
 - If their numbers of allocated resources are equal, the first eligible node listed in the CIB gets consumed first.
 - If `placement-strategy` is `minimal`, the first eligible node listed in the CIB gets consumed first.

Which node has more free capacity?

If only one type of utilization attribute has been defined, free capacity is a simple numeric comparison.

If multiple types of utilization attributes have been defined, then the node that is numerically highest in the the most attribute types has the most free capacity. For example:

- If `nodeA` has more free `cpus`, and `nodeB` has more free `memory`, then their free capacities are equal.
- If `nodeA` has more free `cpus`, while `nodeB` has more free `memory` and `storage`, then `nodeB` has more free capacity.

Which resource is preferred to be assigned first?

- The resource that has the highest `priority` (see the section called “Resource Options”) gets allocated first.
- If their priorities are equal, check whether they are already running. The resource that has the highest score on the node where it’s running gets allocated first, to prevent resource shuffling.
- If the scores above are equal or the resources are not running, the resource has the highest score on the preferred node gets allocated first.
- If the scores above are equal, the first runnable resource listed in the CIB gets allocated first.

Limitations and Workarounds

The type of problem Pacemaker is dealing with here is known as the knapsack problem [http://en.wikipedia.org/wiki/Knapsack_problem] and falls into the NP-complete [<http://en.wikipedia.org/wiki/NP-complete>] category of computer science problems — a fancy way of saying "it takes a really long time to solve".

Clearly in a HA cluster, it's not acceptable to spend minutes, let alone hours or days, finding an optimal solution while services remain unavailable.

So instead of trying to solve the problem completely, Pacemaker uses a *best effort* algorithm for determining which node should host a particular service. This means it arrives at a solution much faster than traditional linear programming algorithms, but by doing so at the price of leaving some services stopped.

In the contrived example at the start of this section:

- `rsc-small` would be allocated to `node1`
- `rsc-medium` would be allocated to `node2`
- `rsc-large` would remain inactive

Which is not ideal.

There are various approaches to dealing with the limitations of pacemaker's placement strategy:

Ensure you have sufficient physical capacity.

It might sound obvious, but if the physical capacity of your nodes is (close to) maxed out by the cluster under normal conditions, then failover isn't going to go well. Even without the utilization feature, you'll start hitting timeouts and getting secondary failures.

Build some buffer into the capabilities advertised by the nodes.

Advertise slightly more resources than we physically have, on the (usually valid) assumption that a resource will not use 100% of the configured amount of CPU, memory and so forth *all* the time. This practice is sometimes called *overcommit*.

Specify resource priorities.

If the cluster is going to sacrifice services, it should be the ones you care about (comparatively) the least. Ensure that resource priorities are properly set so that your most important resources are scheduled first.

Chapter 13. Access Control Lists (ACLs)

By default, the `root` user or any user in the `haclient` group can modify Pacemaker's CIB without restriction. Pacemaker offers *access control lists (ACLs)* to provide more fine-grained authorization.

ACL Prerequisites

In order to use ACLs:

- The Pacemaker software must have been compiled with ACL support. If the output of the command `pacemakerd --features` contains `acls`, your installation supports ACLs.
- Desired users must have user accounts in the `haclient` group on all nodes in the cluster.
- If your CIB was created before Pacemaker 1.1.12, it may need to be updated to the current schema using `cibadmin --upgrade` in order to use the syntax documented here.
- The `enable-acl` cluster option must be set to `true`.

ACL Configuration

ACLs are specified within an `acls` element of the CIB. The `acls` element may contain any number of `acl_role`, `acl_target`, and `acl_group` elements.

ACL Roles

An ACL role is a collection of permissions allowing or denying access to particular portions of the CIB.

Table 13.1. Properties of an ACL Role

Attribute	Description
<code>id</code>	A unique name for the role (required)
<code>description</code>	Arbitrary text (not used by Pacemaker)

An `acl_role` element may contain any number of `acl_permission` elements.

Table 13.2. Properties of an ACL Permission

Attribute	Description
<code>id</code>	A unique name for the permission (required)
<code>description</code>	Arbitrary text (not used by Pacemaker)
<code>kind</code>	The access being granted. Allowed values are <code>read</code> , <code>write</code> , and <code>deny</code> . A value of <code>write</code> grants both read and write access.

Attribute	Description
<code>object-type</code>	The name of an XML element in the CIB to which the permission applies. (Exactly one of <code>object-type</code> , <code>xpath</code> , and <code>reference</code> must be specified for a permission.)
<code>attribute</code>	If specified, the permission applies only to <code>object-type</code> elements that have this attribute set (to any value). If not specified, the permission applies to all <code>object-type</code> elements. May only be used with <code>object-type</code> .
<code>reference</code>	The ID of an XML element in the CIB to which the permission applies. (Exactly one of <code>object-type</code> , <code>xpath</code> , and <code>reference</code> must be specified for a permission.)
<code>xpath</code>	An XPath [https://www.w3.org/TR/xpath-10/] specification selecting an XML element in the CIB to which the permission applies. Attributes may be specified in the XPath to select particular elements, but the permissions apply to the entire element. (Exactly one of <code>object-type</code> , <code>xpath</code> , and <code>reference</code> must be specified for a permission.)

Important

- Permissions are applied to the selected XML element's entire XML subtree (all elements enclosed within it).
- Write permission grants the ability to create, modify, or remove the element and its subtree, and also the ability to create any "scaffolding" elements (enclosing elements that do not have attributes other than an ID).
- Permissions for more specific matches (more deeply nested elements) take precedence over more general ones.
- If multiple permissions are configured for the same match (for example, in different roles applied to the same user), any deny permission takes precedence, then write, then lastly read.

ACL Targets and Groups

ACL targets correspond to user accounts on the system.

Table 13.3. Properties of an ACL Target

Attribute	Description
<code>id</code>	The name of a user on the system (required)

ACL groups may be specified, but are not currently used by Pacemaker. This is expected to change in a future version.

Table 13.4. Properties of an ACL Group

Attribute	Description
<code>id</code>	The name of a group on the system (required)

Each `acl_target` and `acl_group` element may contain any number of `role` elements.

Table 13.5. Properties of an ACL Role Reference

Attribute	Description
id	The id of an <code>acl_role</code> element that specifies permissions granted to the enclosing target or group

Important

The root and hacluster user accounts always have full access to the CIB, regardless of ACLs. For other user accounts, when `enable-acl` is true, permission to all parts of the CIB is denied by default (permissions must be explicitly granted).

ACL Examples

```
<acls>
```

```
  <acl_role id="read_all">
    <acl_permission id="read_all-cib" kind="read" xpath="/cib" />
  </acl_role>
```

```
  <acl_role id="operator">

    <acl_permission id="operator-maintenance-mode" kind="write"
      xpath="//crm_config/nvpair[@name='maintenance-mode']" />

    <acl_permission id="operator-maintenance-attr" kind="write"
      xpath="//nvpair[@name='maintenance']" />

    <acl_permission id="operator-target-role" kind="write"
      xpath="//resources/meta_attributes/nvpair[@name='target-role']" />

    <acl_permission id="operator-is-managed" kind="write"
      xpath="//resources/nvpair[@name='is-managed']" />

    <acl_permission id="operator-rsc_location" kind="write"
      object-type="rsc_location" />

  </acl_role>
```

```
  <acl_role id="administrator">
    <acl_permission id="administrator-cib" kind="write" xpath="/cib" />
  </acl_role>
```

```
  <acl_role id="minimal">

    <acl_permission id="minimal-standby" kind="read"
      description="allow reading standby node attribute (permanent or transie
      xpath="//instance_attributes/nvpair[@name='standby']" />

    <acl_permission id="minimal-maintenance" kind="read"
      description="allow reading maintenance node attribute (permanent or tra
      xpath="//nvpair[@name='maintenance']" />

  </acl_role>
```

```
<acl_permission id="minimal-target-role" kind="read"
  description="allow reading resource target roles"
  xpath="//resources//meta_attributes/nvpair[@name='target-role']"/>

<acl_permission id="minimal-is-managed" kind="read"
  description="allow reading resource managed status"
  xpath="//resources//meta_attributes/nvpair[@name='is-managed']"/>

<acl_permission id="minimal-deny-instance-attributes" kind="deny"
  xpath="//instance_attributes"/>

<acl_permission id="minimal-deny-meta-attributes" kind="deny"
  xpath="//meta_attributes"/>

<acl_permission id="minimal-deny-operations" kind="deny"
  xpath="//operations"/>

<acl_permission id="minimal-deny-utilization" kind="deny"
  xpath="//utilization"/>

<acl_permission id="minimal-nodes" kind="read"
  description="allow reading node names/IDs (attributes are denied separa
  xpath="/cib/configuration/nodes"/>

<acl_permission id="minimal-resources" kind="read"
  description="allow reading resource names/agents (parameters are denied
  xpath="/cib/configuration/resources"/>

<acl_permission id="minimal-deny-constraints" kind="deny"
  xpath="/cib/configuration/constraints"/>

<acl_permission id="minimal-deny-topology" kind="deny"
  xpath="/cib/configuration/fencing-topology"/>

<acl_permission id="minimal-deny-op_defaults" kind="deny"
  xpath="/cib/configuration/op_defaults"/>

<acl_permission id="minimal-deny-rsc_defaults" kind="deny"
  xpath="/cib/configuration/rsc_defaults"/>

<acl_permission id="minimal-deny-alerts" kind="deny"
  xpath="/cib/configuration/alerts"/>

<acl_permission id="minimal-deny-acls" kind="deny"
  xpath="/cib/configuration/acls"/>

<acl_permission id="minimal-cib" kind="read"
  description="allow reading cib element and crm_config/status sections"
  xpath="/cib"/>

</acl_role>

<acl_target id="alice">
  <role id="minimal"/>
```

```
</acl_target>

<acl_target id="bob">
  <role id="read_all"/>
</acl_target>

<acl_target id="carol">
  <role id="read_all"/>
  <role id="operator"/>
</acl_target>

<acl_target id="dave">
  <role id="administrator"/>
</acl_target>

</acls>
```

In the above example, the user `alice` has the minimal permissions necessary to run basic Pacemaker CLI tools, including using `crm_mon` to view the cluster status, without being able to modify anything. The user `bob` can view the entire configuration and status of the cluster, but not make any changes. The user `carol` can read everything, and change selected cluster properties as well as resource roles and location constraints. Finally, `dave` has full read and write access to the entire CIB.

Looking at the `minimal` role in more depth, it is designed to allow read access to the `cib` tag itself, while denying access to particular portions of its subtree (which is the entire CIB).

This is because the DC node is indicated in the `cib` tag, so `crm_mon` will not be able to report the DC otherwise. However, this does change the security model to allow by default, since any portions of the CIB not explicitly denied will be readable. The `cib` read access could be removed and replaced with read access to just the `crm_config` and `status` sections, for a safer approach at the cost of not seeing the DC in status output.

For a simpler configuration, the `minimal` role allows read access to the entire `crm_config` section, which contains cluster properties. It would be possible to allow read access to specific properties instead (such as `stonith-enabled`, `dc-uuid`, `have-quorum`, and `cluster-name`) to restrict access further while still allowing status output, but cluster properties are unlikely to be considered sensitive.

Chapter 14. Status — Here be dragons

Most users never need to understand the contents of the status section and can be happy with the output from `crm_mon`.

However for those with a curious inclination, this section attempts to provide an overview of its contents.

Node Status

In addition to the cluster's configuration, the CIB holds an up-to-date representation of each cluster node in the `status` section.

Example 14.1. A bare-bones status entry for a healthy node `cl-virt-1`

```
<node_state id="1" uname="cl-virt-1" in_ccm="true" crmd="online" crm-debug-origi
  <transient_attributes id="1"/>
  <lrmd id="1"/>
</node_state>
```

Users are highly recommended *not* to modify any part of a node's state *directly*. The cluster will periodically regenerate the entire section from authoritative sources, so any changes should be done with the tools appropriate to those sources.

Table 14.1. Authoritative Sources for State Information

CIB Object	Authoritative Source
<code>node_state</code>	<code>pacemaker-controld</code>
<code>transient_attributes</code>	<code>pacemaker-attd</code>
<code>lrmd</code>	<code>pacemaker-execd</code>

The fields used in the `node_state` objects are named as they are largely for historical reasons and are rooted in Pacemaker's origins as the resource manager for the older Heartbeat project. They have remained unchanged to preserve compatibility with older versions.

Table 14.2. Node Status Fields

Field	Description
<code>id</code>	Unique identifier for the node. Corosync-based clusters use a numeric counter.
<code>uname</code>	The node's name as known by the cluster
<code>in_ccm</code>	Is the node a member at the cluster communication layer? Allowed values: <code>true</code> , <code>false</code> .
<code>crmd</code>	Is the node a member at the pacemaker layer? Allowed values: <code>online</code> , <code>offline</code> .
<code>crm-debug-origin</code>	The name of the source function that made the most recent change (for debugging purposes).
<code>join</code>	Does the node participate in hosting resources? Allowed values: <code>down</code> , <code>pending</code> , <code>member</code> , <code>banned</code> .

Field	Description
expected	Expected value for join.

The cluster uses these fields to determine whether, at the node level, the node is healthy or is in a failed state and needs to be fenced.

Transient Node Attributes

Like regular node attributes, the name/value pairs listed in the `transient_attributes` section help to describe the node. However they are forgotten by the cluster when the node goes offline. This can be useful, for instance, when you want a node to be in standby mode (not able to run resources) just until the next reboot.

In addition to any values the administrator sets, the cluster will also store information about failed resources here.

Example 14.2. A set of transient node attributes for node `cl-virt-1`

```
<transient_attributes id="cl-virt-1">
  <instance_attributes id="status-cl-virt-1">
    <nvpair id="status-cl-virt-1-pingd" name="pingd" value="3"/>
    <nvpair id="status-cl-virt-1-probe_complete" name="probe_complete" value="true"/>
    <nvpair id="status-cl-virt-1-fail-count-pingd:0.monitor_30000" name="fail-count-pingd:0.monitor_30000" value="1"/>
    <nvpair id="status-cl-virt-1-last-failure-pingd:0" name="last-failure-pingd:0" value="09:22:22 UTC 6 April 2009"/>
  </instance_attributes>
</transient_attributes>
```

In the above example, we can see that a monitor on the `pingd:0` resource has failed once, at 09:22:22 UTC 6 April 2009.¹ We also see that the node is connected to three **pingd** peers and that all known resources have been checked for on this machine (`probe_complete`).

Operation History

A node's resource history is held in the `lrm_resources` tag (a child of the `lrm` tag). The information stored here includes enough information for the cluster to stop the resource safely if it is removed from the configuration section. Specifically, the resource's `id`, `class`, `type` and `provider` are stored.

Example 14.3. A record of the `apcstonith` resource

```
<lrm_resource id="apcstonith" type="apcmastersnmp" class="stonith"/>
```

Additionally, we store the last job for every combination of resource, action and interval. The concatenation of the values in this tuple are used to create the `id` of the `lrm_rsc_op` object.

Table 14.3. Contents of an `lrm_rsc_op` job

Field	Description
id	

¹ You can use the standard `date` command to print a human-readable version of any seconds-since-epoch value, for example `date -d @1239009742`.

Field	Description
	Identifier for the job constructed from the resource's <code>id</code> , <code>operation</code> and <code>interval</code> .
<code>call-id</code>	The job's ticket number. Used as a sort key to determine the order in which the jobs were executed.
<code>operation</code>	The action the resource agent was invoked with.
<code>interval</code>	The frequency, in milliseconds, at which the operation will be repeated. A one-off job is indicated by 0.
<code>op-status</code>	The job's status. Generally this will be either 0 (done) or -1 (pending). Rarely used in favor of <code>rc-code</code> .
<code>rc-code</code>	The job's result. Refer to the <i>Resource Agents</i> section of <i>Pacemaker Administration</i> for details on what the values here mean and how they are interpreted.
<code>last-run</code>	Machine-local date/time, in seconds since epoch, at which the job was executed. For diagnostic purposes.
<code>last-rc-change</code>	Machine-local date/time, in seconds since epoch, at which the job first returned the current value of <code>rc-code</code> . For diagnostic purposes.
<code>exec-time</code>	Time, in milliseconds, that the job was running for. For diagnostic purposes.
<code>queue-time</code>	Time, in seconds, that the job was queued for in the LRMD. For diagnostic purposes.
<code>crm_feature_set</code>	The version which this job description conforms to. Used when processing <code>op-digest</code> .
<code>transition-key</code>	A concatenation of the job's graph action number, the graph number, the expected result and the UUID of the controller instance that scheduled it. This is used to construct <code>transition-magic</code> (below).
<code>transition-magic</code>	

Field	Description
	A concatenation of the job's <code>op-status</code> , <code>rc-code</code> and <code>transition-key</code> . Guaranteed to be unique for the life of the cluster (which ensures it is part of CIB update notifications) and contains all the information needed for the controller to correctly analyze and process the completed job. Most importantly, the decomposed elements tell the controller if the job entry was expected and whether it failed.
<code>op-digest</code>	An MD5 sum representing the parameters passed to the job. Used to detect changes to the configuration, to restart resources if necessary.
<code>crm-debug-origin</code>	The origin of the current values. For diagnostic purposes.

Simple Operation History Example

Example 14.4. A monitor operation (determines current state of the `apcstonith` resource)

```
<lrn_resource id="apcstonith" type="apcmastersnmp" class="stonith">
  <lrn_rsc_op id="apcstonith_monitor_0" operation="monitor" call-id="2"
    rc-code="7" op-status="0" interval="0"
    crm-debug-origin="do_update_resource" crm_feature_set="3.0.1"
    op-digest="2e3da9274d3550dc6526fb24bfcba0"
    transition-key="22:2:7:2668bbeb-06d5-40f9-936d-24cb7f87006a"
    transition-magic="0:7:22:2:7:2668bbeb-06d5-40f9-936d-24cb7f87006a"
    last-run="1239008085" last-rc-change="1239008085" exec-time="10" queue-time="0"
  </lrn_rsc_op>
</lrn_resource>
```

In the above example, the job is a non-recurring monitor operation often referred to as a "probe" for the `apcstonith` resource.

The cluster schedules probes for every configured resource on a node when the node first starts, in order to determine the resource's current state before it takes any further action.

From the `transition-key`, we can see that this was the 22nd action of the 2nd graph produced by this instance of the controller (2668bbeb-06d5-40f9-936d-24cb7f87006a).

The third field of the `transition-key` contains a 7, which indicates that the job expects to find the resource inactive. By looking at the `rc-code` property, we see that this was the case.

As that is the only job recorded for this node, we can conclude that the cluster started the resource elsewhere.

Complex Operation History Example

Example 14.5. Resource history of a `pingd` clone with multiple jobs

```
<lrn_resource id="pingd:0" type="pingd" class="ocf" provider="pacemaker">
  <lrn_rsc_op id="pingd:0_monitor_30000" operation="monitor" call-id="34"
```

```
rc-code="0" op-status="0" interval="30000"
crm-debug-origin="do_update_resource" crm_feature_set="3.0.1"
transition-key="10:11:0:2668bbeb-06d5-40f9-936d-24cb7f87006a"
...
last-run="1239009741" last-rc-change="1239009741" exec-time="10" queue-time="0"
<lrmsrc_op id="pingd:0_stop_0" operation="stop"
  crm-debug-origin="do_update_resource" crm_feature_set="3.0.1" call-id="32"
  rc-code="0" op-status="0" interval="0"
  transition-key="11:11:0:2668bbeb-06d5-40f9-936d-24cb7f87006a"
  ...
  last-run="1239009741" last-rc-change="1239009741" exec-time="10" queue-time="0"
<lrmsrc_op id="pingd:0_start_0" operation="start" call-id="33"
  rc-code="0" op-status="0" interval="0"
  crm-debug-origin="do_update_resource" crm_feature_set="3.0.1"
  transition-key="31:11:0:2668bbeb-06d5-40f9-936d-24cb7f87006a"
  ...
  last-run="1239009741" last-rc-change="1239009741" exec-time="10" queue-time="0"
<lrmsrc_op id="pingd:0_monitor_0" operation="monitor" call-id="3"
  rc-code="0" op-status="0" interval="0"
  crm-debug-origin="do_update_resource" crm_feature_set="3.0.1"
  transition-key="23:2:7:2668bbeb-06d5-40f9-936d-24cb7f87006a"
  ...
  last-run="1239008085" last-rc-change="1239008085" exec-time="20" queue-time="0"
</lrmsrc_op>
```

When more than one job record exists, it is important to first sort them by `call-id` before interpreting them.

Once sorted, the above example can be summarized as:

1. A non-recurring monitor operation returning 7 (not running), with a `call-id` of 3
2. A stop operation returning 0 (success), with a `call-id` of 32
3. A start operation returning 0 (success), with a `call-id` of 33
4. A recurring monitor returning 0 (success), with a `call-id` of 34

The cluster processes each job record to build up a picture of the resource's state. After the first and second entries, it is considered stopped, and after the third it considered active.

Based on the last operation, we can tell that the resource is currently active.

Additionally, from the presence of a `stop` operation with a lower `call-id` than that of the `start` operation, we can conclude that the resource has been restarted. Specifically this occurred as part of actions 11 and 31 of transition 11 from the controller instance with the key 2668bbeb.... This information can be helpful for locating the relevant section of the logs when looking for the source of a failure.

Chapter 15. Multi-Site Clusters and Tickets

Apart from local clusters, Pacemaker also supports multi-site clusters. That means you can have multiple, geographically dispersed sites, each with a local cluster. Failover between these clusters can be coordinated manually by the administrator, or automatically by a higher-level entity called a *Cluster Ticket Registry (CTR)*.

Challenges for Multi-Site Clusters

Typically, multi-site environments are too far apart to support synchronous communication and data replication between the sites. That leads to significant challenges:

- How do we make sure that a cluster site is up and running?
- How do we make sure that resources are only started once?
- How do we make sure that quorum can be reached between the different sites and a split-brain scenario avoided?
- How do we manage failover between sites?
- How do we deal with high latency in case of resources that need to be stopped?

In the following sections, learn how to meet these challenges.

Conceptual Overview

Multi-site clusters can be considered as “overlay” clusters where each cluster site corresponds to a cluster node in a traditional cluster. The overlay cluster can be managed by a CTR in order to guarantee that any cluster resource will be active on no more than one cluster site. This is achieved by using *tickets* that are treated as failover domain between cluster sites, in case a site should be down.

The following sections explain the individual components and mechanisms that were introduced for multi-site clusters in more detail.

Ticket

Tickets are, essentially, cluster-wide attributes. A ticket grants the right to run certain resources on a specific cluster site. Resources can be bound to a certain ticket by `rsc_ticket` constraints. Only if the ticket is available at a site can the respective resources be started there. Vice versa, if the ticket is revoked, the resources depending on that ticket must be stopped.

The ticket thus is similar to a *site quorum*, i.e. the permission to manage/own resources associated with that site. (One can also think of the current `have-quorum` flag as a special, cluster-wide ticket that is granted in case of node majority.)

Tickets can be granted and revoked either manually by administrators (which could be the default for classic enterprise clusters), or via the automated CTR mechanism described below.

A ticket can only be owned by one site at a time. Initially, none of the sites has a ticket. Each ticket must be granted once by the cluster administrator.

The presence or absence of tickets for a site is stored in the CIB as a cluster status. With regards to a certain ticket, there are only two states for a site: `true` (the site has the ticket) or `false` (the site does not have the ticket). The absence of a certain ticket (during the initial state of the multi-site cluster) is the same as the value `false`.

Dead Man Dependency

A site can only activate resources safely if it can be sure that the other site has deactivated them. However after a ticket is revoked, it can take a long time until all resources depending on that ticket are stopped "cleanly", especially in case of cascaded resources. To cut that process short, the concept of a *Dead Man Dependency* was introduced.

If a dead man dependency is in force, if a ticket is revoked from a site, the nodes that are hosting dependent resources are fenced. This considerably speeds up the recovery process of the cluster and makes sure that resources can be migrated more quickly.

This can be configured by specifying a `loss-policy="fence"` in `rsc_ticket` constraints.

Cluster Ticket Registry

A CTR is a coordinated group of network daemons that automatically handles granting, revoking, and timing out tickets (instead of the administrator revoking the ticket somewhere, waiting for everything to stop, and then granting it on the desired site).

Pacemaker does not implement its own CTR, but interoperates with external software designed for that purpose (similar to how resource and fencing agents are not directly part of pacemaker).

Participating clusters run the CTR daemons, which connect to each other, exchange information about their connectivity, and vote on which sites gets which tickets.

A ticket is granted to a site only once the CTR is sure that the ticket has been relinquished by the previous owner, implemented via a timer in most scenarios. If a site loses connection to its peers, its tickets time out and recovery occurs. After the connection timeout plus the recovery timeout has passed, the other sites are allowed to re-acquire the ticket and start the resources again.

This can also be thought of as a "quorum server", except that it is not a single quorum ticket, but several.

Configuration Replication

As usual, the CIB is synchronized within each cluster, but it is *not* synchronized across cluster sites of a multi-site cluster. You have to configure the resources that will be highly available across the multi-site cluster for every site accordingly.

Configuring Ticket Dependencies

The `rsc_ticket` constraint lets you specify the resources depending on a certain ticket. Together with the constraint, you can set a `loss-policy` that defines what should happen to the respective resources if the ticket is revoked.

The attribute `loss-policy` can have the following values:

- `fence`: Fence the nodes that are running the relevant resources.
- `stop`: Stop the relevant resources.

- `freeze`: Do nothing to the relevant resources.
- `demote`: Demote relevant resources that are running in master mode to slave mode.

Example 15.1. Constraint that fences node if `ticketA` is revoked

```
<rsc_ticket id="rsc1-req-ticketA" rsc="rsc1" ticket="ticketA" loss-policy="fence" />
```

The example above creates a constraint with the ID `rsc1-req-ticketA`. It defines that the resource `rsc1` depends on `ticketA` and that the node running the resource should be fenced if `ticketA` is revoked.

If resource `rsc1` were a promotable resource (i.e. it could run in master or slave mode), you might want to configure that only master mode depends on `ticketA`. With the following configuration, `rsc1` will be demoted to slave mode if `ticketA` is revoked:

Example 15.2. Constraint that demotes `rsc1` if `ticketA` is revoked

```
<rsc_ticket id="rsc1-req-ticketA" rsc="rsc1" rsc-role="Master" ticket="ticketA" lo
```

You can create multiple `rsc_ticket` constraints to let multiple resources depend on the same ticket. However, `rsc_ticket` also supports resource sets (see the section called “Resource Sets”), so one can easily list all the resources in one `rsc_ticket` constraint instead.

Example 15.3. Ticket constraint for multiple resources

```
<rsc_ticket id="resources-dep-ticketA" ticket="ticketA" loss-policy="fence">
  <resource_set id="resources-dep-ticketA-0" role="Started">
    <resource_ref id="rsc1"/>
    <resource_ref id="group1"/>
    <resource_ref id="clone1"/>
  </resource_set>
  <resource_set id="resources-dep-ticketA-1" role="Master">
    <resource_ref id="ms1"/>
  </resource_set>
</rsc_ticket>
```

In the example above, there are two resource sets, so we can list resources with different roles in a single `rsc_ticket` constraint. There’s no dependency between the two resource sets, and there’s no dependency among the resources within a resource set. Each of the resources just depends on `ticketA`.

Referencing resource templates in `rsc_ticket` constraints, and even referencing them within resource sets, is also supported.

If you want other resources to depend on further tickets, create as many constraints as necessary with `rsc_ticket`.

Managing Multi-Site Clusters

Granting and Revoking Tickets Manually

You can grant tickets to sites or revoke them from sites manually. If you want to re-distribute a ticket, you should wait for the dependent resources to stop cleanly at the previous site before you grant the ticket to the new site.

Use the `crm_ticket` command line tool to grant and revoke tickets.

To grant a ticket to this site:

```
# crm_ticket --ticket ticketA --grant
```

To revoke a ticket from this site:

```
# crm_ticket --ticket ticketA --revoke
```

Important

If you are managing tickets manually, use the `crm_ticket` command with great care, because it cannot check whether the same ticket is already granted elsewhere.

Granting and Revoking Tickets via a Cluster Ticket Registry

We will use Booth [<https://github.com/ClusterLabs/booth>] here as an example of software that can be used with pacemaker as a Cluster Ticket Registry. Booth implements the Raft [http://en.wikipedia.org/wiki/Raft_%28computer_science%29] algorithm to guarantee the distributed consensus among different cluster sites, and manages the ticket distribution (and thus the failover process between sites).

Each of the participating clusters and *arbitrators* runs the Booth daemon `boothd`.

An *arbitrator* is the multi-site equivalent of a quorum-only node in a local cluster. If you have a setup with an even number of sites, you need an additional instance to reach consensus about decisions such as failover of resources across sites. In this case, add one or more arbitrators running at additional sites. Arbitrators are single machines that run a booth instance in a special mode. An arbitrator is especially important for a two-site scenario, otherwise there is no way for one site to distinguish between a network failure between it and the other site, and a failure of the other site.

The most common multi-site scenario is probably a multi-site cluster with two sites and a single arbitrator on a third site. However, technically, there are no limitations with regards to the number of sites and the number of arbitrators involved.

`Boothd` at each site connects to its peers running at the other sites and exchanges connectivity details. Once a ticket is granted to a site, the booth mechanism will manage the ticket automatically: If the site which holds the ticket is out of service, the booth daemons will vote which of the other sites will get the ticket. To protect against brief connection failures, sites that lose the vote (either explicitly or implicitly by being disconnected from the voting body) need to relinquish the ticket after a time-out. Thus, it is made sure that a ticket will only be re-distributed after it has been relinquished by the previous site. The resources that depend on that ticket will fail over to the new site holding the ticket. The nodes that have run the resources before will be treated according to the `loss-policy` you set within the `rsc_ticket` constraint.

Before the booth can manage a certain ticket within the multi-site cluster, you initially need to grant it to a site manually via the `booth` command-line tool. After you have initially granted a ticket to a site, `boothd` will take over and manage the ticket automatically.

Important

The `booth` command-line tool can be used to grant, list, or revoke tickets and can be run on any machine where `boothd` is running. If you are managing tickets via Booth, use only `booth` for manual intervention, not `crm_ticket`. That ensures the same ticket will only be owned by one cluster site at a time.

Booth Requirements

- All clusters that will be part of the multi-site cluster must be based on Pacemaker.
- Booth must be installed on all cluster nodes and on all arbitrators that will be part of the multi-site cluster.
- Nodes belonging to the same cluster site should be synchronized via NTP. However, time synchronization is not required between the individual cluster sites.

General Management of Tickets

Display the information of tickets:

```
# crm_ticket --info
```

Or you can monitor them with:

```
# crm_mon --tickets
```

Display the `rsc_ticket` constraints that apply to a ticket:

```
# crm_ticket --ticket ticketA --constraints
```

When you want to do maintenance or manual switch-over of a ticket, revoking the ticket would trigger the loss policies. If `loss-policy="fence"`, the dependent resources could not be gracefully stopped/demoted, and other unrelated resources could even be affected.

The proper way is making the ticket *standby* first with:

```
# crm_ticket --ticket ticketA --standby
```

Then the dependent resources will be stopped or demoted gracefully without triggering the loss policies.

If you have finished the maintenance and want to activate the ticket again, you can run:

```
# crm_ticket --ticket ticketA --activate
```

For more information

- SUSE's Geo Clustering quick start [https://www.suse.com/documentation/sle-ha-geo-12/art_ha_geo_quick/data/art_ha_geo_quick.html]
- Booth [<https://github.com/ClusterLabs/booth>]

Appendix A. Sample Configurations

Empty

Example A.1. An Empty Configuration

```
<cib crm_feature_set="3.0.7" validate-with="pacemaker-1.2" admin_epoch="1" epoch="1">
  <configuration>
    <crm_config/>
    <nodes/>
    <resources/>
    <constraints/>
  </configuration>
  <status/>
</cib>
```

Simple

Example A.2. A simple configuration with two nodes, some cluster options and a resource

```
<cib crm_feature_set="3.0.7" validate-with="pacemaker-1.2" admin_epoch="1" epoch="1">
  <configuration>
    <crm_config>
      <cluster_property_set id="cib-bootstrap-options">
        <nvpair id="option-1" name="symmetric-cluster" value="true"/>
        <nvpair id="option-2" name="no-quorum-policy" value="stop"/>
        <nvpair id="option-3" name="stonith-enabled" value="0"/>
      </cluster_property_set>
    </crm_config>
    <nodes>
      <node id="xxx" uname="c001n01" type="normal"/>
      <node id="yyy" uname="c001n02" type="normal"/>
    </nodes>
    <resources>
      <primitive id="myAddr" class="ocf" provider="heartbeat" type="IPAddr">
        <operations>
          <op id="myAddr-monitor" name="monitor" interval="300s"/>
        </operations>
        <instance_attributes id="myAddr-params">
          <nvpair id="myAddr-ip" name="ip" value="192.0.2.10"/>
        </instance_attributes>
      </primitive>
    </resources>
    <constraints>
      <rsc_location id="myAddr-prefer" rsc="myAddr" node="c001n01" score="INFINITY"/>
    </constraints>
    <rsc_defaults>
      <meta_attributes id="rsc_defaults-options">
        <nvpair id="rsc-default-1" name="resource-stickiness" value="100"/>
      </meta_attributes>
    </rsc_defaults>
  </configuration>
  <status/>
</cib>
```



```
        <nvpair id="rsc-default-2" name="migration-threshold" value="10"/>
    </meta_attributes>
</rsc_defaults>
<op_defaults>
    <meta_attributes id="op_defaults-options">
        <nvpair id="op-default-1" name="timeout" value="30s"/>
    </meta_attributes>
</op_defaults>
</configuration>
<status/>
</cib>
```

In the above example, we have one resource (an IP address) that we check every five minutes and will run on host c001n01 until either the resource fails 10 times or the host shuts down.

Advanced Configuration

Example A.3. An advanced configuration with groups, clones and STONITH

```
<cib crm_feature_set="3.0.7" validate-with="pacemaker-1.2" admin_epoch="1" epoch="1">
  <configuration>
    <crm_config>
      <cluster_property_set id="cib-bootstrap-options">
        <nvpair id="option-1" name="symmetric-cluster" value="true"/>
        <nvpair id="option-2" name="no-quorum-policy" value="stop"/>
        <nvpair id="option-3" name="stonith-enabled" value="true"/>
      </cluster_property_set>
    </crm_config>
    <nodes>
      <node id="xxx" uname="c001n01" type="normal"/>
      <node id="yyy" uname="c001n02" type="normal"/>
      <node id="zzz" uname="c001n03" type="normal"/>
    </nodes>
    <resources>
      <primitive id="myAddr" class="ocf" provider="heartbeat" type="IPaddr">
        <operations>
          <op id="myAddr-monitor" name="monitor" interval="300s"/>
        </operations>
        <instance_attributes id="myAddr-attrs">
          <nvpair id="myAddr-attr-1" name="ip" value="192.0.2.10"/>
        </instance_attributes>
      </primitive>
      <group id="myGroup">
        <primitive id="database" class="lsb" type="oracle">
          <operations>
            <op id="database-monitor" name="monitor" interval="300s"/>
          </operations>
        </primitive>
        <primitive id="webserver" class="lsb" type="apache">
          <operations>
            <op id="webserver-monitor" name="monitor" interval="300s"/>
          </operations>
        </primitive>
      </group>
    </resources>
  </configuration>
</cib>
```

```
</group>
<clone id="STONITH">
  <meta_attributes id="stonith-options">
    <nvpair id="stonith-option-1" name="globally-unique" value="false"/>
  </meta_attributes>
  <primitive id="stonithclone" class="stonith" type="external/ssh">
    <operations>
      <op id="stonith-op-mon" name="monitor" interval="5s"/>
    </operations>
    <instance_attributes id="stonith-attrs">
      <nvpair id="stonith-attr-1" name="hostlist" value="c001n01,c001n02"/>
    </instance_attributes>
  </primitive>
</clone>
</resources>
<constraints>
  <rsc_location id="myAddr-prefer" rsc="myAddr" node="c001n01"
    score="INFINITY"/>
  <rsc_colocation id="group-with-ip" rsc="myGroup" with-rsc="myAddr"
    score="INFINITY"/>
</constraints>
<op_defaults>
  <meta_attributes id="op_defaults-options">
    <nvpair id="op-default-1" name="timeout" value="30s"/>
  </meta_attributes>
</op_defaults>
<rsc_defaults>
  <meta_attributes id="rsc_defaults-options">
    <nvpair id="rsc-default-1" name="resource-stickiness" value="100"/>
    <nvpair id="rsc-default-2" name="migration-threshold" value="10"/>
  </meta_attributes>
</rsc_defaults>
</configuration>
<status/>
</cib>
```

Appendix B. Revision History

Revision History

Revision 1-0	19 Oct 2009	AndrewBeekhof<andrew@beekhof.net>
Import from Pages.app		
Revision 2-0	26 Oct 2009	AndrewBeekhof<andrew@beekhof.net>
Cleanup and reformatting of docbook xml complete		
Revision 3-0	Tue Nov 12 2009	AndrewBeekhof<andrew@beekhof.net>
Split book into chapters and pass validation		
Re-organize book for use with Publican [https://fedorahosted.org/publican/]		
Revision 3-1	Tue Nov 12 2009	TanjaRothSUSE<taroth@suse.com>, LarsMarowsky-BreeSUSE<lmb@suse.com>, YanGaoSUSE<ygao@suse.com>, ThomasSchraitleSUSE<toms@suse.com>, DejanMuhamedagicSUSE<dmuhamedagic@suse.com>
Utilization chapter		
Resource Templates chapter		
Multi-Site Clusters chapter		
Revision 3-2	Fri Nov 4 2011	PhilippMarekLINBit<philipp.marek@linbit.com>
Extensive style, formatting, and indexing updates		
Revision 4-0	Mon Oct 8 2012	AndrewBeekhof<andrew@beekhof.net>
Converted to asciidoc [http://www.methods.co.nz/asciidoc] (which is converted to docbook for use with Publican [https://fedorahosted.org/publican/])		
Revision 5-0	Mon Feb 23 2015	KenGaillot<kgaillet@redhat.com>
Update for clarity, stylistic consistency and current command-line syntax		
Revision 6-0	Tue Dec 8 2015	KenGaillot<kgaillet@redhat.com>
Update for Pacemaker 1.1.14		
Revision 7-0	Tue May 3 2016	KenGaillot<kgaillet@redhat.com>
Update for Pacemaker 1.1.15		
Revision 7-1	Fri Oct 28 2016	KenGaillot<kgaillet@redhat.com>
Overhaul upgrade documentation, and document node health strategies		
Revision 8-0	Tue Oct 25 2016	KenGaillot<kgaillet@redhat.com>
Update for Pacemaker 1.1.16		
Revision 9-0	Tue Jul 11 2017	KenGaillot<kgaillet@redhat.com>
Update for Pacemaker 1.1.17		
Revision 10-0	Fri Oct 6 2017	KenGaillot<kgaillet@redhat.com>
Update for Pacemaker 1.1.18		
Revision 11-0	Fri Jan 12 2018	KenGaillot<kgaillet@redhat.com>
Update for Pacemaker 2.0.0		
Revision 12-0	Mon Jan 28 2019	KenGaillot<kgaillet@redhat.com>, ReidWahl<nwahl@redhat.com>, JanPokorný<jpokorny@redhat.com>
Update for Pacemaker 2.0.1, remove "Further Reading" and "FAQ" sections, and add minor clarifications and reformatting		
Revision 12-1	Mon May 13 2019	KenGaillot<kgaillet@redhat.com>, MaciejSobkowiak<msobkowiak@egnyte.com>
Document podman support, cluster-name cluster option, and new HealthIOWait agent, with other minor clarifications and corrections		
Revision 12-2	Wed Jun 19 2019	KenGaillot<kgaillet@redhat.com>
Add chapter for ACLs		
Revision 13-0	Tue Oct 15 2019	KenGaillot<kgaillet@redhat.com>

Overhaul fencing, rules, and constraints chapters, elaborate on various options, and update for Pacemaker 2.0.3

Index

Symbols

#digests-, 14
Node attribute, 14
#node-unfenced, 14
Node attribute, 14

A

access control list, 116
acl_group
id, 117
acl_permission
attribute, 117
description, 116
id, 116
kind, 116
object-type, 117
reference, 117
xpath, 117
acl_role
description, 116
id, 116
acl_target
id, 117
role
id, 118
ACL, 116
acl_group, 117
id, 117
acl_permission, 116, 116, 116, 117, 117, 117, 117
attribute, 117
description, 116
id, 116
kind, 116
object-type, 117
reference, 117
xpath, 117
acl_role, 116, 116
description, 116
id, 116
acl_target, 117
id, 117
Action, 23
Property
enabled, 24, 24
id, 24
interval, 24
name, 24
on-fail, 24
role, 25
timeout, 24

Status
call-id, 123
crm-debug-origin, 124
crm_feature_set, 123
exec-time, 123
id, 122
interval, 123
last-rc-change, 123
last-run, 123
op-digest, 124
op-status, 123
operation, 123
queue-time, 123
rc-code, 123
transition-key, 123
transition-magic, 123

action, 91
Ordering Constraints, 91
action attribute, 35
resource_set element, 35
Action Property, 24, 24, 24, 24, 24, 24, 24, 25
Action Status, 122, 123, 123, 123, 123, 123, 123, 123, 123, 123, 123, 123, 124, 124
active_resource, 93
Notification Environment Variable, 93
active_uname, 93
Notification Environment Variable, 93
add-host attribute, 101
network element, 101
admin_epoch, 6
Cluster Option, 6
Alert
Option
timeout, 59
timestamp-format, 59
Alerts, 58
Asymmetrical Clusters, 30
attribute, 12, 117
#digests-, 14
#node-unfenced, 14
acl_permission, 117
fail-count-, 13
last-failure-, 13
maintenance, 13
probe_complete, 13
resource-discovery-enabled, 14
shutdown, 14
site-name, 14
standby, 14
terminate, 14
attribute attribute, 65
expression element, 65
attribute_name, 62
attribute_value, 62

B

- batch-limit, 7
 - Cluster Option, 7
- boolean-op attribute, 64
 - rule element, 64
- Bundle, 99, 99, 99, 99
 - Container, 100
 - Meta-attributes, 104
 - Networking, 101
 - Node Attributes, 104
 - Prerequisites, 99
 - Primitive, 103
 - Storage, 102
- bundle element, 99, 100, 100
 - description attribute, 100
 - id attribute, 100

C

- call-id, 123
 - Action Status, 123
- cib-last-written, 6
 - Cluster Property, 6
- class, 15, 18
 - Resource, 18
- Clone
 - Option
 - clone-max, 87
 - clone-min, 88
 - clone-node-max, 88
 - globally-unique, 87
 - interleave, 88
 - notify, 88
 - ordered, 88
 - promotable, 88
 - promoted-max, 88
 - promoted-node-max, 88
 - Property
 - id, 87
- Clone Option, 87, 87, 88, 88, 88, 88, 88, 88, 88, 88
- Clone Property, 87
- Clone Resources, 86
- clone-max, 87
 - Clone Option, 87
- clone-min, 88
 - Clone Option, 88
- clone-node-max, 88
 - Clone Option, 88
- Clones, 86, 91
- Cluster, 6
 - Option
 - admin_epoch, 6
 - batch-limit, 7
 - cluster-delay, 9

- cluster-ipc-limit, 9
- cluster-recheck-interval, 10
- concurrent-fencing, 8
- Configuration Version, 6
- dc-deadtime, 9
- election-timeout, 11
- enable-acl, 9
- enable-startup-probes, 8
- epoch, 6
- fence-reaction, 8
- join-finalization-timeout, 11
- join-integration-timeout, 11
- maintenance-mode, 8
- migration-limit, 7
- no-quorum-policy, 7
- node-health-base, 9
- node-health-green, 10
- node-health-red, 10
- node-health-strategy, 9
- node-health-yellow, 10
- num_updates, 6
- pe-error-series-max, 9
- pe-input-series-max, 9
- pe-warn-series-max, 9
- placement-strategy, 9
- priority-fencing-delay, 8
- remove-after-stop, 10
- shutdown-escalation, 11
- shutdown-lock, 10
- shutdown-lock-limit, 10
- start-failure-is-fatal, 7
- startup-fencing, 10
- stonith-action, 8
- stonith-enabled, 8
- stonith-max-attempts, 8
- stonith-timeout, 8
- stonith-watchdog-timeout, 8
- stop-all-resources, 7
- stop-orphan-actions, 7
- stop-orphan-resources, 7
- symmetric-cluster, 7
- transition-delay, 11
- validate-with, 6

- Property
 - cib-last-written, 6
 - cluster-infrastructure, 7
 - cluster-name, 7
 - dc-uuid, 6
 - dc-version, 7
 - have-quorum, 6
- Setting Options with Rules, 72
- Cluster Option, 6, 6, 6, 6, 7, 7, 7, 7, 7, 7, 7, 7, 8, 8, 8, 8, 8, 8, 8, 8, 9, 9, 9, 9, 9, 9, 9, 9, 9, 9, 10, 10, 10, 10, 10, 10, 10, 10, 11, 11, 11, 11, 11

- Cluster Property, 6, 6, 6, 7, 7, 7
- cluster-delay, 9
 - Cluster Option, 9
- cluster-infrastructure, 7
 - Cluster Property, 7
- cluster-ipc-limit, 9
 - Cluster Option, 9
- cluster-name, 7
 - Cluster Property, 7
- cluster-recheck-interval, 10
 - Cluster Option, 10
- Colocation Constraint, 32
 - rsc_colocation element, 33
- concurrent-fencing, 8
 - Cluster Option, 8
- Configuration, 41, 41
- Configuration Version, 6
 - Cluster, 6
- Constraint, 28
 - Colocation Constraint, 32
 - rsc_colocation element, 33
 - Location Constraint, 28
 - Resource Discovery, 30
 - rsc_location element, 28
 - Ordering Constraint, 31
 - rsc_order element, 31
 - Resource Set, 34
 - Rule, 64
- Constraints
 - Ordering
 - action, 91
 - role, 90
 - rsc-role, 89
 - with-rsc-role, 89
- Container, 100
 - Docker
 - Bundle, 99
 - podman
 - Bundle, 99
 - rkt
 - Bundle, 99
- control-port attribute, 101
 - network element, 101
- Controlling Cluster Options, 72
- crm-debug-origin, 121, 124
 - Action Status, 124
 - Node Status, 121
- crmd, 121
 - Node Status, 121
- CRM_alert_
 - attribute_name, 62
 - attribute_value, 62
 - desc, 62
 - exec_time, 62

- interval, 62
- kind, 61
- node, 62
- nodeid, 62
- rc, 62
- recipient, 61
- rsc, 62
- status, 62
- target_rc, 62
- task, 62
- timestamp, 62
- timestamp_epoch, 62
- timestamp_usec, 62
- version, 61
- CRM_alert_node_
 - sequence, 62
- crm_feature_set, 123
 - Action Status, 123
- custom, 82

D

- dampen, 78
 - Ping Resource Option, 78
- Date Specification, 67
- Date/Time Expression, 66
 - Date Specification, 67
 - Duration, 67
- date_expression element, 66, 66, 66, 66, 66
 - end attribute, 66
 - id attribute, 66
 - operation attribute, 66
 - start attribute, 66
- date_spec element, 67, 67, 67, 67, 67, 67, 67, 67, 67, 67, 67
 - hours attribute, 67
 - id attribute, 67
 - monthdays attribute, 67
 - months attribute, 67
 - moon attribute, 67
 - weekdays attribute, 67
 - weeks attribute, 67
 - weekyears attribute, 67
 - yeardays attribute, 67
 - years attribute, 67
- dc-deadtime, 9
 - Cluster Option, 9
- dc-uuid, 6
 - Cluster Property, 6
- dc-version, 7
 - Cluster Property, 7
- demote_resource, 94
 - Notification Environment Variable, 94
- demote_uname, 94

- Notification Environment Variable, 94
- desc, 62
- description, 116, 116
 - acl_permission, 116
 - acl_role, 116
- description attribute, 100
 - bundle element, 100
- Determine by Rules, 69
- Determine Resource Location, 69
- devices, 52
 - fencing-level, 52
- Docker
 - Bundle, 99
- docker element, 100, 100, 100, 100, 100, 100, 100, 100
 - image attribute, 100
 - network attribute, 100
 - options attribute, 100
 - promoted-max attribute, 100
 - replicas attribute, 100
 - replicas-per-host attribute, 100
 - run-command attribute, 100
- Duration, 67
- duration element, 67, 68, 68, 68, 68, 68, 68, 68
 - hours attribute, 68
 - id attribute, 68
 - minutes attribute, 68
 - months attribute, 68
 - seconds attribute, 68
 - weeks attribute, 68
 - years attribute, 68

E

- election-timeout, 11
 - Cluster Option, 11
- enable-acl, 9
 - Cluster Option, 9
- enable-startup-probes, 8
 - Cluster Option, 8
- enabled, 24, 24
 - Action Property, 24, 24
- end attribute, 66
 - date_expression element, 66
- Environment Variable
 - CRM_alert_
 - attribute_name, 62
 - attribute_value, 62
 - desc, 62
 - exec_time, 62
 - interval, 62
 - kind, 61
 - node, 62
 - nodeid, 62
 - rc, 62

- recipient, 61
- rsc, 62
- status, 62
- target_rc, 62
- task, 62
- timestamp, 62
- timestamp_epoch, 62
- timestamp_usec, 62
- version, 61
- CRM_alert_node_
 - sequence, 62
- OCF_RESKEY_CRM_meta_notify_
 - active_resource, 93
 - active_uname, 93
 - demote_resource, 94
 - demote_uname, 94
 - inactive_resource, 93
 - master_resource, 94
 - master_uname, 94
 - operation, 93
 - promote_resource, 94
 - promote_uname, 94
 - slave_resource, 94
 - slave_uname, 94
 - start_resource, 93
 - start_uname, 93
 - stop_resource, 93
 - stop_uname, 93
 - type, 92
- epoch, 6
 - Cluster Option, 6
- exec-time, 123
 - Action Status, 123
- exec_time, 62
- expected, 122
 - Node Status, 122
- expression element, 64, 65, 65, 65, 65, 65, 65
 - attribute attribute, 65
 - id attribute, 65
 - operation attribute, 65
 - type attribute, 65
 - value attribute, 65
 - value-source attribute, 65

F

- fail-count-, 13
 - Node attribute, 13
- failure-timeout, 20
 - Resource Option, 20
- feedback
 - contact information for this manual, xiii
- fence-reaction, 8
 - Cluster Option, 8

Fencing, 43, 43, 43, 43, 43, 44, 44, 44, 44, 44, 44, 45, 45, 45, 45, 45, 45, 46, 46, 46, 46, 46, 46, 47

Configuration, 41

fencing-level

devices, 52

id, 51

index, 52

target, 51

target-attribute, 52, 52

target-pattern, 52

Property

pcmk_action_limit, 44

pcmk_delay_base, 44

pcmk_delay_max, 44

pcmk_host_argument, 44

pcmk_host_check, 43

pcmk_host_list, 43

pcmk_host_map, 43

pcmk_list_action, 45

pcmk_list_retries, 46

pcmk_list_timeout, 45

pcmk_monitor_action, 46

pcmk_monitor_retries, 46

pcmk_monitor_timeout, 46

pcmk_off_action, 45

pcmk_off_retries, 45

pcmk_off_timeout, 45

pcmk_reboot_action, 44

pcmk_reboot_retries, 45

pcmk_reboot_timeout, 44

pcmk_status_action, 46

pcmk_status_retries, 47

pcmk_status_timeout, 46

provides, 43

stonith-timeout, 43, 44, 45

fencing-level, 51, 51, 52, 52, 52, 52, 52

devices, 52

id, 51

index, 52

target, 51

target-attribute, 52, 52

target-pattern, 52

first attribute, 31

rsc_order element, 31

first-action attribute, 31

rsc_order element, 31

G

globally-unique, 87

Clone Option, 87

green, 81

Group Property

id, 86

Group Resource Property, 86

Group Resources, 85

Groups, 85, 86

H

have-quorum, 6

Cluster Property, 6

host-interface attribute, 101

network element, 101

host-netmask attribute, 101

network element, 101

host_list, 78

Ping Resource Option, 78

hours attribute, 67, 68

date_spec element, 67

duration element, 68

I

id, 18, 24, 51, 86, 87, 116, 116, 117, 117, 118, 121, 122

acl_group, 117

acl_permission, 116

acl_role, 116

acl_target, 117

Action Property, 24

Action Status, 122

Clone Property, 87

fencing-level, 51

Group Resource Property, 86

Node Status, 121

Resource, 18

role, 118

id attribute, 29, 31, 33, 35, 64, 65, 66, 67, 68, 100, 101, 102

bundle element, 100

date_expression element, 66

date_spec element, 67

duration element, 68

expression element, 65

port-mapping element, 101

resource_set element, 35

rsc_colocation element, 33

rsc_location element, 29

rsc_order element, 31

rule element, 64

storage-mapping element, 102

image attribute, 100, 100, 100

docker element, 100

podman element, 100

rkt element, 100

inactive_resource, 93

Notification Environment Variable, 93

index, 52

fencing-level, 52

- interleave, 88
 - Clone Option, 88
- internal-port attribute, 102
 - port-mapping element, 102
- interval, 24, 62, 123
 - Action Property, 24
 - Action Status, 123
- in_ccm, 121
 - Node Status, 121
- ip-range-start attribute, 101
 - network element, 101
- is-managed, 19
 - Resource Option, 19

J

- join, 121
 - Node Status, 121
- join-finalization-timeout, 11
 - Cluster Option, 11
- join-integration-timeout, 11
 - Cluster Option, 11

K

- kind, 61, 116
 - acl_permission, 116
- kind attribute, 32
 - rsc_order element, 32

L

- last-failure-, 13
 - Node attribute, 13
- last-rc-change, 123
 - Action Status, 123
- last-run, 123
 - Action Status, 123
- Linux Standard Base
 - Resources, 16
- Location
 - Determine by Rules, 69
- Location Constraint, 28
 - Resource Discovery, 30
 - rsc_location element, 28
- Location Relative to Other Resources, 32
- LSB, 16
 - Resources, 16

M

- maintenance, 13, 19
 - Node attribute, 13
 - Resource Option, 19
- maintenance-mode, 8
 - Cluster Option, 8
- master_resource, 94

- Notification Environment Variable, 94
- master_uname, 94
 - Notification Environment Variable, 94
- Meta-attributes, 104
- migrate-on-red, 81
- migration-limit, 7
 - Cluster Option, 7
- migration-threshold, 20
 - Resource Option, 20
- minutes attribute, 68
 - duration element, 68
- monthdays attribute, 67
 - date_spec element, 67
- months attribute, 67, 68
 - date_spec element, 67
 - duration element, 68
- moon attribute, 67
 - date_spec element, 67
- Moving, 76
 - Resources, 76
- multiple-active, 20
 - Resource Option, 20
- multiplier, 78
 - Ping Resource Option, 78

N

- Nagios Plugins, 17
 - Resources, 17
- name, 24
 - Action Property, 24
- network attribute, 100, 100, 100
 - docker element, 100
 - podman element, 100
 - rkt element, 100
- network element, 101, 101, 101, 101, 101, 101
 - add-host attribute, 101
 - control-port attribute, 101
 - host-interface attribute, 101
 - host-netmask attribute, 101
 - ip-range-start attribute, 101
- Networking, 101
- no-quorum-policy, 7
 - Cluster Option, 7
- Node
 - attribute, 12
 - #digests-, 14
 - #node-unfenced, 14
 - fail-count-, 13
 - last-failure-, 13
 - maintenance, 13
 - probe_complete, 13
 - resource-discovery-enabled, 14
 - shutdown, 14

- site-name, 14
 - standby, 14
 - terminate, 14
 - Score, 28
 - Status, 121
 - crm-debug-origin, 121
 - crmd, 121
 - expected, 122
 - id, 121
 - in_ccm, 121
 - join, 121
 - uname, 121
 - node, 62
 - Node attribute, 13, 13, 13, 13, 14, 14, 14, 14, 14, 14
 - node attribute, 29
 - rsc_location element, 29
 - Node Attribute Expression, 64
 - Node Attributes, 104
 - Node health
 - custom, 82
 - green, 81
 - migrate-on-red, 81
 - none, 81
 - only-green, 81
 - progressive, 81
 - red, 81
 - score, 81
 - yellow, 81
 - Node Status, 121, 121, 121, 121, 121, 121, 122
 - node-attribute attribute, 33
 - rsc_colocation element, 33
 - node-health-base, 9
 - Cluster Option, 9
 - node-health-green, 10
 - Cluster Option, 10
 - node-health-red, 10
 - Cluster Option, 10
 - node-health-strategy, 9
 - Cluster Option, 9
 - node-health-yellow, 10
 - Cluster Option, 10
 - nodeid, 62
 - none, 81
 - Notification Environment Variable, 92, 93, 93, 93, 93, 93, 93, 93, 93, 94, 94, 94, 94, 94, 94, 94, 94
 - notify, 88
 - Clone Option, 88
 - num_updates, 6
 - Cluster Option, 6
- O**
- object-type, 117
 - acl_permission, 117
 - OCF, 15
 - Resources, 15
 - OCF_FAILED_MASTER, 92
 - OCF_NOT_RUNNING, 92
 - OCF_RESKEY_CRM_meta_notify_
 - active_resource, 93
 - active_uname, 93
 - demote_resource, 94
 - demote_uname, 94
 - inactive_resource, 93
 - master_resource, 94
 - master_uname, 94
 - operation, 93
 - promote_resource, 94
 - promote_uname, 94
 - slave_resource, 94
 - slave_uname, 94
 - start_resource, 93
 - start_uname, 93
 - stop_resource, 93
 - stop_uname, 93
 - type, 92
 - OCF_RUNNING_MASTER, 92
 - OCF_SUCCESS, 92
 - on-fail, 24
 - Action Property, 24
 - only-green, 81
 - op-digest, 124
 - Action Status, 124
 - op-status, 123
 - Action Status, 123
 - Open Cluster Framework
 - Resources, 15
 - operation, 93, 123
 - Action Status, 123
 - Notification Environment Variable, 93
 - operation attribute, 65, 66
 - date_expression element, 66
 - expression element, 65
 - Operation History, 122
 - Opt-In Clusters, 30
 - Opt-Out Clusters, 30
 - Option
 - admin_epoch, 6
 - batch-limit, 7
 - clone-max, 87
 - clone-min, 88
 - clone-node-max, 88
 - cluster-delay, 9
 - cluster-ipc-limit, 9
 - cluster-recheck-interval, 10
 - concurrent-fencing, 8
 - Configuration Version, 6
 - dampen, 78

- dc-deadtime, 9
- election-timeout, 11
- enable-acl, 9
- enable-startup-probes, 8
- epoch, 6
- failure-timeout, 20
- fence-reaction, 8
- globally-unique, 87
- host_list, 78
- interleave, 88
- is-managed, 19
- join-finalization-timeout, 11
- join-integration-timeout, 11
- maintenance, 19
- maintenance-mode, 8
- migration-limit, 7
- migration-threshold, 20
- multiple-active, 20
- multiplier, 78
- no-quorum-policy, 7
- node-health-base, 9
- node-health-green, 10
- node-health-red, 10
- node-health-strategy, 9
- node-health-yellow, 10
- notify, 88
- num_updates, 6
- ordered, 88
- pe-error-series-max, 9
- pe-input-series-max, 9
- pe-warn-series-max, 9
- placement-strategy, 9
- priority, 19
- priority-fencing-delay, 8
- promotable, 88
- promoted-max, 88
- promoted-node-max, 88
- remove-after-stop, 10
- requires, 20
- resource-stickiness, 19
- shutdown-escalation, 11
- shutdown-lock, 10
- shutdown-lock-limit, 10
- start-failure-is-fatal, 7
- startup-fencing, 10
- stonith-action, 8
- stonith-enabled, 8
- stonith-max-attempts, 8
- stonith-timeout, 8
- stonith-watchdog-timeout, 8
- stop-all-resources, 7
- stop-orphan-actions, 7
- stop-orphan-resources, 7
- symmetric-cluster, 7

- target-role, 19
- timeout, 59
- timestamp-format, 59
- transition-delay, 11
- validate-with, 6
- options attribute, 100, 100, 100, 102
 - docker element, 100
 - podman element, 100
 - rkt element, 100
 - storage-mapping element, 102
- ordered, 88
 - Clone Option, 88
- Ordering
 - action, 91
 - role, 90
 - rsc-role, 89
 - with-rsc-role, 89
- Ordering Constraint, 31
 - rsc_order element, 31
- Ordering Constraints, 89, 89, 90, 91

P

- pcmk_action_limit, 44
 - Fencing, 44
- pcmk_delay_base, 44
 - Fencing, 44
- pcmk_delay_max, 44
 - Fencing, 44
- pcmk_host_argument, 44
 - Fencing, 44
- pcmk_host_check, 43
 - Fencing, 43
- pcmk_host_list, 43
 - Fencing, 43
- pcmk_host_map, 43
 - Fencing, 43
- pcmk_list_action, 45
 - Fencing, 45
- pcmk_list_retries, 46
 - Fencing, 46
- pcmk_list_timeout, 45
 - Fencing, 45
- pcmk_monitor_action, 46
 - Fencing, 46
- pcmk_monitor_retries, 46
 - Fencing, 46
- pcmk_monitor_timeout, 46
 - Fencing, 46
- pcmk_off_action, 45
 - Fencing, 45
- pcmk_off_retries, 45
 - Fencing, 45
- pcmk_off_timeout, 45

- Fencing, 45
- pcm_k_reboot_action, 44
 - Fencing, 44
- pcm_k_reboot_retries, 45
 - Fencing, 45
- pcm_k_reboot_timeout, 44
 - Fencing, 44
- pcm_k_status_action, 46
 - Fencing, 46
- pcm_k_status_retries, 47
 - Fencing, 47
- pcm_k_status_timeout, 46
 - Fencing, 46
- pe-error-series-max, 9
 - Cluster Option, 9
- pe-input-series-max, 9
 - Cluster Option, 9
- pe-warn-series-max, 9
 - Cluster Option, 9
- Ping Resource
 - Option
 - dampen, 78
 - host_list, 78
 - multiplier, 78
- Ping Resource Option, 78, 78, 78
- placement-strategy, 9
 - Cluster Option, 9
- podman
 - Bundle, 99
- podman element, 100, 100, 100, 100, 100, 100, 100, 100, 100
 - image attribute, 100
 - network attribute, 100
 - options attribute, 100
 - promoted-max attribute, 100
 - replicas attribute, 100
 - replicas-per-host attribute, 100
 - run-command attribute, 100
- port attribute, 102
 - port-mapping element, 102
- port-mapping, 101
- port-mapping element, 101, 102, 102, 102
 - id attribute, 101
 - internal-port attribute, 102
 - port attribute, 102
 - range attribute, 102
- Prerequisites, 99
- Primitive, 103
- priority, 19
 - Resource Option, 19
- priority-fencing-delay, 8
 - Cluster Option, 8
- probe_complete, 13
 - Node attribute, 13
- progressive, 81
- Promotable, 87
- promotable, 88
 - Clone Option, 88
- Promotable Clone Resources, 87
- promoted-max, 88
 - Clone Option, 88
- promoted-max attribute, 100, 100, 100
 - docker element, 100
 - podman element, 100
 - rkt element, 100
- promoted-node-max, 88
 - Clone Option, 88
- promote_resource, 94
 - Notification Environment Variable, 94
- promote_uname, 94
 - Notification Environment Variable, 94
- Property
 - cib-last-written, 6
 - class, 18
 - cluster-infrastructure, 7
 - cluster-name, 7
 - dc-uuid, 6
 - dc-version, 7
 - enabled, 24, 24
 - have-quorum, 6
 - id, 18, 24, 87
 - interval, 24
 - name, 24
 - on-fail, 24
 - pcm_k_action_limit, 44
 - pcm_k_delay_base, 44
 - pcm_k_delay_max, 44
 - pcm_k_host_argument, 44
 - pcm_k_host_check, 43
 - pcm_k_host_list, 43
 - pcm_k_host_map, 43
 - pcm_k_list_action, 45
 - pcm_k_list_retries, 46
 - pcm_k_list_timeout, 45
 - pcm_k_monitor_action, 46
 - pcm_k_monitor_retries, 46
 - pcm_k_monitor_timeout, 46
 - pcm_k_off_action, 45
 - pcm_k_off_retries, 45
 - pcm_k_off_timeout, 45
 - pcm_k_reboot_action, 44
 - pcm_k_reboot_retries, 45
 - pcm_k_reboot_timeout, 44
 - pcm_k_status_action, 46
 - pcm_k_status_retries, 47
 - pcm_k_status_timeout, 46
 - provider, 18
 - provides, 43
 - role, 25

- stonith-timeout, 43, 44, 45
- timeout, 24
- type, 18
- provider, 18
 - Resource, 18
- provides, 43
 - Fencing, 43

Q

- queue-time, 123
 - Action Status, 123

R

- range attribute, 102
 - port-mapping element, 102
- rc, 62
- rc-code, 123
 - Action Status, 123
- recipient, 61
- red, 81
- reference, 117
 - acl_permission, 117
- remove-after-stop, 10
 - Cluster Option, 10
- replicas attribute, 100, 100, 100
 - docker element, 100
 - podman element, 100
 - rkt element, 100
- replicas-per-host attribute, 100, 100, 100
 - docker element, 100
 - podman element, 100
 - rkt element, 100
- require-all attribute, 35
 - resource_set element, 35
- requires, 20
 - Resource Option, 20
- Resource, 15, 18, 18, 18, 18
 - Action, 23
 - Alerts, 58
 - Bundle, 99
 - Container, 100
 - Meta-attributes, 104
 - Networking, 101
 - Node Attributes, 104
 - Prerequisites, 99
 - Primitive, 103
 - Storage, 102
 - class, 15
 - Clones, 86
 - Constraint, 28
 - Group Property
 - id, 86
 - Groups, 85

- Location
 - Determine by Rules, 69
- Location Relative to Other Resources, 32
- LSB, 16
- Moving, 76
- Nagios Plugins, 17
- OCF, 15
- Option
 - failure-timeout, 20
 - is-managed, 19
 - maintenance, 19
 - migration-threshold, 20
 - multiple-active, 20
 - priority, 19
 - requires, 20
 - resource-stickiness, 19
 - target-role, 19
- Promotable, 87
- Property
 - class, 18
 - id, 18
 - provider, 18
 - type, 18
- Resource Set, 34
- Score, 28
- Start Order, 31
- STONITH, 17
- System Services, 17
- Systemd, 16
- Upstart, 17
- Resource Discovery, 30
- Resource Option, 19, 19, 19, 19, 19, 20, 20, 20, 20
- Resource Set, 34, 34
- resource-discovery attribute, 30
 - rsc_location element, 30
- resource-discovery-enabled, 14
 - Node attribute, 14
- resource-stickiness, 19
 - Clones, 91
 - Groups, 86
 - Resource Option, 19
- Resources, 15, 15, 16, 16, 16, 17, 17, 17, 17, 76
- resource_set element, 35, 35, 35, 35, 35, 35
 - action attribute, 35
 - id attribute, 35
 - require-all attribute, 35
 - role attribute, 35
 - score attribute, 35
 - sequential attribute, 35
- Return Code
 - OCF_FAILED_MASTER, 92
 - OCF_NOT_RUNNING, 92
 - OCF_RUNNING_MASTER, 92
 - OCF_SUCCESS, 92

rkt
 Bundle, 99
rkt element, 100, 100, 100, 100, 100, 100, 100, 100
 image attribute, 100
 network attribute, 100
 options attribute, 100
 promoted-max attribute, 100
 replicas attribute, 100
 replicas-per-host attribute, 100
 run-command attribute, 100
role, 25, 90, 118
 Action Property, 25
 id, 118
 Ordering Constraints, 90
role attribute, 35, 64
 resource_set element, 35
 rule element, 64
rsc, 62
rsc attribute, 29, 33
 rsc_colocation element, 33
 rsc_location element, 29
rsc-pattern attribute, 29
 rsc_location element, 29
rsc-role, 89
 Ordering Constraints, 89
rsc_colocation element, 33, 33, 33, 33, 33, 33, 33, 33
 id attribute, 33
 node-attribute attribute, 33
 rsc attribute, 33
 score attribute, 33
 with-rsc attribute, 33
rsc_location element, 28, 28, 29, 29, 29, 29, 29, 30
 id attribute, 29
 node attribute, 29
 resource-discovery attribute, 30
 rsc attribute, 29
 rsc-pattern attribute, 29
 score attribute, 29
rsc_order element, 31, 31, 31, 31, 31, 31, 31, 32, 32
 first attribute, 31
 first-action attribute, 31
 id attribute, 31
 kind attribute, 32
 symmetrical attribute, 32
 then attribute, 31
 then-action attribute, 31
Rule, 64
 Controlling Cluster Options, 72
 Date/Time Expression, 66
 Date Specification, 67
 Duration, 67
 Determine Resource Location, 69
 Node Attribute Expression, 64
rule element, 64, 64, 64, 64, 64, 64

 boolean-op attribute, 64
 id attribute, 64
 role attribute, 64
 score attribute, 64
 score-attribute attribute, 64
run-command attribute, 100, 100, 100
 docker element, 100
 podman element, 100
 rkt element, 100

S

Score, 28, 28
score, 81
score attribute, 29, 33, 35, 64
 resource_set element, 35
 rsc_colocation element, 33
 rsc_location element, 29
 rule element, 64
score-attribute attribute, 64
 rule element, 64
seconds attribute, 68
 duration element, 68
sequence, 62
sequential attribute, 35
 resource_set element, 35
Setting Options with Rules, 72
shutdown, 14
 Node attribute, 14
shutdown-escalation, 11
 Cluster Option, 11
shutdown-lock, 10
 Cluster Option, 10
shutdown-lock-limit, 10
 Cluster Option, 10
site-name, 14
 Node attribute, 14
slave_resource, 94
 Notification Environment Variable, 94
slave_uname, 94
 Notification Environment Variable, 94
source-dir attribute, 102
 storage-mapping element, 102
source-dir-root attribute, 102
 storage-mapping element, 102
standby, 14
 Node attribute, 14
start attribute, 66
 date_expression element, 66
Start Order, 31
start-failure-is-fatal, 7
 Cluster Option, 7
startup-fencing, 10
 Cluster Option, 10

- start_resource, 93
 - Notification Environment Variable, 93
- start_uname, 93
 - Notification Environment Variable, 93
- status, 62
- Status, 121
 - call-id, 123
 - crm-debug-origin, 121, 124
 - crmd, 121
 - crm_feature_set, 123
 - exec-time, 123
 - expected, 122
 - id, 121, 122
 - interval, 123
 - in_ccm, 121
 - join, 121
 - last-rc-change, 123
 - last-run, 123
 - op-digest, 124
 - op-status, 123
 - operation, 123
 - queue-time, 123
 - rc-code, 123
 - transition-key, 123
 - transition-magic, 123
 - uname, 121
- Status of a Node, 121
- STONITH, 17
 - Configuration, 41
 - Resources, 17
- stonith-action, 8
 - Cluster Option, 8
- stonith-enabled, 8
 - Cluster Option, 8
- stonith-max-attempts, 8
 - Cluster Option, 8
- stonith-timeout, 8, 43, 44, 45
 - Cluster Option, 8
 - Fencing, 43, 44, 45
- stonith-watchdog-timeout, 8
 - Cluster Option, 8
- stop-all-resources, 7
 - Cluster Option, 7
- stop-orphan-actions, 7
 - Cluster Option, 7
- stop-orphan-resources, 7
 - Cluster Option, 7
- stop_resource, 93
 - Notification Environment Variable, 93
- stop_uname, 93
 - Notification Environment Variable, 93
- Storage, 102
- storage element, 102
- storage-mapping element, 102, 102, 102, 102, 102, 102

- id attribute, 102
- options attribute, 102
- source-dir attribute, 102
- source-dir-root attribute, 102
- target-dir attribute, 102
- symmetric-cluster, 7
 - Cluster Option, 7
- symmetrical attribute, 32
 - rsc_order element, 32
- Symmetrical Clusters, 30
- System Service
 - Resources, 17
- System Services, 17
- Systemd, 16
 - Resources, 16

T

- target, 51
 - fencing-level, 51
- target-attribute, 52, 52
 - fencing-level, 52, 52
- target-dir attribute, 102
 - storage-mapping element, 102
- target-pattern, 52
 - fencing-level, 52
- target-role, 19
 - Resource Option, 19
- target_rc, 62
- task, 62
- terminate, 14
 - Node attribute, 14
- then attribute, 31
 - rsc_order element, 31
- then-action attribute, 31
 - rsc_order element, 31
- timeout, 24, 59
 - Action Property, 24
- timestamp, 62
- timestamp-format, 59
- timestamp_epoch, 62
- timestamp_usec, 62
- transition-delay, 11
 - Cluster Option, 11
- transition-key, 123
 - Action Status, 123
- transition-magic, 123
 - Action Status, 123
- type, 18, 92
 - Notification Environment Variable, 92
 - Resource, 18
- type attribute, 65
 - expression element, 65

U

- uname, 121
 - Node Status, 121
- Upstart, 17
 - Resources, 17

V

- validate-with, 6
 - Cluster Option, 6
- value attribute, 65
 - expression element, 65
- value-source attribute, 65
 - expression element, 65
- version, 61

W

- weekdays attribute, 67
 - date_spec element, 67
- weeks attribute, 67, 68
 - date_spec element, 67
 - duration element, 68
- weekyears attribute, 67
 - date_spec element, 67
- with-rsc attribute, 33
 - rsc_colocation element, 33
- with-rsc-role, 89
 - Ordering Constraints, 89

X

- XML attribute
 - action attribute
 - resource_set element, 35
 - attribute attribute
 - expression element, 65
 - boolean-op attribute
 - rule element, 64
 - description attribute
 - bundle element, 100
 - end attribute
 - date_expression element, 66
 - first attribute
 - rsc_order element, 31
 - first-action attribute
 - rsc_order element, 31
 - hours attribute
 - date_spec element, 67
 - duration element, 68
 - id attribute
 - bundle element, 100
 - date_expression element, 66
 - date_spec element, 67
 - duration element, 68
 - expression element, 65

- port-mapping element, 101
- resource_set element, 35
- rsc_colocation element, 33
- rsc_location element, 29
- rsc_order element, 31
- rule element, 64
- storage-mapping element, 102
- image attribute
 - docker element, 100
 - podman element, 100
 - rkt element, 100
- internal-port attribute
 - port-mapping element, 102
- kind attribute
 - rsc_order element, 32
- minutes attribute
 - duration element, 68
- monthdays attribute
 - date_spec element, 67
- months attribute
 - date_spec element, 67
 - duration element, 68
- moon attribute
 - date_spec element, 67
- network attribute
 - docker element, 100
 - podman element, 100
 - rkt element, 100
- network element
 - add-host attribute, 101
 - control-port attribute, 101
 - host-interface attribute, 101
 - host-netmask attribute, 101
 - ip-range-start attribute, 101
- node attribute
 - rsc_location element, 29
- node-attribute attribute
 - rsc_colocation element, 33
- operation attribute
 - date_expression element, 66
 - expression element, 65
- options attribute
 - docker element, 100
 - podman element, 100
 - rkt element, 100
 - storage-mapping element, 102
- port attribute
 - port-mapping element, 102
- promoted-max attribute
 - docker element, 100
 - podman element, 100
 - rkt element, 100
- range attribute
 - port-mapping element, 102

- replicas attribute
 - docker element, 100
 - podman element, 100
 - rkt element, 100
- replicas-per-host attribute
 - docker element, 100
 - podman element, 100
 - rkt element, 100
- require-all attribute
 - resource_set element, 35
- resource-discovery attribute
 - rsc_location element, 30
- role attribute
 - resource_set element, 35
 - rule element, 64
- rsc attribute
 - rsc_colocation element, 33
 - rsc_location element, 29
- rsc-pattern attribute
 - rsc_location element, 29
- run-command attribute
 - docker element, 100
 - podman element, 100
 - rkt element, 100
- score attribute
 - resource_set element, 35
 - rsc_colocation element, 33
 - rsc_location element, 29
 - rule element, 64
- score-attribute attribute
 - rule element, 64
- seconds attribute
 - duration element, 68
- sequential attribute
 - resource_set element, 35
- source-dir attribute
 - storage-mapping element, 102
- source-dir-root attribute
 - storage-mapping element, 102
- start attribute
 - date_expression element, 66
- symmetrical attribute
 - rsc_order element, 32
- target-dir attribute
 - storage-mapping element, 102
- then attribute
 - rsc_order element, 31
- then-action attribute
 - rsc_order element, 31
- type attribute
 - expression element, 65
- value attribute
 - expression element, 65
- value-source attribute
 - expression element, 65
- weekdays attribute
 - date_spec element, 67
- weeks attribute
 - date_spec element, 67
 - duration element, 68
- weekyears attribute
 - date_spec element, 67
- with-rsc attribute
 - rsc_colocation element, 33
- yeardays attribute
 - date_spec element, 67
- years attribute
 - date_spec element, 67
 - duration element, 68
- XML element
 - add-host attribute
 - network element, 101
 - bundle element, 99
 - description attribute, 100
 - id attribute, 100
 - control-port attribute
 - network element, 101
 - date_expression element, 66
 - end attribute, 66
 - id attribute, 66
 - operation attribute, 66
 - start attribute, 66
 - date_spec element, 67
 - hours attribute, 67
 - id attribute, 67
 - monthdays attribute, 67
 - months attribute, 67
 - moon attribute, 67
 - weekdays attribute, 67
 - weeks attribute, 67
 - weekyears attribute, 67
 - yeardays attribute, 67
 - years attribute, 67
 - docker element, 100
 - image attribute, 100
 - network attribute, 100
 - options attribute, 100
 - promoted-max attribute, 100
 - replicas attribute, 100
 - replicas-per-host attribute, 100
 - run-command attribute, 100
 - duration element, 67
 - hours attribute, 68
 - id attribute, 68
 - minutes attribute, 68
 - months attribute, 68
 - seconds attribute, 68
 - weeks attribute, 68

- years attribute, 68
- expression element, 64
 - attribute attribute, 65
 - id attribute, 65
 - operation attribute, 65
 - type attribute, 65
 - value attribute, 65
 - value-source attribute, 65
- host-interface attribute
 - network element, 101
- host-netmask attribute
 - network element, 101
- ip-range-start attribute
 - network element, 101
- network element, 101
- podman element, 100
 - image attribute, 100
 - network attribute, 100
 - options attribute, 100
 - promoted-max attribute, 100
 - replicas attribute, 100
 - replicas-per-host attribute, 100
 - run-command attribute, 100
- port-mapping, 101
- port-mapping element
 - id attribute, 101
 - internal-port attribute, 102
 - port attribute, 102
 - range attribute, 102
- resource_set element
 - action attribute, 35
 - id attribute, 35
 - require-all attribute, 35
 - role attribute, 35
 - score attribute, 35
 - sequential attribute, 35
- rkt element, 100
 - image attribute, 100
 - network attribute, 100
 - options attribute, 100
 - promoted-max attribute, 100
 - replicas attribute, 100
 - replicas-per-host attribute, 100
 - run-command attribute, 100
- rsc_colocation element, 33
 - id attribute, 33
 - node-attribute attribute, 33
 - rsc attribute, 33
 - score attribute, 33
 - with-rsc attribute, 33
- rsc_location element, 28
 - id attribute, 29
 - node attribute, 29
 - resource-discovery attribute, 30

- rsc attribute, 29
- rsc-pattern attribute, 29
- score attribute, 29
- rsc_order element, 31
 - first attribute, 31
 - first-action attribute, 31
 - id attribute, 31
 - kind attribute, 32
 - symmetrical attribute, 32
 - then attribute, 31
 - then-action attribute, 31
- rule element, 64
 - boolean-op attribute, 64
 - id attribute, 64
 - role attribute, 64
 - score attribute, 64
 - score-attribute attribute, 64
- storage element, 102
- storage-mapping element, 102
 - id attribute, 102
 - options attribute, 102
 - source-dir attribute, 102
 - source-dir-root attribute, 102
 - target-dir attribute, 102
- xpath, 117
 - acl_permission, 117

Y

- yeardays attribute, 67
 - date_spec element, 67
- years attribute, 67, 68
 - date_spec element, 67
 - duration element, 68
- yellow, 81